

VMware Tanzu Kubernetes Grid Integrated Edition Validation Guide



vmware®

You can find the most up-to-date technical documentation on the VMware website at:

<https://docs.vmware.com/>

If you have comments about this documentation, submit your feedback to

docfeedback@vmware.com

VMware, Inc.
3401 Hillview Ave.
Palo Alto, CA 94304
www.vmware.com

Copyright © 2020 VMware, Inc. All rights reserved. [Copyright and trademark information.](#)

Contents

About This Guide	6
1 Introduction	7
Validation Prerequisites	7
Validation Process	7
Revalidation Requirements	9
Validation Workflow	9
Validation Support	10
2 Accessing the Lab	12
Accessing VMware Integration Validation Service	12
Creating a Certification Test Session	13
Lab Selection	14
Main Console	15
Top Menu Bar	15
3 TKGI Validation Tests Overview	18
4 Application Packaging and Conformance Validation	19
Compliance and Security Tests	19
PKS.Security.Image.Sign.PushToHarbor	19
PKS.Security.Image.Harbor.Scan	23
PKS.Security.Image.Harbor.Pull	24
PKS.Security.Image.Harbor.Delete	26
PKS.Security.Image.Harbor.Modify	29
PKS.Security.Image.DockerBenchSecurity	31
PKS.Security.Image.Setuid.Setgid	33
PKS.Security.SignedImage.Cluster	35
5 Application Life Cycle Management	39
Operational Tests	39
PKS.LCM.Image.Build	39
PKS.LCM.App.Install	41
PKS.LCM.App.Uninstall	43
Kubernetes Life Cycle Tests	44
K8s.Deployment.Create	45
K8s.Deployment.Delete.Deploy	47
K8s.Service.Deploy	50

	K8s.Service.Delete.Expose	52
	K8s.ReplicaSet.Scale	54
	K8s.Pod.Delete	56
6	Partner Validation	59
	Functionality Tests	59
	App.Functionality (Partner Validation)	59
7	TKGI Centric Platform Validation	61
	Resilience/Negative Tests	61
	PKS.Resilience.ClusterExpand	61
	PKS.Worker.Reboot	64
	PKS.Master.Reboot	68
	PKS.Worker.PowerOffOn	71
	PKS.Master.PowerOffOn	75
8	Collecting Logs	79
9	Submitting Validation Results	83
10	Requesting Partner-Ready VMware Tanzu Logo	84
11	Working with the Lab	85
	Utilities Installed on the Lab	85
	Downloading and Transferring Files	86
	Accessing CLI VM	87
	Accessing TKGI Environment	87
	Creating a PKS Cluster	88
	Checking PKS Cluster Status	88
	Validating Nodes for Cluster	89
	Installing Helm Package Manager	90
	Accessing Harbor Repository	91
	Pushing Images into Harbor	91
	Pushing Securely Signed Images to Harbor	93
	Steps to Add Storage Class	97
	Accessing vSphere Client	98
	Load Balancers	100
	Adding Persistent Storage	100
	Adding Unmounted Block Devices to the Worker and Master Nodes	100
	Adding Persistent Storage to a Pod Using Persistent Volume Claim	102
	Dynamic Persistent Volumes	102

[Statically Allocated Persistent Volumes](#) 104

12 Tips and Troubleshooting 106

[Troubleshooting PKS Cluster Creation](#) 107

[Restart PKS Instance](#) 107

[Get UAAC Token and Log In to PKS](#) 108

[Accessing PCF Ops Manager](#) 110

[Accessing NSX-T Manager](#) 111

[Accessing Kubernetes Dashboard](#) 112

13 Appendix A: Sample Docker Bench Security Script Output 115

14 Appendix B: TKGI Software BOM 119

[TKGI Software BOM – 1.7 Lab](#) 119

[TKGI Software BOM – 1.8 Lab](#) 119

15 Appendix C: TKGI Cluster Plans 121

About This Guide

The *VMware Tanzu Kubernetes Grid Integrated Edition (TKGI) Validation Guide* provides information about validation of the CNA (Cloud Native Application) solution running in the TKGI environment on vSphere.

Intended Audience

This guide is for experienced Windows and Linux system engineers who are familiar with VMware virtual machine technologies, data center operations, and Tanzu Kubernetes Grid Integrated Edition (TKGI). The user must have familiarity with deploying VC and ESXi, and installing and configuring required software and services.

VMware Technical Publications Glossary

VMware Technical Publications provides a glossary of terms that might be unfamiliar to you. For definitions of terms as they are used in the VMware technical documentation, go to <https://www.vmware.com/support/pubs/>.

VMware Documentation

Related VMware product documentation that might be useful can be found at on the VMware website at <http://www.vmware.com/support/pubs>.

Kubernetes Documentation

Related Kubernetes product documentation that might be useful can be found at on the Kubernetes website at <https://kubernetes.io/docs/tutorials/kubernetes-basics/deploy-app/deploy-interactive/> and <https://kubernetes.io/blog/2018/05/01/developing-on-kubernetes/>.

Document Feedback

VMware welcomes your suggestions for improving our documentation. If you have comments, submit your feedback through the VMware {code} site (<https://code.vmware.com/certifications>).

Introduction

1

The TKGI ISV Validation Program provides a pre-configured and ready-to-use validation testbed in the cloud.

The objective of this validation program is to ensure that the Independent Software Vendor (ISV) CNA (Cloud Native Application) solution has passed a detailed evaluation and testing process maintained by VMware for listing on the VMware Marketplace at <https://marketplace.vmware.com>.

Listing Application on VMware Marketplace as a validated solution provides customer assurance that the application is vetted on TKGI.

This chapter includes the following topics:

- [Validation Prerequisites](#)
- [Validation Process](#)
- [Revalidation Requirements](#)
- [Validation Workflow](#)
- [Validation Support](#)

Validation Prerequisites

Ensure that the following prerequisites are met before starting the TKGI validation:

- You must be enrolled in the VMware Technology Alliance Partner (TAP) program at any level.
- You must assure that the Application is validated (tested) on the Kubernetes platform.
- Confirm that the Application meets security requirements by filing the *TKGI Validation-in-Cloud - Security Intake Questionnaire*.
- Confirm that the Application Cluster requirements meet the VMware Learning Platform (VLP) lab resources by filing the *TKGI Lab - Kubernetes Environment Intake Questionnaire*.

Validation Process

This section lists the validation process steps that you must perform to get your solution validated.

Procedure

1 Create a DCPN request.

Before validating your application, answer the following questionnaires and submit them by opening a Developer Center Partner Network (DCPN) request:

- TKGI Validation-in-Cloud - Security Intake Questionnaire
- TKGI Lab - Kubernetes Environment Intake Questionnaire

Both questionnaires are available on the **VMware {code}** site > **Certification** > **PKS Partner Application Certification**, under **Documentation and Reference** (<https://code.vmware.com/group/cert/1.0/pks-partner-application-certification>).

2 VMware reviews and approves the request.

After the documentation from the preceding step is reviewed and approved, you are notified on the DCPN case to go ahead with access to the cloud environment to proceed with your validation.

3 Access the Validation-in-Cloud.

To begin, log in to the VMware {code} site (<https://code.vmware.com>) using your My VMware or Partner Central credentials. For more information, see [Chapter 2 Accessing the Lab](#).

4 Deploy your Application solution and run the Partner tests.

After deployment, the Application solution is powered on and configured (see [Accessing TKGI Environment](#) and [Operational Tests](#)). Perform an initial sanity test by logging in to the solution and validating its functionality (see [Chapter 6 Partner Validation](#)). Ensure that the IP address and host names are configured appropriately. Before running the VMware Validation Tests, you run a set of tests that are specific to your Application solution. It ensures that your solution operates normally before running the VMware Validation Tests.

5 Run the VMware Validation tests.

After configuring your solution for operation and running specific tests to determine your solution's stability, run the VMware Validation Tests.

6 Collect the logs and submit validation results.

After the validation tests are run successfully, collect the logs (see [Chapter 8 Collecting Logs](#) for details), and submit the validation results with the Validation Session ID and *TKGI Validation-in-Cloud - Validation Tests Check List* (see [Chapter 9 Submitting Validation Results](#) for details) in the DCPN case.

Ensure that you do not decommission the Validation-in-Cloud test environment for at least 24 hours. VMware updates the DCPN case with test completion status or instructions for any follow-up items.

Also, ensure that the Application solution is deployed and running in a good state after the validation tests are done. VMware might need to verify the Application solution deployment during the validation results review.

If the tests are successful, you are instructed on getting the Partner Ready VMware Tanzu Logo. At this time, the Validation-In-Cloud environment can be decommissioned by clicking the **FINISH** button. If the tests are unsuccessful, rerun the tests and submit new results.

7 Decommission Validation-in-Cloud.

The Validation-in-Cloud environment is automatically decommissioned after the lease expires. All server, networking, and storage configurations associated with the lab that are built on IaaS are deleted.

8 Get Partner Ready VMware Tanzu Logo.

For more information, see [Chapter 10 Requesting Partner-Ready VMware Tanzu Logo](#).

Revalidation Requirements

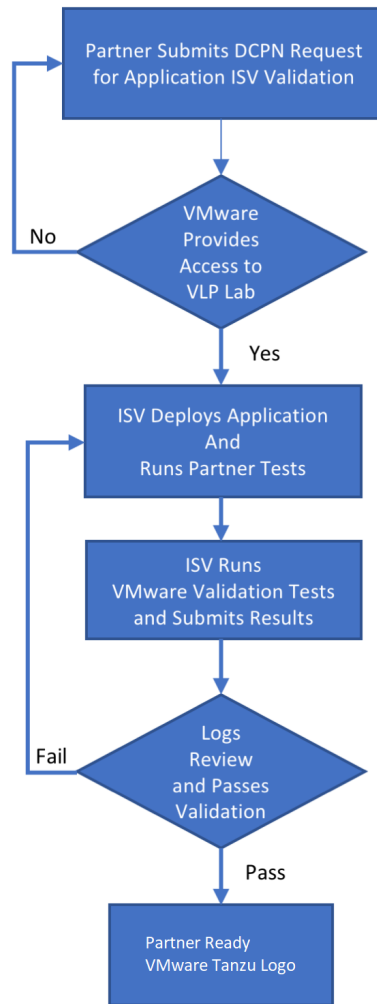
VMware generally requires revalidation for:

- Every major TKGI version update
- Every major partner application version update

VMware pre-announces TKGI version revalidation requirements as much as possible to allow partner resource planning.

Validation Workflow

The following figure illustrates the TKGI validation workflow:

Figure 1-1. TKGI Validation Workflow

Validation Support

To report issues with the validation environment or to ask questions regarding the validation, create a case using Developer Center Partner Network (DCPN).

Note You must create a new case for every unique issue.

To create a DCPN case, perform the following steps:

- 1 On the DCPN, navigate to **Cases**.
- 2 Click **New** and provide the necessary information:
 - **Request Type:** Program Request
 - **Project:** priv-<company>-tanzu_PR

Where <company> is your company name in the DCPN.

- Under **Program Information**:
 - **Summary**: Enter a short descriptive summary about the issue reported.
 - **Description**: Enter a detailed description about the issue or the questions being asked. The description must match the summary of the case as closely as possible.

Avoid case descriptions that are not descriptive, such as “I cannot do X. Why not?”.
 - Under **VMware Details**:
 - **Sub Status**: **Waiting on VMware**.
- 3 Ensure to attach any supplemental documentation, screenshots, or files to the case (if necessary) in the **Case Files** section, which is located towards the bottom of the case page.

Accessing the Lab

2

This chapter provides information about accessing the VMware Validation Lab from the VMware Integration Validation (VIVa) platform.

After starting the lab, you get access to the control center Windows VM of the lab. The control center VM acts as the jump host to the lab and from here you can access the pre-configured TKGI environment.

This chapter includes the following topics:

- [Accessing VMware Integration Validation Service](#)
- [Creating a Certification Test Session](#)
- [Lab Selection](#)
- [Main Console](#)
- [Top Menu Bar](#)

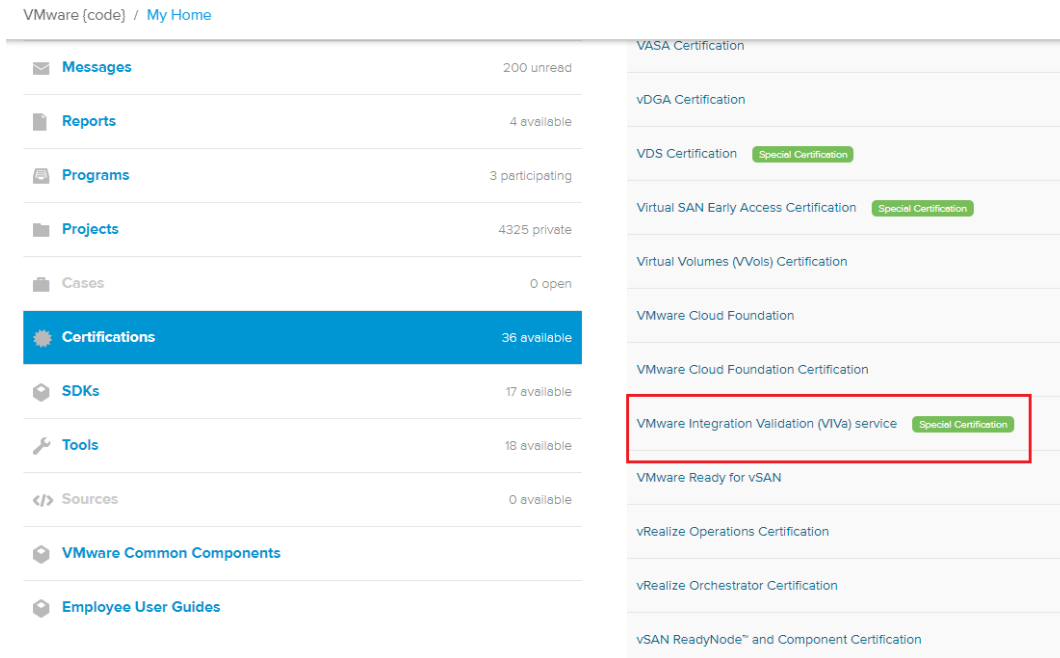
Accessing VMware Integration Validation Service

To access the VMware Integration Validation service on a web browser, perform the following steps:

Procedure

- 1 Go to the VMware Code site (<https://code.vmware.com>) and log in using your **My VMware** or **Partner Central** credentials.
- 2 Click **MY HOME > Certifications**.

3 Under **My Certifications**, scroll down and click **VMware Integration Validation (VIVA) service**.



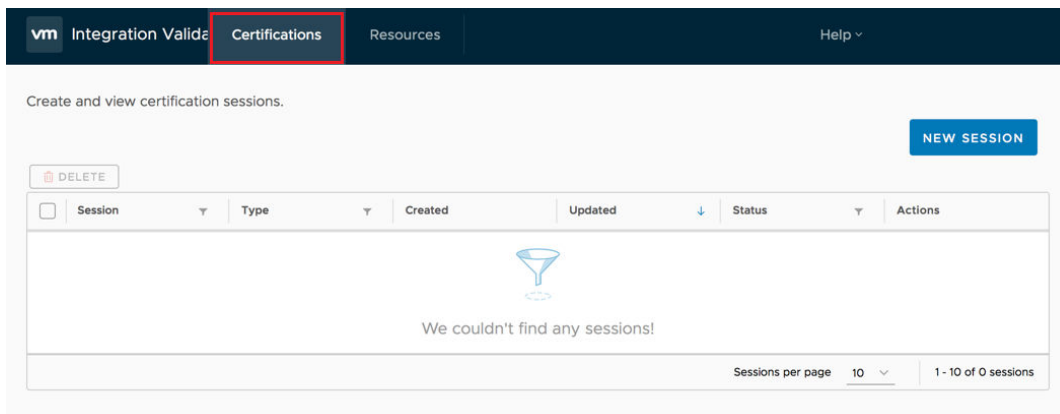
Creating a Certification Test Session

To create a certification test session for TKGI validation, perform the following steps:

Note If you are a first-time user, you must accept the End User License Agreement to proceed.

Procedure

- 1 On the VMware Integration Validation landing page, click **Certifications**.



- 2 Click **NEW SESSION**.
- 3 From the **Certification Type** drop-down menu, select **Tanzu Kubernetes Grid Integrated Edition**.
- 4 From the **VMware Release** drop-down menu, select the PKS lab version.

5 Click **CREATE**.

Lab Selection

After creating a TKGI validation session, you must select the lab in the **Lab Selection** page.

Important Make sure that you note down the string in the URL after `https://cert.vmware.com/caas/cert/`. This string is the VIVA Session ID. You must provide this ID in the DCPN case, with the logs after validation is completed.

1 Click **CREATE LAB**.

A cloud-based lab environment is created on the VMware Learning Platform (VLP).

Note The lab environment creation might take some time if the lab resources are not immediately available for the queued request.

2 Click **OPEN LAB**.

A new browser tab opens on the VLP site where the session is automatically authenticated with the VMware {code} user account.


Note Ensure that the pop-up blocker on your browser is disabled.

Main Console

The Main Console is where you start working on the VMware Validation Lab. Ensure that the **Lab Status** is displayed as **Ready** on the desktop of the Main Console before you start.

- The Main Console Windows Client password is *VMware1!*
- Use the Toggle Full-Screen and Maximize options to expand the Main Console window.
- Copy and Paste operations between the console window and your desktop are disabled for security reasons. However, you can use the **SEND TEXT** button on the top left of the console to send text from your desktop to the VM.
- If you are a first-time user of the system, click the Guide icon for a guided tour of the console interface.

Figure 2-1. Lab Main Console




Important After you enroll in the Validation-in-Cloud program, you are provided five days to complete the validation process. If you require additional time to complete the process, you must open a request for an extension of the TKGI lab 24 hours before the Validation-in-Cloud expires.

Top Menu Bar

The top menu bar on the console displays the time remaining to complete the validation process.

The menu bar has the following buttons:


[▼ HELP](#) | [PRIVACY](#) | [MY PROFILE](#) | [LOG OUT](#) | [ENGLISH](#) 

EXIT

FINISH

RESET

TIME REMAINING: 5 DAYS 47 MINUTES



- **EXIT:** Click to exit your lab. A pop-up window appears to confirm your action. If you click **CONFIRM**, the VMware Learning Platform (VLP) considers that you want to exit the lab, and returns to the VMware Learning Platform (VLP) lab console login page. Click **CANCEL** to resume your lab session.

|| ARE YOU SURE YOU WANT TO EXIT THIS LAB?

You are about to leave this lab. You can return to finish at a later time, your VMs will remain in a running state, but you should save any unsaved work. Are you sure you want to exit this lab?

CANCEL

CONFIRM

- **FINISH:** Click to complete your lab session. A pop-up window appears to confirm your action. If you click **CONFIRM**, the VMware Learning Platform (VLP) considers that you have completed the lab and returns to the VMware Learning Platform (VLP) lab console login page. You are not allowed to resume your lab after confirming the **FINISH** action. Click **CANCEL** to resume your lab session.

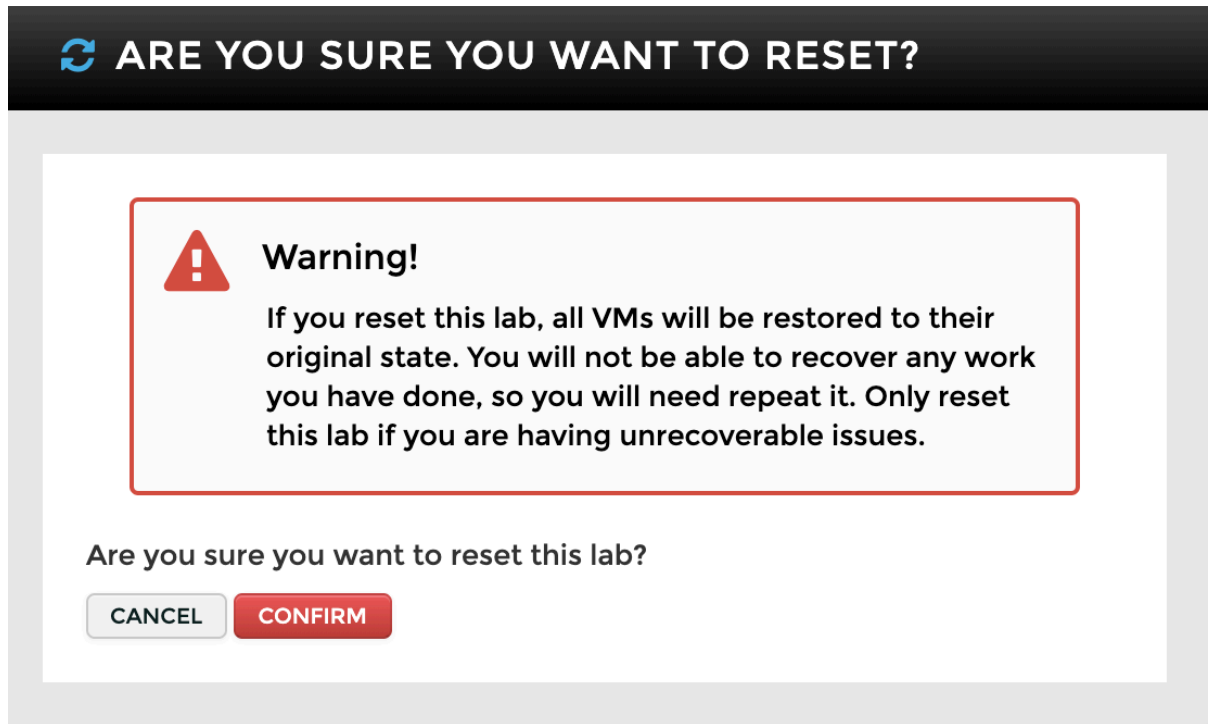
✔ ARE YOU SURE YOU WANT TO FINISH YOUR LAB?

Caution! Once you end the lab, you will not be able to return. If you would like to exit so you can return to the lab at a later time, please choose to Exit the lab instead. Are you sure you want to finish this lab and submit for grading?

CANCEL

CONFIRM

- **RESET:** Click to reset your lab. This action restores all your virtual machines to their original state. You cannot recover any work after you click **CONFIRM**.



For more information about using the lab console effectively, see [Chapter 12 Tips and Troubleshooting](#).

TKGI Validation Tests Overview

3

The TKGI validation tests are categorized into the following test areas:

- Application Packaging and Conformance Validation ([Chapter 4 Application Packaging and Conformance Validation](#))
- Application Life Cycle Management ([Chapter 5 Application Life Cycle Management](#))
- Partner Validation ([Chapter 6 Partner Validation](#))
- TKGI Centric Platform Validation ([Chapter 7 TKGI Centric Platform Validation](#))

Note Google sample images are used as examples in the commands in this guide.

Application Packaging and Conformance Validation

4

This test area focuses on the Harbor push, pull, signing, and scanning operations.

This chapter includes the following topics:

- [Compliance and Security Tests](#)

Compliance and Security Tests

The Compliance and Security Tests group has the following test cases:

- [PKS.Security.Image.Sign.PushToHarbor](#)
- [PKS.Security.Image.Harbor.Scan](#)
- [PKS.Security.Image.Harbor.Pull](#)
- [PKS.Security.Image.Harbor.Delete](#)
- [PKS.Security.Image.Harbor.Modify](#)
- [PKS.Security.Image.DockerBenchSecurity](#)
- [PKS.Security.Image.Setuid.Setgid](#)
- [PKS.Security.SignedImage.Cluster](#)

PKS.Security.Image.Sign.PushToHarbor

This test verifies that the image can be signed and pushed to Harbor with signature recorded in notary. Signing provides protection against any container tampering that can happen in transport.

Note This test is optional.

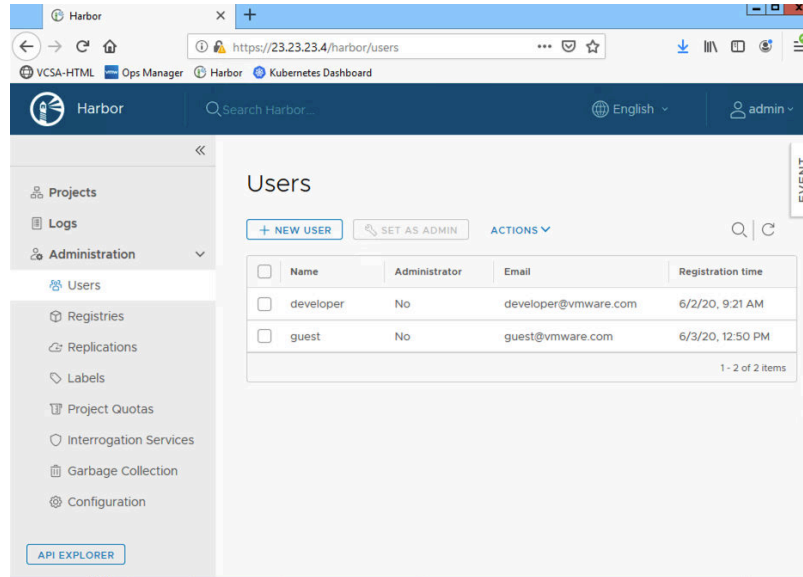
Optional	
Test ID	PKS.Security.Image.Sign.PushToHarbor
Configuration	Follow the steps detailed in Accessing TKGI Environment .

Test ID	PKS.Security.Image.Sign.PushToHarbor
Configuration	Follow the steps detailed in Accessing TKGI Environment .

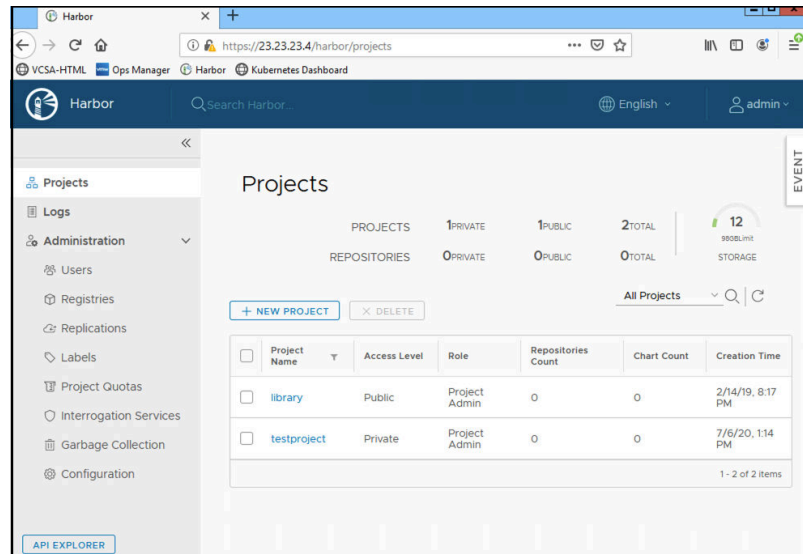
Optional

Test Script

- 1 Open Firefox and click the **Harbor** bookmark (<https://23.23.23.4>).
Log in using credentials (user name: *admin* ; password: *VMware1!*).
- 2 Under the **Users** tab, create two users - *developer* and *guest*.



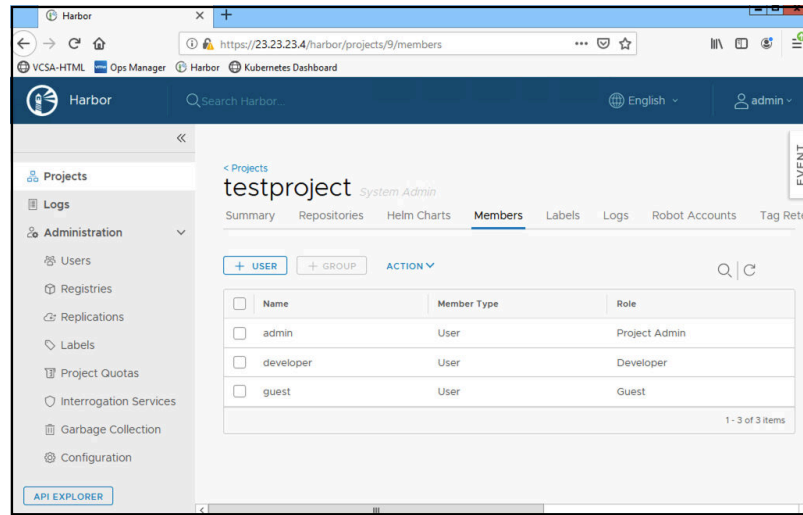
- 3 Under the **Projects** tab, create a new project - *testproject*.



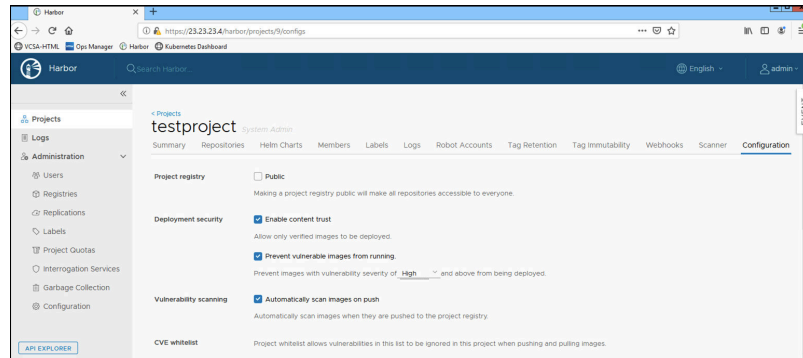
- 4 Go to **Projects > testproject**. Under the **Members** tab, click the **+User** button and add the *developer* and the *guest* users to the project.

Ensure that the *developer* is assigned the Developer role and *guest* is assigned the Guest role.

Optional



- 5 Go to the **Configuration** tab, set the values as follows, and click **SAVE**:
 - Select **Enable Content Trust**.
 - Select **Prevent vulnerable images from running** and set the severity to "High".
 - Select **Automatically scan images on push**.



- 6 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 7 Pull your Application image from your repository (Google sample images are used as examples).

```
root@cli-vm:~# docker pull gcr.io/google-samples/gb-frontend:v4
root@cli-vm:~# docker pull gcr.io/google_containers/redis:e2e
root@cli-vm:~# docker pull gcr.io/google-samples/gb-redis-slave:v1
```

- 8 Log in to Harbor registry as *developer*.

```
root@cli-vm:~# docker login harbor.corp.local -u developer
```

- 9 Set up two environment variables and tag the images (Google sample images are used as examples).

```
root@cli-vm:~# export DOCKER_CONTENT_TRUST=1
root@cli-vm:~# export DOCKER_CONTENT_TRUST_SERVER=https://harbor.corp.local:4443
root@cli-vm:~# docker tag gcr.io/google-samples/gb-frontend:v4
harbor.corp.local/testproject/gb-frontend:v4
root@cli-vm:~# docker tag gcr.io/google-samples/gb-redis-slave:v1
harbor.corp.local/testproject/gb-redis-slave:v1
```

Optional

```
root@cli-vm:~# docker tag gcr.io/google_containers/redis:e2e harbor.corp.local/
testproject/redis:e2e
```

- 10 Push the images into the Harbor repository. You are asked to enter a root key passphrase. You must enter the passphrase (*VMware1!*) every time you sign the image.

```
root@cli-vm:~# docker push harbor.corp.local/testproject/gb-frontend:v4
```

```
root@cli-vm:~# docker push harbor.corp.local/testproject/gb-redisslave:v1
```

```
root@cli-vm:~# docker push harbor.corp.local/testproject/redis:e2e
```

- 11 Verify that the tags are marked as Signed in the Harbor UI. Go to **Projects > Repositories > testproject/gb-redisslave** (testproject/gb-frontend, or testproject/redis).

Test Output

- 1 Images are successfully pushed in to Harbor (see Test script step 10).

```
root@cli-vm:~# docker push harbor.corp.local/testproject/gb-frontend:v4
```

The push refers to repository [harbor.corp.local/testproject/gb-frontend]

```
3a31f3bf94a2: Layer already exists
cdc990c9b585: Layer already exists
...
816f1903c60f: Layer already exists
c12ecfd4861d: Layer already exists
v4: digest:
sha256:aaa5b327ef3b4cb705513ab674fa40df66981616950c7de4912a621f9ee03dd4
size: 6968
Signing and pushing trust metadata
Enter passphrase for root key with ID ec1a991: VMware1!
Repeat passphrase for root key with ID ec1a991: VMware1!
Enter passphrase for new repository key with ID 7284e73: VMware1!
Repeat passphrase for new repository key with ID 7284e73: VMware1!
Finished initializing "harbor.corp.local/testproject/gb-frontend"
Successfully signed harbor.corp.local/testproject/gb-frontend:v4
```

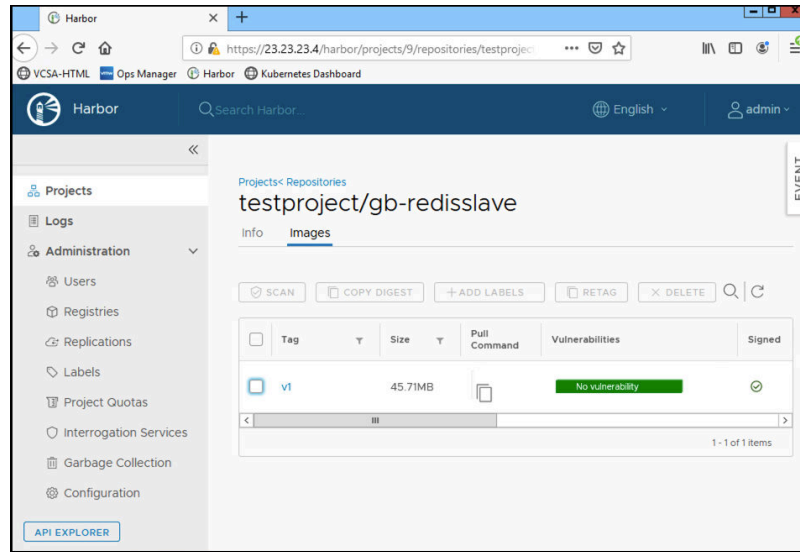
```
root@cli-vm:~# docker push harbor.corp.local/testproject/gb-redisslave:v1
```

```
...
Signing and pushing trust metadata
Enter passphrase for root key with ID ec1a991: VMware1!
Enter passphrase for new repository key with ID 63e9a04: VMware1!
Repeat passphrase for new repository key with ID 63e9a04: VMware1!
Finished initializing "harbor.corp.local/testproject/gb-redisslave"
Successfully signed harbor.corp.local/testproject/gb-redisslave:v1
```

```
root@cli-vm:~# docker push harbor.corp.local/testproject/redis:e2e
```

```
...
Signing and pushing trust metadata
Enter passphrase for root key with ID ec1a991: VMware1!
Enter passphrase for new repository key with ID e7477bb: VMware1!
Repeat passphrase for new repository key with ID e7477bb: VMware1!
Finished initializing "harbor.corp.local/testproject/redis"
Successfully signed harbor.corp.local/testproject/redis:e2e
```

- 2 The tags (v4, v1, e2e in the example) are marked as Signed in the Harbor UI. Go to **Projects > Repositories > testproject/gb-redisslave** (testproject/gb-frontend, or testproject/redis).

Optional**PKS.Security.Image.Harbor.Scan**

This test verifies that the image does not have any vulnerability or virus by scanning the image. Vulnerability scanning and Notary services give the confidence to trust container images.

Note This test is optional.

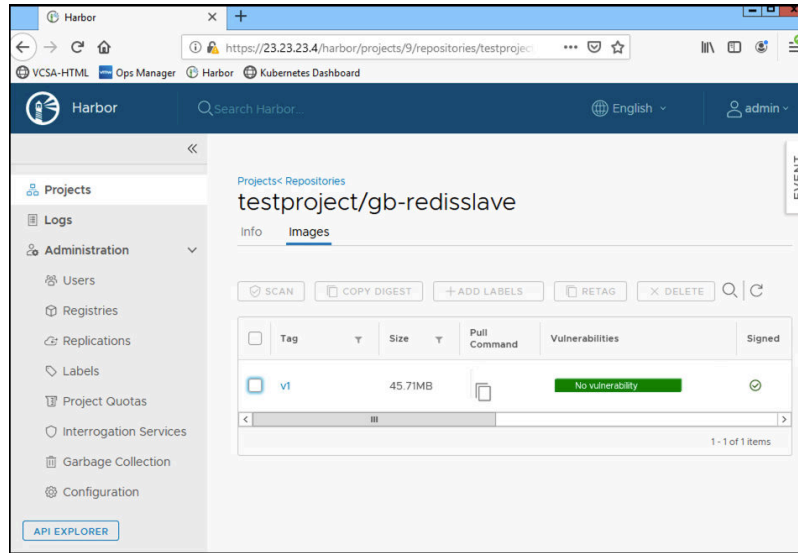
Optional

Test ID	PKS.Security.Image.Harbor.Scan
Configuration	Run the PKS.Security.Image.Sign.PushToHarbor test.

Optional

- Test Script
- 1 Select the tag from the Harbor UI. Go to **Projects > Repositories > testproject/gb-redisslave** (testproject/gb-frontend, or testproject/redis).
 - 2 Click **Scan**.
 - 3 Verify that the **Vulnerability** bar is green.

Test Output Vulnerability bar is green in the Harbor UI (see the Test script step 2).



Note If the Vulnerability bar is not green, fix vulnerability issues and rerun the test (in this case, application image has to be rebuilt and reloaded into Harbor to pass the validation).

Note Harbor uses Clair project for this vulnerability scanning feature. It automatically checks for updates from a few CVE sources. When an update in a CVE database is found, it downloads the data and uses it in the next scan.

PKS.Security.Image.Harbor.Pull

This test verifies that the signed image can be pulled from Harbor using the associated signature.

Note This test is optional.

Optional

- | | |
|---------------|--|
| Test ID | PKS.Security.Image.Harbor.Pull |
| Configuration | Run the PKS.Security.Image.Sign.PushToHarbor and PKS.Security.Image.Harbor.Scan tests. |

Optional

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Create K8s secret.


```
root@cli-vm:~# kubectl create secret docker-registry regsecret --docker-
server=harbor.corp.local --docker-username=developer --docker-password=VMware1!
--docker-email=developer@vmware.com

secret/regsecret created
```
- 3 Download YAML file from <https://github.com/kubernetes/examples/archive/master.zip>. This file is used to deploy your Application (K8S examples master.zip is used for Google sample images example).
- 4 Modify your Application YAML file locally (example: Google sample YAML) to use the image location in Harbor instead of using remote repository.

Example:

<https://github.com/kubernetes/examples/blob/master/guestbook/all-in-one/guestbook-all-in-one.yaml>

- Line 38: replace 'image: k8s.gcr.io/redis:e2e' with 'image: harbor.corp.local/testproject/redis:e2e'
- Line 82: replace 'image: gcr.io/google_samples/gb-redisslave:v1' with 'image: harbor.corp.local/testproject/gb-redisslave:v1'
- Line 133: replace 'image: gcr.io/google-samples/gb-frontend:v4' with 'image: harbor.corp.local/testproject/gb-frontend:v4'

Add `imagePullSecrets` with name `regsecret` for spec for all images to be used:

```
spec:
  imagePullSecrets:
  - name: regsecret
  containers:
  - name: master
    image: harbor.corp.local/testproject/redis:e2e
```

- 5 Copy the modified YAML file into /root on CLI-VM. See [Downloading and Transferring Files](#) for details.
- 6 Deploy your Application from previously signed Harbor images using modified YAML file (Google sample images are used as examples).

```
root@cli-vm:~# kubectl apply -f ./guestbook-all-in-one.yaml
```

```
service/redis-master created
deployment.apps/redis-master created
service/redis-slave created
deployment.apps/redis-slave created
service/frontend created
deployment.apps/frontend created
```

- 7 Verify that the Application is deployed.

```
root@cli-vm:~# kubectl get all
```

Test Output

Application is deployed successfully (see Test script step 7).

Note It takes up to two minutes to get the pods status to *Running*. At the beginning, the status is *Container creating*.

Optional

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/frontend-7776b5c947-254z6	1/1	Running	0	1m
pod/frontend-7776b5c947-9pv22	1/1	Running	0	1m
pod/frontend-7776b5c947-ccvwx	1/1	Running	0	1m
pod/redis-master-5c95c89d9-488gd	1/1	Running	0	1m
pod/redis-slave-6548b4f8b9-hpcv9	1/1	Running	0	1m
pod/redis-slave-6548b4f8b9-m4hts	1/1	Running	0	1m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/frontend	ClusterIP	10.100.200.133	<none>	80/TCP	1m
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP	3d
service/redis-master	ClusterIP	10.100.200.205	<none>	6379/TCP	1m
service/redis-slave	ClusterIP	10.100.200.29	<none>	6379/TCP	1m

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/frontend	3	3	3	3	1m
deployment.apps/redis-master	1	1	1	1	1m
deployment.apps/redis-slave	2	2	2	2	1m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/frontend-685d76557	3	3	3	1m
replicaset.apps/redis-master-6497dd4b79	1	1	1	1m
replicaset.apps/redis-slave-5d8774ddbc	2	2	2	1m

PKS.Security.Image.Harbor.Delete

This test verifies that the signed image can be deleted from Harbor.

Note This test is optional.

Optional

Test ID PKS.Security.Image.Harbor.Delete

Configuration Run the [PKS.Security.Image.Sign.PushToHarbor](#) and [PKS.Security.Image.Harbor.Scan](#) tests.

Optional

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.

- 2 Install 'notary' CLI client.

```
root@cli-vm# wget https://github.com/theupdateframework/notary/releases/download/
v0.6.0/notary-Linux-amd64
```

```
--2018-11-15 11:45:57-- https://github.com/theupdateframework/notary/releases/
download/v0.6.0/notary-Linux-amd64
Resolving github.com (github.com)... 192.30.255.113, 192.30.255.112
...
HTTP request sent, awaiting response... 200 OK
...
2018-11-15 11:45:59 (9.00 MB/s) - 'notary-Linux-amd64' saved [11109184/11109184]
```

```
root@cli-vm:~# chmod +x notary-Linux-amd64
```

```
root@cli-vm:~# mv ./notary-Linux-amd64 /usr/bin/notary
```

- 3 Remove tag from previously signed Application image (Google sample images are used as examples).

```
root@cli-vm:~# notary -s https://harbor.corp.local:4443 -d ~/.docker/trust remove
harbor.corp.local/testproject/gb-frontend v4 --tlscacert /etc/docker/certs.d/
harbor.corp.local/ca.crt
```

```
Removal of v4 from harbor.corp.local/testproject/gb-frontend staged for next
publish.
```

- 4 Verify the status of changes for the previously signed Application image.

```
root@cli-vm:~# notary -s https://harbor.corp.local:4443 -d ~/.docker/trust status
harbor.corp.local/testproject/gb-frontend --tlscacert /etc/docker/certs.d/
harbor.corp.local/ca.crt
```

- 5 Publish pending changes for the previously signed Application image.

```
root@cli-vm:~# notary -s https://harbor.corp.local:4443 -d ~/.docker/trust publish
harbor.corp.local/testproject/gb-frontend --tlscacert /etc/docker/certs.d/
harbor.corp.local/ca.crt
```

```
Pushing changes to harbor.corp.local/testproject/gb-frontend
Enter username: developer
Enter password: VMware1!
Enter passphrase for targets key with ID 7284e73: VMware1!
Successfully published changes for repository harbor.corp.local/testproject/gb-
frontend
```

- 6 Select tag (v4 in example) from the Harbor UI: **Projects > Repositories > testproject/gb-frontend** to make sure that it is shown as unsigned.

- 7 Select tag (v4 in example) and click **DELETE**. Confirm deletion by clicking **DELETE** in the pop-up dialog-box.

- 8 Delete (untag) image from Harbor using the Docker command.

```
root@cli-vm:~# docker image rm harbor.corp.local/testproject/gb-frontend:v4
```

```
Untagged: harbor.corp.local/testproject/gb-frontend:v4
Untagged: harbor.corp.local/testproject/
gb-frontend@sha256:12e31f7f4f8fa2f63c338bdf3e1b1fe04e95246ebfc5bacd7fa75125e7255a7e
```

- 9 Delete the Application image using the Docker command.

Optional

```
root@cli-vm:~# docker image rm gcr.io/google-samples/gb-frontend:v4
```

```
Untagged: gcr.io/google-samples/gb-frontend:v4
Untagged: gcr.io/google-samples/gb-frontend@sha256:
d44e7d7491a537f822e7fe8615437e4a8a08f3a7a1d7d4cb9066b6b
Deleted: sha256:e2b3e8542af735080e6bda06873ce666e2319eea353884a88e45f3c9ef996846
...
Deleted: sha256:c12ecfd4861d454a39fc17d8ef351b183425657e607def95bfa75e482d49fdce
```

- 10 Verify that the Docker image is deleted.

```
root@cli-vm:~# docker images
```

Test Output

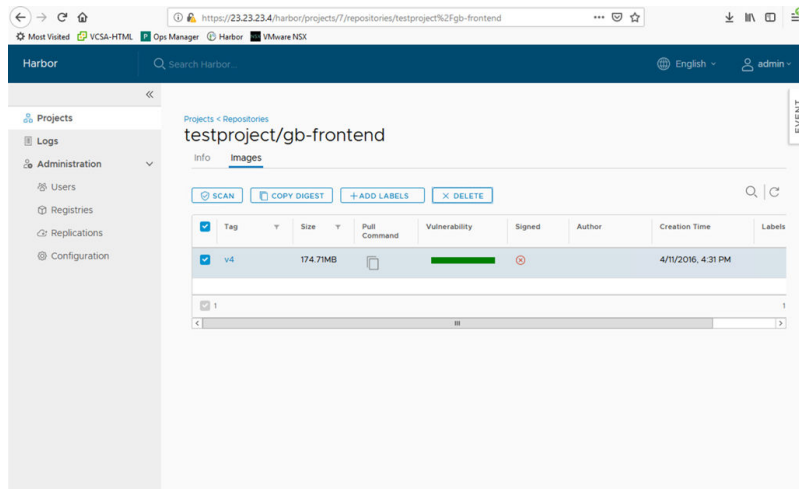
- 1 Previously signed Application image has unpublished changes (see Test script step 4).

```
root@cli-vm:~# notary -s https://harbor.corp.local:4443 -d ~/.docker/trust status
harbor.corp.local/testproject/gb-frontend --tlscacert /etc/docker/certs.d/
harbor.corp.local/ca.crt
```

Unpublished changes for harbor.corp.local/testproject/gb-frontend:

#	ACTION	SCOPE	TYPE	PATH
0	delete	targets	target	v4

- 2 Tag (v4 in example) is shown as unsigned in the Harbor UI: **Projects > Repositories > testproject/gb-frontend** (see Test script step 6).



- 3 Docker image is deleted (see Test script step 10).

```
root@cli-vm:~# docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
------------	-----	----------	---------	------

```
root@cli-vm:~#
```

PKS.Security.Image.Harbor.Modify

This test modifies a container and verifies that the modified image can be created, signed, and pushed into Harbor.

Note This test is optional.

Optional	
Test ID	PKS.Security.Image.Harbor.Modify
Configuration	Run the PKS.Security.Image.Sign.PushToHarbor and PKS.Security.Image.Harbor.Scan tests.

Optional

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Run the container from a previously signed image (Google sample images are used as examples).

```
root@cli-vm:~# docker run -i -t harbor.corp.local/testproject/gb-frontend:v4 /bin/bash
```

```
root@1e395146ecc1:/var/www/html#
```

```
root@1e395146ecc1:/var/www/html# ls -l
```

```
total 12
-rw-r--r-- 1 root root 966 Sep  9 2015 controllers.js
-rw-r--r-- 1 root root 932 Feb 20 2016 guestbook.php
-rw-r--r-- 1 root root 921 Sep  9 2015 index.html
```

- 3 Modify the container by adding a *testfile*.

```
root@1e395146ecc1:/var/www/html# touch testfile
```

```
root@1e395146ecc1:/var/www/html# ls -l
```

```
total 12
-rw-r--r-- 1 root root 966 Sep  9 2015 controllers.js
-rw-r--r-- 1 root root 932 Feb 20 2016 guestbook.php
-rw-r--r-- 1 root root 921 Sep  9 2015 index.html
-rw-r--r-- 1 root root   0 Nov 29 22:08 testfile
```

- 4 Exit from container.

```
root@1e395146ecc1:/var/www/html# exit
```

```
exit
```

- 5 Save changes to the image (use container ID *1e395146ecc1* from the interactive prompt in steps 2–4).

```
root@cli-vm:~# docker commit 1e395146ecc1 harbor.corp.local/testproject/gb-frontend:testfile
```

```
sha256:ac2dc5354fbb73099bf75f1a4ce4b3dc041d8276129ab8d8babb55372ab87298
```

- 6 Push the modified image into Harbor.

```
root@cli-vm:~# docker push harbor.corp.local/testproject/gb-frontend:testfile
```

```
The push refers to repository [harbor.corp.local/testproject/gb-frontend]
ec13c8302e63: Pushed
...
testfile: digest:
sha256:ed9cbb8f7366cfc2435bd83b36c07bfc5f113b2ee628d288b1e53d81d0cd5e35
size: 7175
```

- 7 Run the container from the modified image.

```
root@cli-vm:~# docker run -i -t harbor.corp.local/testproject/gb-frontend:testfile /bin/bash
```

- 8 Verify that the container includes *testfile* added to the modified image.

```
root@6974e8f281b9:/var/www/html# ls -l
```

- 9 Exit from container.

Optional

```
root@6974e8f281b9:/var/www/html# exit
```

```
exit
```

Test Output

Container includes the *testfile* added to the modified image (see Test script step 8).

```
root@6974e8f281b9:/var/www/html# ls -l
```

```
total 12
-rw-r--r-- 1 root root 966 Sep  9 2015 controllers.js
-rw-r--r-- 1 root root 932 Feb 20 2016 guestbook.php
-rw-r--r-- 1 root root 921 Sep  9 2015 index.html
-rw-r--r-- 1 root root   0 Nov 29 22:08 testfile
```

PKS.Security.Image.DockerBenchSecurity

This test verifies that the Application container runtime does not have any security issues.

Note It is recommended to run this test.

Recommended

Test ID PKS.Security.Image.DockerBenchSecurity

Configuration Run the [PKS.Security.Image.Sign.PushToHarbor](#) and [PKS.Security.Image.Harbor.Scan](#) tests.

Recommended

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Install Docker Bench Security script.

This script scans your Docker containers for various aspects and reports the results on which you can act upon. You can perform the following tests using the Docker bench security:

- Host configuration
- Docker daemon configuration
- Container images
- Container runtime, and so on.

You can download the Docker bench script from Git clone by running the following command:

```
root@cli-vm:~# git clone https://github.com/docker/docker-bench-security.git
```

- 3 Run the container from the Application image stored in Harbor to get the Container Runtime output in step 5.

```
root@cli-vm:~# docker run -i -t harbor.corp.local/testproject/gb-frontend:v4 /bin/bash &
```

```
[1] 10380
```

- 4 Verify that the container is running on background.

```
root@cli-vm:~# docker container ls -a
```

- 5 Start the container.

```
root@cli-vm:~# docker start 3f423882a7e1
```

```
3f423882a7e1
```

- 6 Verify that the container started.

```
root@cli-vm:~# docker container ls
```

- 7 Run the docker-bench-security script.

```
root@cli-vm:~# cd docker-bench-security/
```

```
root@cli-vm:~/docker-bench-security# ./docker-bench-security.sh
```

- 8 The docker-bench-security command checks your Application (running on Docker in the PKS environment) for various security aspects. A detailed report is displayed, which shows tests that are passed and tests that have warnings. See [Chapter 13 Appendix A: Sample Docker Bench Security Script Output](#) for details.

Total score calculation in test output:

```
[WARN] Scored -1
```

```
[PASS] Scored +1
```

```
Not score 0
```

- 9 Clean up by stopping and removing container.

```
root@cli-vm:~# docker stop 3f423882a7e1
```

```
3f423882a7e1
```

```
root@cli-vm:~# docker rm 3f423882a7e1
```

```
3f423882a7e1
```


Recommended

- 10 Check the warnings marked as [WARN]. Address security recommendations to increase the Total score.

Rerun and restart the container. Rerun docker-bench-security script and check the report.

Test Output

- 1 Container is running on background (see Test script step 4).

```
root@cli-vm:~# docker container ls -a
```

CONTAINER ID	IMAGE	STATUS	PORTS	NAMES
3f423882a7e1	harbor.corp.local/testproject/gb-frontend	Created		"/bin/bash"

- 2 Container has started (see Test script step 6).

```
root@cli-vm:~# docker container ls
```

CONTAINER ID	IMAGE	STATUS	PORTS	NAMES
3f423882a7e1	harbor.corp.local/testproject/gb-frontend	Up 9 seconds	80/tcp	"/bin/bash"

PKS.Security.Image.Setuid.Setgid

This test verifies that Setuid and Setgid binaries are removable from the container created from the Application image.

Note This test is mandatory.

Mandatory

Test ID PKS.Security.Image.Setuid.Setgid

Configuration Run the [PKS.Security.Image.Sign.PushToHarbor](#) and [PKS.Security.Image.Harbor.Scan](#) tests.

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 setuid and setgid are UNIX access rights flags that allow you to run an executable file with the permissions of the executable's owner or group respectively, and to change behavior in directories. To get a list of binaries with special permissions in the Application container image, start the container from image stored in Harbor (Google sample images are used as examples).

```
root@cli-vm:~# docker run harbor.corp.local/testproject/gb-frontend:v4 find / -perm +6000 -type f -exec ls -ld {} \; 2> /dev/null
```

Mandatory

Test Output

Collect the list of binaries with special permissions in the Application container image.

```
root@cli-vm:~# docker run harbor.corp.local/testproject/gb-frontend:v4 find / -perm +6000 -type f -exec ls -ld {} \; 2> /dev/null
```

```
-rwxr-sr-x 1 root shadow 35408 Jan 28 2016 /sbin/unix_chkpwd
-rwsr-xr-x 1 root root 39912 Nov 18 2015 /usr/bin/newgrp
...
-rwsr-xr-x 1 root root 40000 Mar 29 2015 /bin/mount
-rwsr-xr-x 1 root root 61392 Oct 28 2014 /bin/ping6
```

If the test fails, perform the following steps to unset special permission flags (in this case, application image has to be rebuilt and reloaded into Harbor to pass the validation):

- 1 Unset SUID (special permissions flags):

```
root@cli-vm:~# docker run -i -t harbor.corp.local/testproject/gb-frontend:v4 /bin/bash
```

```
root@43f4fad7578f:/var/www/html# for i in `find / -perm +6000 -type f`; do chmod a-s $i; done
```

```
find: `/proc/21/task/21/fd/5': No such file or directory
find: `/proc/21/task/21/fdinfo/5': No such file or directory
find: `/proc/21/fd/5': No such file or directory
find: `/proc/21/fdinfo/5': No such file or directory
```

- 2 Verify that SUID is not set.

```
root@43f4fad7578f:/var/www/html# find / -perm +6000 -type f -exec ls -ld {} \; 2> /dev/null
```

```
root@43f4fad7578f:/var/www/html#
```

- 3 To save changes to the image, exit from container.

```
root@43f4fad7578f:/var/www/html# exit
```

```
exit
```

- 4 Commit changes into image with the new tag.

```
root@cli-vm:~# docker commit 43f4fad7578f harbor.corp.local/testproject/gb-frontend:nosuid
```

```
sha256:ac2dc5354fbb73099bf75f1a4ce4b3dc041d8276129ab8d8babb55372ab87298
```

- 5 Push the modified image into Harbor.

```
root@cli-vm:~# docker push harbor.corp.local/testproject/gb-frontend:nosuid
```

```
The push refers to repository [harbor.corp.local/testproject/gb-frontend]
ec13c8302e63: Pushed
...
testfile: digest:
sha256:ed9cbb8f7366cfc2435bd83b36c07bfc5f113b2ee628d288b1e53d81d0cd5e35
size: 7175
```

- 6 Run the container from the modified image.

```
root@cli-vm:~# docker run -i -t harbor.corp.local/testproject/gb-frontend:nosuid /bin/bash
```

```
root@ f0154e22dbc4 :/var/www/html#
```

Mandatory

- 7 Start another PuTTY session to CLI-VM and verify that the container from the modified image is running.

```
root@cli-vm:~# docker container ls
```

CONTAINER ID	IMAGE	STATUS
COMMAND	CREATED	
PORTS	NAMES	
f0154e22dbc4	harbor.corp.local/testproject/gb-frontend	Up 9 seconds
bash"	38 seconds ago	80/tcp
suspicious_tesla		

PKS.Security.SignedImage.Cluster

This test verifies that the Application image can be used from Harbor in a cluster.

Note This test is mandatory.

Mandatory

Test ID	PKS.Security.SignedImage.Cluster
---------	----------------------------------

Configuration	Run the PKS.Security.Image.Sign.PushToHarbor and PKS.Security.Image.Harbor.Scan tests.
---------------	--

Mandatory

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Verify that the Application image is already pushed into Harbor *testproject* (Google sample images are used as examples).

```
root@cli-vm:~# docker images
```

- 3 Log in to PKS.

```
root@cli-vm:~# pks login -a uaa.corp.local -u vmware -p VMware1! -k
```

```
API Endpoint: uaa.corp.local
User: vmware
```

- 4 Create a cluster (test-cluster). See [Accessing TKGI Environment](#) for details.

```
root@cli-vm:~# pks create-cluster test-cluster --external-hostname ext-test-cluster.corp.local --plan small --num-nodes 3
```

Note The number of nodes and plan configuration might differ based on your application requirements.

You can skip this step if you have already created a PKS cluster.

- 5 Configure *test-cluster* to use Harbor.

```
root@cli-vm:~# pks cluster test-cluster
```

- Check output for Kubernetes Master IP(s).
- Edit the */etc/hosts* file and add *Kubernetes_Master_IP(s) ext-test-cluster.corp.local* line at the end of the file.

```
root@cli-vm:~# vi /etc/hosts
```

- Get credentials for the cluster - *test-cluster*.

```
root@cli-vm:~# pks get-credentials test-cluster
```

```
Fetching credentials for cluster test-cluster.
```

```
Context set for cluster test-cluster.
```

```
User can now switch between clusters by using:
$kubectl config use-context <cluster-name>
```

```
root@cli-vm:~# kubectl config use-context test-cluster
```

```
Switched to context test-cluster.
```

- 6 Create a YAML file (example: *your-app.yaml*).

```
root@cli-vm~# vi your-app.yaml
```

```
---
apiVersion: v1
kind: Pod
metadata:
  name: your-app
spec:
  containers:
  - name: your-app-container
    image: harbor.corp.local/testproject/your-app-name
    imagePullSecrets:
    - name: regsecret
```

Mandatory

- 7 Deploy the Application image from Harbor using *test-cluster*.

```
root@cli-vm~# kubectl apply -f ./your-app.yaml
```

```
deployment.apps/your-app created
```

- 8 Verify the deployment. Check that the pods with Application are up within two minutes.

```
root@cli-vm:~# kubectl get all
```

- 9 Delete the deployment.

```
root@cli-vm~# kubectl delete deployment.apps/your-app
```

```
deployment.apps "your-app" deleted
```

- 10 Verify that the Application deployment is deleted.

```
root@cli-vm~# kubectl get all
```

Test Output

- 1 Application image is pushed into Harbor *testproject* (Google sample images are used as example) - see Test script step 2.

```
root@cli-vm:~# docker images
```

REPOSITORY	SIZE	TAG	IMAGE ID	
gcr.io/google-samples/gb-frontend	512MB	v4	e2b3e8542af7	2
years ago				
harbor.corp.local/testproject/gb-frontend	512MB	v4	e2b3e8542af7	2
years ago				
gcr.io/google-samples/gb-redislave	110MB	v1	5f026ddffa27	3
years ago				
harbor.corp.local/library/gb-redislave	110MB	v1	5f026ddffa27	3
years ago				
gcr.io/google-containers/redis	419MB	e2e	e5e67996c442	3
years ago				
harbor.corp.local/library/redis	419MB	e2e	e5e67996c442	3
years ago				

- 2 Application is deployed. Pods with your application become available within two minutes (see Test script step 8).

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/your-app-865974f8d5-xt2vp	1/1	Running	0	2m
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/your-app	ClusterIP	10.100.200.141	<none>	80/TCP
2m				
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP
3d				
NAME	DESIRED	CURRENT	UP-TO-DATE	
deployment.apps/your-app	1	1	1	
1				2m
NAME	DESIRED	CURRENT	READY	
replicaset.apps/your-app-865974f8d5	1	1		

- 3 Verify that the application deployment is deleted (see Test script step 10).

Mandatory

```
root@cli-vm~# kubectl get all
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP

Application Life Cycle Management

5

This test area focuses on LCM (application install and uninstall operations), and verifies Kubernetes life cycle commands with the Application.

This chapter includes the following topics:

- [Operational Tests](#)
- [Kubernetes Life Cycle Tests](#)

Operational Tests

This group has the following test cases:

- [PKS.LCM.Image.Build](#)
- [PKS.LCM.App.Install](#)
- [PKS.LCM.App.Uninstall](#)

PKS.LCM.Image.Build

This test verifies that the Application image can be built and pushed to Harbor.

Note This test is optional.

Optional	
Test ID	PKS.LCM.Image.Build
Configuration	Follow the steps detailed in Accessing TKGI Environment .

Test ID	PKS.LCM.Image.Build
Configuration	Follow the steps detailed in Accessing TKGI Environment .

Optional

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Clone the Application source from external Git repository. For example:


```
root@cli-vm:~# git clone https://github.com/your-repository-name/your-project-name
```
- 3 Build the Docker image from the cloned Git source of application. For example: your-project-name/your-app name has Dockerfile:


```
root@cli-vm:~/your-project-name # docker build -t your-project-name/your-app-name:your_tag
```

```
Sending build context to Docker daemon 24.58MB
...
Successfully built f3a1658dd930
Successfully tagged your-project-name/your-app-name:your_tag
```
- 4 Verify that the Application image is successfully built.


```
root@cli-vm:~/your-project-name # docker images
```
- 5 Create a project - '*your-project-name*' for your Application solution in Harbor. See step 3 in [PKS.Security.Image.Sign.PushToHarbor](#).
- 6 Tag the newly built Docker image.


```
root@cli-vm:~/ your-project-name # docker tag your-project-name/your-app-name:your_tag harbor.corp.local/your-project_name/your-app-name:your_tag
```
- 7 Push the tagged Docker image into Harbor.


```
root@cli-vm:~/your-app-name # docker push harbor.corp.local/your-project_name/your-app-name:your_tag
```

```
The push refers to repository [harbor.corp.local/your-project_name/your-app-name]
aee035870ae6: Pushed
3e8a55eec1e1: Pushed
bae9908faa30: Pushed
your_tag: digest:
sha256:ba23b869f5d813480704f11e52a7a03b3bebf65648b5dbe3bac02c46f6ed5124
size: 1160
```
- 8 Open Firefox and click the **Harbor** bookmark (<https://23.23.23.4>).

Log in using credentials (user name: *admin* ; password: *VMware1!*).
- 9 Verify that the Application image is successfully pushed into Harbor.

Test Output

- 1 Application image is successfully built from Dockerfile (see Test script step 4).


```
root@cli-vm:~/your-project-name # docker images
```

REPOSITORY	SIZE	TAG	IMAGE ID	
your-project-name	/your-app-name	your_tag	f3a1658dd930	28 seconds ago
	45.8MB			
- 2 Application image is successfully pushed into Harbor (see Test script step 9).

Tag *your_tag* is shown in the Harbor UI. Go to **Projects > Repositories > your-project-name/your-app-name**.

PKS.LCM.App.Install

This test verifies that the Application image can be pulled from Harbor and installed.

Note This test is mandatory.

Mandatory	
Test ID	PKS.LCM.App.Install
Configuration	Application image is uploaded into Harbor.

Mandatory

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Log in to PKS.

```
root@cli-vm:~# pks login -a uaa.corp.local -u vmware -p VMware1! -k
```

```
API Endpoint: uaa.corp.local
User: vmware
```

- 3 Create a cluster - test-cluster (see [Accessing TKGI Environment](#) for details) or use the test-cluster that you created in [PKS.Security.SignedImage.Cluster](#).

```
root@cli-vm:~# pks create-cluster test-cluster --external-hostname ext-test-cluster.corp.local --plan small --num-nodes 3
```

Note The number of nodes and plan configuration might differ based on your application requirements.

You can skip this step if you have already created a PKS cluster.

- 4 Configure *test-cluster* to use Harbor.

```
root@cli-vm:~# pks cluster test-cluster
```

- Check output for Kubernetes Master IP(s).
- Edit the `/etc/hosts` file and add Kubernetes_Master_IP(s) `ext-test-cluster.corp.local` line at the end of the file.

```
root@cli-vm:~# vi /etc/hosts
```

- Get credentials for the cluster - *test-cluster*.

```
root@cli-vm:~# pks get-credentials test-cluster
```

```
Fetching credentials for cluster test-cluster.
```

```
Context set for cluster test-cluster.
```

```
User can now switch between clusters by using:
$kubectl config use-context <cluster-name>
```

```
root@cli-vm:~# kubectl config use-context test-cluster
```

```
Switched to context "test-cluster".
```

- 5 Deploy the Application image from Harbor using *test-cluster*.

```
root@cli-vm:~# kubectl run your-app --image=harbor.corp.local/ your-project_name/your-app-solution-name:your_tag
```

```
deployment.apps/your-app created
```

- 6 Verify the deployment. Check that the pods with Application are up within two minutes.

```
root@cli-vm:~# kubectl get all
```

Test Output

Application image is successfully deployed (see Test script step 6).

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/your-app-865974f8d5-xt2vp	1/1	Running	0	2m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/your-app	ClusterIP	10.100.200.141	<none>	80/TCP	2m

Mandatory

service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP	3d
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/your-app	1	1	1	1	2m
NAME	DESIRED	CURRENT	READY	AGE	
replicaset.apps/your-app-854d7dcd67	1	1	1	2m	

PKS.LCM.App.Uninstall

This test verifies that the Application image can be deleted.

Note This test is mandatory.

Mandatory

Test ID	PKS.LCM.App.Uninstall
Configuration	Run the PKS.LCM.App.Install test.

Mandatory

- Test Script
- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
 - 2 Verify existing Application deployment.

```
root@cli-vm:~# kubectl get all
```

- 3 Delete deployment.

```
root@cli-vm~# kubectl delete deployment.apps/your-app
```

```
deployment.apps "your-app" deleted
```

- 4 Verify that your Application deployment is deleted.

```
root@cli-vm~# kubectl get all
```

- Test Output
- 1 Existing application deployment is verified (see Test script step 2).

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/your-app-865974f8d5-xt2vp	1/1	Running	0	1h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE				
service/your-app	ClusterIP	10.100.200.141	<none>	80/TCP
1h				
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP
3d				

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE
AGE				
deployment.apps/your-app	1	1	1	1
1h				

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/your-app-854d7dcd67	1	1	1	2m

- 2 Application deployment is successfully deleted (see Test script step 4).

```
root@cli-vm~# kubectl get all
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE				
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP
3d				

Kubernetes Life Cycle Tests

This group has the following test cases:

- [K8s.Deployment.Create](#)
- [K8s.Deployment.Delete.Deploy](#)
- [K8s.Service.Deploy](#)
- [K8s.Service.Delete.Expose](#)
- [K8s.ReplicaSet.Scale](#)

■ [K8s.Pod.Delete](#)

K8s.Deployment.Create

This test verifies that the Application can be deployed with "*kind: Deployment*" specified in the YAML file.

Note This test is mandatory.

Mandatory	
Test ID	K8s.Deployment.Create
Configuration	NA

Mandatory

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Log in to PKS.

```
root@cli-vm:~# pks login -a uaa.corp.local -u vmware -p VMware1! -k
```

```
API Endpoint: uaa.corp.local
User: vmware
```

- 3 Create a cluster - *test-cluster* (see [Accessing TKGI Environment](#) for details) or use the test-cluster that you created in [PKS.Security.SignedImage.Cluster](#).

```
root@cli-vm:~# pks create-cluster test-cluster --external-hostname ext-test-cluster.corp.local --plan small --num-nodes 3
```

Note The number of nodes and plan configuration might differ based on your application requirements.

You can skip this step if you have already created a PKS cluster.

- 4 Configure *test-cluster* to use Harbor.

```
root@cli-vm:~# pks cluster test-cluster
```

- Check output for Kubernetes Master IP(s).
- Edit the `/etc/hosts` file and add `Kubernetes_Master_IP(s) ext-test-cluster.corp.local` line at the end of the file.

```
root@cli-vm:~# vi /etc/hosts
```

- Get credentials for the cluster - *test-cluster*.

```
root@cli-vm:~# pks get-credentials test-cluster
```

```
Fetching credentials for cluster test-cluster.
Context set for cluster test-cluster.
User can now switch between clusters by using:
    $kubectl config use-context <cluster-name>
```

```
root@cli-vm:~# kubectl config use-context test-cluster
```

```
Switched to context "test-cluster".
```

- 5 Clone the Application Git repository (Google sample images from K8s are used in examples).

```
root@cli-vm:~# git clone https://github.com/kubernetes/examples.git
```

```
Cloning into 'examples'...
remote: Enumerating objects: 25, done.
remote: Counting objects: 100% (25/25), done.
remote: Compressing objects: 100% (22/22), done.
remote: Total 11507 (delta 5), reused 15 (delta 3), pack-reused 11482
Receiving objects: 100% (11507/11507), 16.94 MiB | 8.21 MiB/s, done.
Resolving deltas: 100% (6131/6131), done.
Checking connectivity... done.
```

- 6 Deploy your Application using the YAML file with “*kind: Deployment*” specified.

```
root@cli-vm:~# cat examples/guestbook/frontend-deployment.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
```

Mandatory

```
spec:
  selector:
    matchLabels:
      app: guestbook
      tier: frontend
  replicas: 3
  template:
    metadata:
      labels:
        app: guestbook
        tier: frontend
    spec:
      containers:
      - name: php-redis
        image: gcr.io/google-samples/gb-frontend:v4
        resources:
          requests:
            cpu: 100m
            memory: 100Mi
        env:
        - name: GET_HOSTS_FROM
          value: dns
        ports:
        - containerPort: 80
```

```
root@cli-vm:~# kubectl apply -f examples/guestbook/frontend-deployment.yaml
```

```
deployment.apps/frontend created
```

- 7 Verify the deployment.

```
root@cli-vm:~# kubectl get all
```

Note Use 'kubectl get all --all-namespaces', if the Application uses its own namespace.

Test Output

Application is successfully deployed deployment is shown in output as available within two minutes (see Test script step 7).

```
root@cli-vm:~# kubectl get all
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/frontend	3	3	3	3	2m

K8s.Deployment.Delete.Deploy

This test verifies that the Application deployment can be deleted and redeployed.

Note This test is mandatory.

Mandatory

Test ID K8s.Deployment.Delete.Deploy

Configuration Run the [K8s.Deployment.Create](#) test.

Mandatory

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Verify that the Application deployment is running (Google sample images deployment from K8s are used as examples).

```
root@cli-vm:~# kubectl get all
```

Note Use 'kubectl get all --all-namespaces', if the Application uses its own namespace.

- 3 Delete the deployment.

```
root@cli-vm:~# kubectl delete deployment.apps/frontend
```

```
deployment.apps "frontend" deleted
```

- 4 Verify that the Application deployment is deleted.

```
root@cli-vm:~# kubectl get all
```

Note Use 'kubectl get all --all-namespaces', if the Application uses its own namespace.

- 5 Redeploy your Application using the YAML file with "*kind: Deployment*" specified (K8S examples frontend-deployment.yaml is used for Google sample images example).

```
root@cli-vm:~# cat examples/guestbook/frontend-deployment.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
spec:
  selector:
    matchLabels:
      app: guestbook
      tier: frontend
  replicas: 3
  template:
    metadata:
      labels:
        app: guestbook
        tier: frontend
    spec:
      containers:
        - name: php-redis
          image: gcr.io/google-samples/gb-frontend:v4
          resources:
            requests:
              cpu: 100m
              memory: 100Mi
          env:
            - name: GET_HOSTS_FROM
              value: dns
          ports:
            - containerPort: 80
```

```
root@cli-vm:~# kubectl apply -f examples/guestbook/frontend-deployment.yaml
```

```
deployment.apps/frontend created
```

- 6 Verify the deployment.

Mandatory

```
root@cli-vm:~# kubectl get all
```

Note Use 'kubectl get all --all-namespaces', if the Application uses its own namespace.

Test Output

- 1 Application deployment is running - *frontend* is shown in the output (see Test script step 2).

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/frontend-56f7975f44-77svf	1/1	Running	0	54m
pod/frontend-56f7975f44-9lwxw	1/1	Running	0	54m
pod/frontend-56f7975f44-qfhsc	1/1	Running	0	54m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP	3d

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/frontend	3	3	3	3	54m

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/frontend-56f7975f44	3	3	3	54m

- 2 Application deployment is deleted (see Test script step 4).

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/frontend-56f7975f44-77svf	0/1	Terminating	0	54m
pod/frontend-56f7975f44-qfhsc	0/1	Terminating	0	54m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP	3d

```
root@cli-vm:~# kubectl get all
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP	3d

- 3 Application is deployed successfully – *deployment* is shown in the output as available within two minutes (see Test script step 6).

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/frontend-56f7975f44-2tbtc	1/1	Running	0	18s
pod/frontend-56f7975f44-ft4kn	1/1	Running	0	18s
pod/frontend-56f7975f44-s8wq5	1/1	Running	0	18s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP	3d

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/frontend	3	3	3	3	19s

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/frontend-56f7975f44	3	3	3	19s

K8s.Service.Deploy

This test verifies that the Application service can be deployed with "kind: Service" specified in the YAML file.

Note This test is mandatory.

Mandatory

Test ID	K8s.Service.Deploy
---------	--------------------

Configuration	NA
---------------	----

Mandatory

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Log in to PKS.

```
root@cli-vm:~# pks login -a uaa.corp.local -u vmware -p VMware1! -k
```

```
API Endpoint: uaa.corp.local
User: vmware
```

- 3 Create a cluster - *test-cluster* (see [Accessing TKGI Environment](#) for details) or use the test-cluster that you created in [PKS.Security.SignedImage.Cluster](#).

```
root@cli-vm:~# pks create-cluster test-cluster --external-hostname ext-test-cluster.corp.local --plan small --num-nodes 3
```

Note The number of nodes and plan configuration might differ based on your application requirements.

You can skip this step if you have already created a PKS cluster.

- 4 Configure *test-cluster* to use Harbor.

```
root@cli-vm:~# pks cluster test-cluster
```

- Check output for Kubernetes Master IP(s).
- Edit the `/etc/hosts` file and add Kubernetes_Master_IP(s) `ext-test-cluster.corp.local` line at the end of the file.

```
root@cli-vm:~# vi /etc/hosts
```

- Get credentials for the cluster - *test-cluster*.

```
root@cli-vm:~# pks get-credentials test-cluster
```

```
Fetching credentials for cluster test-cluster.
Context set for cluster test-cluster.
User can now switch between clusters by using:
    kubectl config use-context <cluster-name>
```

```
root@cli-vm:~# kubectl config use-context test-cluster
```

```
Switched to context "test-cluster".
```

- 5 Clone the Application Git repository (Google sample images from K8s are used in examples).

```
root@cli-vm:~# git clone https://github.com/kubernetes/examples.git
```

```
Cloning into 'examples'...
remote: Enumerating objects: 25, done.
remote: Counting objects: 100% (25/25), done.
remote: Compressing objects: 100% (22/22), done.
remote: Total 11507 (delta 5), reused 15 (delta 3), pack-reused 11482
Receiving objects: 100% (11507/11507), 16.94 MiB | 8.21 MiB/s, done.
Resolving deltas: 100% (6131/6131), done.
Checking connectivity... done.
```

- 6 Deploy your Application using the YAML file.

```
root@cli-vm:~# cat service-frontend.yaml
```

```
apiVersion: v1
kind: Service
metadata:
  name: frontend
```

Mandatory

```
labels:
  app: guestbook
  tier: frontend
spec:
  # type: LoadBalancer
  ports:
  - port: 80
  selector:
    app: guestbook
    tier: frontend
```

```
root@cli-vm:~# kubectl apply -f ./service-frontend.yaml
```

```
service/frontend createde3c94b6f59a1
```

- 7 Verify that the service is running.

```
root@cli-vm:~# kubectl get svc
```

Note Use 'kubectl get svc --all-namespaces', if the Application uses its own namespace.

- 8 Verify that the service can be accessed. (This step depends on the application.)

Test Output

- 1 Service is running - *frontend* is shown in the output (see Test script step 7).

```
root@cli-vm:~# kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
frontend	ClusterIP	10.100.200.98	<none>	80/TCP	17s
kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP	8d

- 2 Service can be accessed. (This step depends on the application.)

K8s.Service.Delete.Expose

This test verifies that the K8s service can be deleted and redeployed.

Note This test is mandatory.

Mandatory

Test ID K8s.Service.Delete.Expose

Configuration Run the [K8s.Service.Deploy](#) test.

Mandatory

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Verify that the previously deployed service is running (Google sample images are used as examples).

```
root@cli-vm:~# kubectl get svc
```

Note Use 'kubectl get svc --all-namespaces', if the Application uses its own namespace.

- 3 Delete the service.

```
root@cli-vm:~# kubectl delete service/frontend
```

```
service "frontend" deleted
```

- 4 Verify that the service is deleted.

```
root@cli-vm:~# kubectl get svc
```

Note Use 'kubectl get svc --all-namespaces', if the Application uses its own namespace.

- 5 Expose the service using the YAML file (example: service for Google sample images).

```
root@cli-vm:~# cat service-frontend.yaml
```

```
apiVersion: v1
kind: Service
metadata:
  name: frontend
  labels:
    app: guestbook
    tier: frontend
spec:
  # type: LoadBalancer
  ports:
  - port: 80
  selector:
    app: guestbook
    tier: frontend
```

```
root@cli-vm:~# kubectl expose -f ./service-frontend.yaml
service/frontend exposed
```

- 6 Verify that the service is running.

```
root@cli-vm:~# kubectl get svc
```

Note Use 'kubectl get svc --all-namespaces', if the Application uses its own namespace.

- 7 Verify that the service can be accessed. (This step depends on the application.)

Test Output

- 1 Service is running - *frontend* is shown in the output (see Test script step 2).

```
root@cli-vm:~# kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
frontend	ClusterIP	10.100.200.6	<none>	80/TCP	7d
kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP	8d
redis-master	ClusterIP	10.100.200.101	<none>	6379/TCP	7d
redis-slave	ClusterIP	10.100.200.194	<none>	6379/TCP	7d

- 2 Service is deleted - *frontend* is not shown in the output (see Test script step 4).

Mandatory

```
root@cli-vm:~# kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP	8d
redis-master	ClusterIP	10.100.200.101	<none>	6379/TCP	7d
redis-slave	ClusterIP	10.100.200.194	<none>	6379/TCP	7d

- 3 Service is running – *frontend* is shown in the output (see Test script step 6).

```
root@cli-vm:~# kubectl get svc
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
frontend	ClusterIP	10.100.200.98	<none>	80/TCP	17s
kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP	8d
redis-master	ClusterIP	10.100.200.101	<none>	6379/TCP	7d
redis-slave	ClusterIP	10.100.200.194	<none>	6379/TCP	7d

- 4 Service can be accessed. (This step depends on the application.)

K8s.ReplicaSet.Scale

This test verifies that the Application deployment can be scaled using ReplicaSets.

Note This test is mandatory.

Mandatory

Test ID	K8s.ReplicaSet.Create.Scale
Configuration	Run the K8s.Deployment.Create test.

Mandatory

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Verify that the Application deployment is running (Google sample images deployment from K8s examples are used) .

```
root@cli-vm:~# kubectl get all
```

- 3 Scale your Application deployment.

```
root@cli-vm:~# kubectl scale --replicas=5 deployment.apps/frontend
```

```
deployment.apps/frontend scaled
```

- 4 Verify that the Application deployment is scaled.

```
root@cli-vm:~# kubectl get all
```

- 5 Rescale your Application deployment.

```
root@cli-vm:~# kubectl scale --replicas=3 deployment.apps/frontend
```

```
deployment.apps/frontend scaled
```

- 6 Verify that the Application deployment is rescaled.

```
root@cli-vm:~# kubectl get all
```

Test Output

- 1 Application deployment is running – *frontend* is shown in the output (see Test script step 2).

```
root@cli-vm:~# kubectl get all
```

```
NAME                                READY   STATUS    RESTARTS   AGE
pod/frontend-56f7975f44-7v65j      1/1     Running   0           13s
pod/frontend-56f7975f44-ktkgk      1/1     Running   0           13s
pod/frontend-56f7975f44-v9s5k      1/1     Running   0           13s

NAME                                TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
service/kubernetes                  ClusterIP     10.100.200.1  <none>       443/TCP    3d

NAME                                DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/frontend            3         3         3             3           13s

NAME                                DESIRED   CURRENT   READY        AGE
replicaset.apps/frontend-56f7975f44 3         3         3            13s
```

- 2 Application deployment is successfully scaled (see Test script step 4).

```
root@cli-vm:~# kubectl get all
```

```
NAME                                READY   STATUS    RESTARTS   AGE
pod/frontend-56f7975f44-7v65j      1/1     Running   0           48s
pod/frontend-56f7975f44-j8llt      1/1     Running   0           5s
pod/frontend-56f7975f44-ktkgk      1/1     Running   0           48s
pod/frontend-56f7975f44-m92d6      1/1     Running   0           5s
pod/frontend-56f7975f44-v9s5k      1/1     Running   0           48s
NAME                                TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
service/kubernetes                  ClusterIP     10.100.200.1  <none>       443/TCP    3d
NAME                                DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/frontend            5         5         5             5           48s
NAME                                DESIRED   CURRENT   READY        AGE
replicaset.apps/frontend-56f7975f44 5         5         5            48s
```

- 3 Application deployment is successfully rescaled (see Test script step 6).

Mandatory

```
root@cli-vm:~# kubectl get all
```

```

NAME                                READY    STATUS    RESTARTS   AGE
pod/frontend-56f7975f44-7v65j      1/1     Running   0           1m
pod/frontend-56f7975f44-j8llt      0/1     Terminating 0           33s
pod/frontend-56f7975f44-ktkgk      1/1     Running   0           1m
pod/frontend-56f7975f44-m92d6      0/1     Terminating 0           33s
pod/frontend-56f7975f44-v9s5k      1/1     Running   0           1m
NAME                                TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
service/kubernetes                  ClusterIP     10.100.200.1  <none>       443/TCP    3d
NAME                                DESIRED    CURRENT    UP-TO-DATE  AVAILABLE  AGE
deployment.apps/frontend            3          3          3           3          1m
NAME                                DESIRED    CURRENT    READY       AGE
replicaset.apps/frontend-56f7975f44 3          3          3           1m

```

K8s.Pod.Delete

This test verifies that K8s pod can be automatically recreated after deletion.

Note This test is mandatory.

Mandatory

Test ID	K8s.Pod.Delete
Configuration	Run the K8s.Deployment.Create test.

Mandatory

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Verify that the Application deployment is running (Google sample images deployment from K8s examples are used).

```
root@cli-vm:~# kubectl get all
```

- 3 Delete pod.

```
root@cli-vm:~# kubectl delete pod/frontend-56f7975f44-7v65j
```

```
pod "frontend-56f7975f44-7v65j" deleted
```

- 4 Verify that the pod is successfully recreated.

```
root@cli-vm:~# kubectl get pods
```

- 5 Delete another pod.

```
root@cli-vm:~# kubectl delete pod frontend-56f7975f44-ktkgk
```

```
pod "frontend-56f7975f44-ktkgk" deleted
```

- 6 Verify that the pod is successfully recreated. Also verify that deployments and replicaset are available, and ready.

```
root@cli-vm:~# kubectl get all
```

Test Output

- 1 Application deployment is running – *frontend* is shown in the output (see Test script step 2).

```
root@cli-vm:~# kubectl get all
```

```
NAME                                READY    STATUS    RESTARTS   AGE
pod/frontend-56f7975f44-7v65j      1/1      Running   0           42m
pod/frontend-56f7975f44-ktkgk      1/1      Running   0           42m
pod/frontend-56f7975f44-v9s5k      1/1      Running   0           42m
NAME                                TYPE     CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
service/kubernetes                  ClusterIP  10.100.200.1  <none>        443/TCP    3d
NAME                                DESIRED  CURRENT  UP-TO-DATE  AVAILABLE  AGE
deployment.apps/frontend            3        3        3            3          42m
NAME                                DESIRED  CURRENT  READY        AGE
replicaset.apps/frontend-56f7975f44 3        3        3           42m
```

- 2 The pod is successfully recreated after deletion (see Test script step 4).

```
root@cli-vm:~# kubectl get pods
```

```
NAME                                READY    STATUS    RESTARTS   AGE
frontend-56f7975f44-bxspz          1/1      Running   0           54s
frontend-56f7975f44-ktkgk          1/1      Running   0           44m
frontend-56f7975f44-v9s5k          1/1      Running   0           44m
```

- 3 The pod is successfully recreated. Deployments and replicaset are available, and ready (see Test script step 6).

```
root@cli-vm:~# kubectl get all
```

```
NAME                                READY    STATUS    RESTARTS   AGE
pod/frontend-56f7975f44-bxspz      1/1      Running   0           1m
pod/frontend-56f7975f44-dwkvf      1/1      Running   0           16s
pod/frontend-56f7975f44-v9s5k      1/1      Running   0           45m
NAME                                TYPE     CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
service/kubernetes                  ClusterIP  10.100.200.1  <none>        443/TCP    3d
```

Mandatory

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/frontend	3	3	3	3	45m
NAME	DESIRED	CURRENT	READY	AGE	
replicaset.apps/frontend-56f7975f44	3	3	3	45m	

Partner Validation

6

This test area focuses on partner validation.

This chapter includes the following topics:

- [Functionality Tests](#)

Functionality Tests

The Functionality Tests group has the following test case:

- [App.Functionality \(Partner Validation\)](#)

App.Functionality (Partner Validation)

This test ensures the container functionality in the TKGI environment.

Note This test is mandatory.

Mandatory	
Test ID	App.Functionality (Partner Validation)
Configuration	PKS cluster must be created.

Mandatory

Test Script

After deployment, the Application is powered on and configured. Perform an initial sanity test by logging in to the solution and validating its functionality. Ensure that the IP address and host names are configured appropriately. Before running the VMware Validation Tests, you must run a set of tests that are specific to your Application. This ensures that your solution operates normally before running the VMware Validation Tests.

Partner Validation Test Scenario Examples:

Example Scenario 1:

Deploy a multi-tier application with a frontend, a database, and a load balancer. For details, see [Load Balancers](#).

- When the deployment "`app-name/deployment.yaml`" is applied:
 - Current context is set to the cluster "`test-cluster`" and namespace "`app-name`".
 - Deployment "`app-name`" is rolled out within two minutes.
 - Application is reachable at "`http://24.24.24.x`" through the load-balancer "`app-name-loadbalancer`".
- When the deployment "`app-name/deployment.yaml`" is deleted, the namespace "`app-name`" does not exist.

Example Scenario 2:

Deleting and recreating workload.

- When the deployment "`app-name/deployment.yaml`" is applied:
 - Current context is set to the cluster "`test-cluster`" and namespace "`app-name`".
 - Deployment "`app-name`" is resized to five pods.
 - Deployment "`app-name`" is rolled out within two minutes.
 - Application is reachable at "`http://24.24.24.x`" through the load-balancer "`app-name-loadbalancer`".
- When the deployment "`app-name/deployment.yaml`" is deleted, the namespace "`app-name`" does not exist.
- When the deployment "`app-name/deployment.yaml`" is applied:
 - Current context is set to the cluster "`test-cluster`" and namespace "`app-name`".
 - Deployment "`app-name`" is resized to five pods.
 - Deployment "`app-name`" is rolled out within two minutes.
 - Application is reachable at "`http://24.24.24.x`" through the load-balancer "`app-name-loadbalancer`".
- When the deployment "`app-name/deployment.yaml`" is deleted, the namespace "`app-name`" does not exist.

Test Output

Application operates normally before running the VMware Validation Tests.

Note Ensure that the Application solution is deployed and running in a good state after the validation tests are done. VMware might need to verify the Application solution deployment during the validation results review.

TKGI Centric Platform Validation

7

This test area focuses on the cluster K8s operations when the Application is running.

This chapter includes the following topics:

- [Resilience/Negative Tests](#)

Resilience/Negative Tests

The Resilience/Negative Tests group has the following test cases:

- [PKS.Resilience.ClusterExpand](#)
- [PKS.Worker.Reboot](#)
- [PKS.Master.Reboot](#)
- [PKS.Worker.PowerOffOn](#)
- [PKS.Master.PowerOffOn](#)

PKS.Resilience.ClusterExpand

This test verifies that the Application can be expanded with the growing K8s resources.

Note This test is optional.

Optional	
Test ID	PKS.Resilience.ClusterExpand
Configuration	PKS cluster is created. Application is running.

Optional

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Check the cluster existence.

```
root@cli-vm:~# pks clusters
```

Name	Plan Name	UUID	Status
k8s-cluster-101	small	53187740-e5cf-4ed7-9181-ab71bbdc1945	succeeded
CREATE			

- 3 Get credentials for the cluster.

```
root@cli-vm:~# pks get-credentials k8s-cluster-101
```

```
Fetching credentials for cluster k8s-cluster-101.
Context set for cluster k8s-cluster-101.
```

You can now switch between clusters by using:

```
$kubectl config use-context <cluster-name>
```

- 4 Set kubectl to use the context of the cluster.

```
root@cli-vm:~# kubectl config use-context k8s-cluster-101
```

```
Switched to context "k8s-cluster-101".
```

- 5 Check the cluster nodes status.

```
root@cli-vm:~# kubectl get nodes -o wide
```

NAME	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE	STATUS	ROLES	AGE	VERSION
CONTAINER-RUNTIME							
afc7bfbd-bea7-4044-8d81-2ed6121eec78	172.23.0.4	172.23.0.4	Ubuntu 16.04.5 LTS	Ready	<none>	25d	v1.11.2
cd48fdbbc-5f15-437c-9a1b-f70235c59a62	172.23.0.3	172.23.0.3	Ubuntu 16.04.5 LTS	Ready	<none>	25d	v1.11.2
dc8ca134-93bb-422a-89d0-89a2ba359685	172.23.0.5	172.23.0.5	Ubuntu 16.04.5 LTS	Ready	<none>	25d	v1.11.2

- 6 Resize the cluster by increasing the number of worker nodes ('--num-nodes 5' is only an example, it depends on the initial number of worker nodes in the cluster).

```
root@cli-vm:~# pks resize k8s-cluster-101 --num-nodes 5
```

```
Are you sure you want to resize cluster k8s-cluster-101 to 5? (y/n): y
```

Use 'pks cluster k8s-cluster-101' to monitor the state of your cluster.

```
root@cli-vm:~# pks cluster k8s-cluster-101
```

```
Name: k8s-cluster-101
Plan Name: small
UUID: 53187740-e5cf-4ed7-9181-ab71bbdc1945
Last Action: UPDATE
Last Action State: in progress
Last Action Description: Instance update in progress
Kubernetes Master Host: pks-cluster-101.corp.local
```

Optional

```
Kubernetes Master Port: 8443
Worker Nodes: 5
Kubernetes Master IP(s): 24.24.24.10
Network Profile Name:
```

- 7 Monitor cluster resizing process by running this command:

```
root@cli-vm:~# pks cluster k8s-cluster-101
```

- 8 Verify that the cluster is successfully resized and number of workers is increased.

```
root@cli-vm:~# pks clusters
```

```
root@cli-vm:~# pks cluster k8s-cluster-101
```

```
root@cli-vm:~# kubectl get nodes -o wide
```

Test Output

Cluster is successfully resized and number of workers is increased (see Test script step 8).

```
root@cli-vm:~# pks clusters
```

Name	Plan Name	UUID	Status
k8s-cluster-101	small	53187740-e5cf-4ed7-9181-ab71bbdc1945	succeeded
UPDATE			

```
root@cli-vm:~# pks cluster k8s-cluster-101
```

```
Name: k8s-cluster-101
Plan Name: small
UUID: 53187740-e5cf-4ed7-9181-ab71bbdc1945
Last Action: UPDATE
Last Action State: succeeded
Last Action Description: Instance update completed
Kubernetes Master Host: pks-cluster-101.corp.local
Kubernetes Master Port: 8443
Worker Nodes: 5
Kubernetes Master IP(s): 24.24.24.10
Network Profile Name:
```

```
root@cli-vm:~# kubectl get nodes -o wide
```

NAME	INTERNAL-IP	EXTERNAL-IP	OS-IMAGE	STATUS	ROLES	AGE	VERSION
					KERNEL-VERSION		CONTAINER-
							RUNTIME
40771b7c-9728-4c35-ba88-5407295b45ab	172.23.0.8	172.23.0.8	Ubuntu 16.04.5 LTS	Ready	<none>	16m	v1.11.2
17.12.1-ce					4.15.0-33-generic		docker://
afc7bfb-d-bea7-4044-8d81-2ed6121eec78	172.23.0.4	172.23.0.4	Ubuntu 16.04.5 LTS	Ready	<none>	26d	v1.11.2
17.12.1-ce					4.15.0-33-generic		docker://
cd48fdb-c-5f15-437c-9a1b-f70235c59a62	172.23.0.3	172.23.0.3	Ubuntu 16.04.5 LTS	Ready	<none>	26d	v1.11.2
17.12.1-ce					4.15.0-33-generic		docker://
dc8ca134-93bb-422a-89d0-89a2ba359685	172.23.0.5	172.23.0.5	Ubuntu 16.04.5 LTS	Ready	<none>	26d	v1.11.2
17.12.1-ce					4.15.0-33-generic		docker://
f1b26c6e-3ae8-4e46-aae6-ec174f3670b0	172.23.0.7	172.23.0.7	Ubuntu 16.04.5 LTS	Ready	<none>	23m	v1.11.2
17.12.1-ce					4.15.0-33-generic		docker://

PKS.Worker.Reboot

This test verifies that the running K8s cluster worker node can be rebooted when the Stateless application is running.

Note This test is mandatory for Stateless applications only.

Mandatory for Stateless Applications Only	
Test ID	PKS.Worker.Reboot
Configuration	PKS cluster is created. Application is running.

Mandatory for Stateless Applications Only

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Check the cluster existence.

```
root@cli-vm:~# pks clusters
```

Name	Plan Name	UUID	Status
k8s-cluster-101	small	53187740-e5cf-4ed7-9181-ab71bbdc1945	succeeded

- 3 Get credentials for the cluster.

```
root@cli-vm:~# pks get-credentials k8s-cluster-101
```

```
Fetching credentials for cluster k8s-cluster-101.
Context set for cluster k8s-cluster-101.
```

You can now switch between clusters by using:

```
$kubectl config use-context <cluster-name>
```

- 4 Set kubectl to use the context of the cluster.

```
root@cli-vm:~# kubectl config use-context k8s-cluster-101
```

```
Switched to context "k8s-cluster-101".
```

- 5 Check the cluster nodes status (cluster with six worker nodes for Google sample images is used as application example).

```
root@cli-vm:~# kubectl get nodes -o wide
```

NAME	VERSION	INTERNAL-IP	EXTERNAL-IP	STATUS	ROLES	AGE
0cc55cd5-ba68-4647-aa30-f0974ee218a5	v1.11.6	172.23.2.7	172.23.2.7	Ready	<none>	2d
generic	docker://17.12.1-ce			Ubuntu	16.04.5 LTS	4.15.0-42-
89f41686-aa51-42b8-8555-7f098e7d7d3e	v1.11.6	172.23.2.8	172.23.2.8	Ready	<none>	2d
generic	docker://17.12.1-ce			Ubuntu	16.04.5 LTS	4.15.0-42-
8bd0e317-fde8-4455-9e2a-d11144bce7db	v1.11.6	172.23.2.6	172.23.2.6	Ready	<none>	2d
generic	docker://17.12.1-ce			Ubuntu	16.04.5 LTS	4.15.0-42-
97a40e95-d49d-420a-8ff2-1f15fb02b247	v1.11.6	172.23.2.5	172.23.2.5	Ready	<none>	2d
generic	docker://17.12.1-ce			Ubuntu	16.04.5 LTS	4.15.0-42-
c3b10035-90e0-4077-a841-38c0db885bdb	v1.11.6	172.23.2.4	172.23.2.4	Ready	<none>	2d
generic	docker://17.12.1-ce			Ubuntu	16.04.5 LTS	4.15.0-42-
e5e041ed-ef25-4ea7-8dc6-9abad079e8f9	v1.11.6	172.23.2.3	172.23.2.3	Ready	<none>	2d
generic	docker://17.12.1-ce			Ubuntu	16.04.5 LTS	4.15.0-42-

- 6 Check pods status.

Mandatory for Stateless Applications Only

```
root@cli-vm:~# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
frontend-7776b5c947-7clrm	1/1	Running	0	6m
frontend-7776b5c947-96jlh	1/1	Running	0	6m
frontend-7776b5c947-xsrv7	1/1	Running	0	6m
redis-master-5c95c89d9-78ds4	1/1	Running	0	6m
redis-slave-6548b4f8b9-5p92d	1/1	Running	0	6m
redis-slave-6548b4f8b9-x9fd9	1/1	Running	0	6m

- 7 Get the pod IP for worker node under reboot test.

```
root@cli-vm:~# kubectl describe pod frontend-7776b5c947-7clrm | grep IP
```

```
IP: 172.16.7.7
```

- 8 Get the VM name for the worker node under the reboot test.

```
root@cli-vm:~# source A-BOSH.env
```

```
root@cli-vm:~# bosh vms
```

```
...
Deployment 'service-instance_a2a02b92-0ac5-4bfd-9706-fcfeae528d2b'

Instance                               Process State AZ
IPs      VM CID                      VM Type    Active
master/4a7a6e3c-fcf3-4fa7-950b-7031c68f67ba running    AZ-COMP
172.23.2.2 vm-04c898ea-6614-4b08-a22a-c50bed645293 medium.disk true
worker/32ea5716-26fe-47cd-966d-b915af7c854f running    AZ-COMP
172.23.2.4 vm-4506fbdd-b021-46fe-95ba-eb8197c22482 medium.disk true
worker/36d21afb-ec09-4c14-8a9c-7e43aac9b375 running    AZ-COMP
172.23.2.5 vm-ad72851e-df0c-4184-a6fe-53d7743242c2 medium.disk true
worker/9835f1fc-087e-4df9-af50-ff535e42c0c2 running    AZ-COMP
172.23.2.3 vm-edc446f3-3827-424e-b9b7-3136f3af75a4 medium.disk true
worker/9face359-f852-427f-9faa-649e3f7b7cb5 running    AZ-COMP
172.23.2.6 vm-79f67e08-72c1-46aa-9032-03fd47a0ff2f medium.disk true
worker/e218da9d-41d6-4098-a560-56669415b66e starting    AZ-COMP
172.23.2.7 vm-843f9e36-dcee-4b69-929b-a5355e3bec74 medium.disk true
worker/f10418ab-3362-4a7d-9968-a6a757a45c78 running    AZ-COMP
172.23.2.8 vm-9163ea29-97c1-494c-b796-fc8a35ef8241 medium.disk true
```

Pods IP block has CIDR 172.16.0.0/16 defined by NSX-T setup.

Nodes IP block has CIDR 172.23.0.0/16 defined by NSX-T setup.

Pod with IP 172.16.7.7 corresponds to node with IP 172.23.2.7.

vm-843f9e36-dcee-4b69-929b-a5355e3bec74 is VM under test.

- 9 Access vSphere Client. See [Accessing vSphere Client](#) for details.
- 10 Highlight vm-843f9e36-dcee-4b69-929b-a5355e3bec74 under RegionA01-COMP01/PKS-COMP cluster and select **Power** > **Reset** from the context menu.
- 11 Go back to the PuTTY session to CLI-VM and monitor the rebooted node status by running the bosh vms command. Wait until **Process State** is shown in green as **Running** (it may take up to five minutes).
- 12 Verify the Application deployment readiness.

```
root@cli-vm:~# kubectl get all
```

Mandatory for Stateless Applications Only

13 Repeat steps 6–12 for each worker node.

Test Output

- 1 Worker node is successfully rebooted and has **Process State** as **Running** (see Test script step 11). The Process State can be *unresponsive agent*, *starting*, or *failing* during reboot.

```
root@cli-vm:~# bosh vms
```

Instance IPs	VM CID	Process State VM Type	AZ Active
worker/e218da9d-41d6-4098-a560-56669415b66e 172.23.2.7	vm-843f9e36-dcee-4b69-929b-a5355e3bec74	'unresponsive agent' medium.disk	AZ-COMP true

```
root@cli-vm:~# bosh vms
```

Instance IPs	VM CID	Process State VM Type	AZ Active
worker/e218da9d-41d6-4098-a560-56669415b66e 172.23.2.7	vm-843f9e36-dcee-4b69-929b-a5355e3bec74	starting medium.disk	AZ-COMP true

```
root@cli-vm:~# bosh vms
```

Instance IPs	VM CID	Process State VM Type	AZ Active
worker/e218da9d-41d6-4098-a560-56669415b66e 172.23.2.7	vm-843f9e36-dcee-4b69-929b-a5355e3bec74	failing medium.disk	AZ-COMP true

```
root@cli-vm:~# bosh vms
```

Instance IPs	Process State VM Type	AZ Active
worker/e218da9d-41d6-4098-a560-56669415b66e 172.23.2.7	running medium.disk	AZ-COMP true

- 2 Application deployment is running and ready.

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/frontend-7776b5c947-7clrm	1/1	Running	1	12m
pod/frontend-7776b5c947-96jlh	1/1	Running	0	12m
pod/frontend-7776b5c947-xsrv7	1/1	Running	0	12m
pod/redis-master-5c95c89d9-78ds4	1/1	Running	0	12m
pod/redis-slave-6548b4f8b9-5p92d	1/1	Running	0	12m
pod/redis-slave-6548b4f8b9-x9fd9	1/1	Running	0	12m

NAME AGE	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
service/frontend 12m	ClusterIP	10.100.200.152	<none>	80/TCP
service/kubernetes 2d	ClusterIP	10.100.200.1	<none>	443/TCP
service/redis-master 12m	ClusterIP	10.100.200.55	<none>	6379/TCP
service/redis-slave 12m	ClusterIP	10.100.200.194	<none>	6379/TCP

NAME AGE	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE
deployment.apps/frontend 12m	3	3	3	3

Mandatory for Stateless Applications Only

deployment.apps/redis-master	1	1	1	1
12m				
deployment.apps/redis-slave	2	2	2	2
12m				

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/frontend-7776b5c947	3	3	3	12m
replicaset.apps/redis-master-5c95c89d9	1	1	1	12m
replicaset.apps/redis-slave-6548b4f8b9	2	2	2	12m

- 3 Every worker node is successfully rebooted and has **Process State** as **Running** (see Test script step 13).

PKS.Master.Reboot

This test verifies that the running K8s cluster master node can be rebooted when the Stateless application is running.

Note This test is mandatory for Stateless applications only.

Mandatory for Stateless Applications Only

Test ID	PKS.Master.Reboot
Configuration	PKS cluster is created. Application is running.

Mandatory for Stateless Applications Only

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Check the cluster existence.

```
root@cli-vm:~# pks clusters
```

Name	Plan Name	UUID	Status
Action k8s-cluster-101 CREATE	small	53187740-e5cf-4ed7-9181-ab71bbdc1945	succeeded

- 3 Get credentials for the cluster.

```
root@cli-vm:~# pks get-credentials k8s-cluster-101
```

```
Fetching credentials for cluster k8s-cluster-101.  
Context set for cluster k8s-cluster-101.
```

You can now switch between clusters by using:

```
$kubectl config use-context <cluster-name>
```

- 4 Set kubectl to use the context of the cluster.

```
root@cli-vm:~# kubectl config use-context k8s-cluster-101
```

```
Switched to context "k8s-cluster-101".
```

- 5 Get the VM name for the master node under reboot test.

```
root@cli-vm:~# source A-BOSH.env
```

```
root@cli-vm:~# bosh vms
```

```
Deployment 'service-instance_a2a02b92-0ac5-4bfd-9706-fcfeae528d2b'
```

Instance	Process State	AZ
IPs VM CID	VM Type	Active
master/4a7a6e3c-fcf3-4fa7-950b-7031c68f67ba	running	AZ-COMP
172.23.2.2 vm-04c898ea-6614-4b08-a22a-c50bed645293	medium.disk	true

- 6 Access vSphere Client. See [Accessing vSphere Client](#) for details.
- 7 Highlight vm-04c898ea-6614-4b08-a22a-c50bed645293 under RegionA01-COMP01/PKS-COMP cluster and select **Power > Reset** from the context menu.
- 8 Go back to the PuTTY session to CLI-VM and monitor the rebooted master node status by running the bosh vms command. Wait until **Process State** is shown in green as **Running** (it might take up to five minutes).
- 9 Verify the Application deployment readiness.

```
root@cli-vm:~# kubectl get all
```

Test Output

- 1 Master node is successfully rebooted and has **Process State** as **Running** (see Test script step 8). The Process State can be *unresponsive agent*, *starting*, or *failing* during reboot.

```
root@cli-vm:~# bosh vms
```

Instance	Process State	AZ
IPs VM CID	VM Type	Active
master/4a7a6e3c-fcf3-4fa7-950b-7031c68f67ba	'unresponsive agent'	AZ-COMP
172.23.2.2 vm-04c898ea-6614-4b08-a22a-c50bed645293	medium.disk	true

Mandatory for Stateless Applications Only

```
root@cli-vm:~# bosh vms
```

Instance IPs	VM CID	Process	State	AZ	VM Type	Active
master/4a7a6e3c-fcf3-4fa7-950b-7031c68f67ba		starting		AZ-COMP		
172.23.2.2	vm-04c898ea-6614-4b08-a22a-c50bed645293				medium.disk	true

```
root@cli-vm:~# bosh vms
```

Instance IPs	VM CID	Process	State	AZ	VM Type	Active
master/4a7a6e3c-fcf3-4fa7-950b-7031c68f67ba		failing		AZ-COMP		
172.23.2.2	vm-04c898ea-6614-4b08-a22a-c50bed645293				medium.disk	true

```
root@cli-vm:~# bosh vms
```

master/4a7a6e3c-fcf3-4fa7-950b-7031c68f67ba	running	AZ-COMP		
172.23.2.2	vm-04c898ea-6614-4b08-a22a-c50bed645293	medium.disk	true	

- 2 Application deployment is running and ready (see Test script step 9).

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/frontend-7776b5c947-7clrm	1/1	Running	1	12m
pod/frontend-7776b5c947-96jlh	1/1	Running	0	12m
pod/frontend-7776b5c947-xsrv7	1/1	Running	0	12m
pod/redis-master-5c95c89d9-78ds4	1/1	Running	0	12m
pod/redis-slave-6548b4f8b9-5p92d	1/1	Running	0	12m
pod/redis-slave-6548b4f8b9-x9fd9	1/1	Running	0	12m

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE				
service/frontend	ClusterIP	10.100.200.152	<none>	80/TCP
12m				
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP
2d				
service/redis-master	ClusterIP	10.100.200.55	<none>	6379/TCP
12m				
service/redis-slave	ClusterIP	10.100.200.194	<none>	6379/TCP
12m				

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE
AGE				
deployment.apps/frontend	3	3	3	3
12m				
deployment.apps/redis-master	1	1	1	1
12m				
deployment.apps/redis-slave	2	2	2	2
12m				

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/frontend-7776b5c947	3	3	3	12m
replicaset.apps/redis-master-5c95c89d9	1	1	1	12m
replicaset.apps/redis-slave-6548b4f8b9	2	2	2	12m

PKS.Worker.PowerOffOn

This test verifies that running K8s cluster worker node can rejoin the cluster after it is powered off/on, when the Stateless application is running.

Note This test is mandatory for Stateless applications only.

Mandatory for Stateless Applications Only

Test ID	PKS.Worker.PowerOffOn
Configuration	PKS cluster is created. Application is running.

Mandatory for Stateless Applications Only

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Check the cluster existence.

```
root@cli-vm:~# pks clusters
```

Name	Plan Name	UUID	Status
k8s-cluster-101	small	53187740-e5cf-4ed7-9181-ab71bbdc1945	succeeded

- 3 Get credentials for the cluster.

```
root@cli-vm:~# pks get-credentials k8s-cluster-101
```

```
Fetching credentials for cluster k8s-cluster-101.
Context set for cluster k8s-cluster-101.
```

You can now switch between clusters by using:

```
$kubectl config use-context <cluster-name>
```

- 4 Set kubectl to use the context of the cluster.

```
root@cli-vm:~# kubectl config use-context k8s-cluster-101
```

```
Switched to context "k8s-cluster-101".
```

- 5 Verify the Application deployment readiness.

```
root@cli-vm:~# kubectl get all
```

- 6 Check the pod status.

```
root@cli-vm:~# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
frontend-7776b5c947-7clrm	1/1	Running	0	6m
frontend-7776b5c947-96jlh	1/1	Running	0	6m
frontend-7776b5c947-xsrv7	1/1	Running	0	6m
redis-master-5c95c89d9-78ds4	1/1	Running	0	6m
redis-slave-6548b4f8b9-5p92d	1/1	Running	0	6m
redis-slave-6548b4f8b9-x9fd9	1/1	Running	0	6m

- 7 Get the pod IP for worker node under the power off/on test.

```
root@cli-vm:~# kubectl describe pod frontend-7776b5c947-7clrm | grep IP
```

```
IP: 172.16.7.7
```

- 8 Get the VM name for the worker node under the power off/on test.

```
root@cli-vm:~# source A-BOSH.env
```

```
root@cli-vm:~# bosh vms
```

```
...
Deployment 'service-instance_a2a02b92-0ac5-4bfd-9706-fcfeae528d2b'

Instance                               Process State AZ
IPs      VM CID                      VM Type      Active
master/4a7a6e3c-fcf3-4fa7-950b-7031c68f67ba running      AZ-COMP
172.23.2.2 vm-04c898ea-6614-4b08-a22a-c50bed645293 medium.disk true
worker/32ea5716-26fe-47cd-966d-b915af7c854f running      AZ-COMP
```


Mandatory for Stateless Applications Only

```

172.23.2.4  vm-4506fbdd-b021-46fe-95ba-eb8197c22482  medium.disk  true
worker/36d21afb-ec09-4c14-8a9c-7e43aac9b375  running      AZ-COMP
172.23.2.5  vm-ad72851e-df0c-4184-a6fe-53d7743242c2  medium.disk  true
worker/9835f1fc-087e-4df9-af50-ff535e42c0c2  running      AZ-COMP
172.23.2.3  vm-edc446f3-3827-424e-b9b7-3136f3af75a4  medium.disk  true
worker/9face359-f852-427f-9faa-649e3f7b7cb5  running      AZ-COMP
172.23.2.6  vm-79f67e08-72c1-46aa-9032-03fd47a0ff2f  medium.disk  true
worker/e218da9d-41d6-4098-a560-56669415b66e  starting     AZ-COMP
172.23.2.7  vm-843f9e36-dcee-4b69-929b-a5355e3bec74  medium.disk  true
worker/f10418ab-3362-4a7d-9968-a6a757a45c78  running      AZ-COMP
172.23.2.8  vm-9163ea29-97c1-494c-b796-fc8a35ef8241  medium.disk  true

```

Pods IP block has CIDR 172.16.0.0/16 defined by NSX-T setup.

Nodes IP block has CIDR 172.23.0.0/16 defined by NSX-T setup.

Pod with IP 172.16.7.7 corresponds to node with IP 172.23.2.7.

vm-843f9e36-dcee-4b69-929b-a5355e3bec74 is VM under test.

- 9 Access vSphere Client. See [Accessing vSphere Client](#) for details.
- 10 Highlight vm-843f9e36-dcee-4b69-929b-a5355e3bec74 under RegionA01-COMP01/PKS-COMP cluster and select **Power** > **Power Off** from the context menu.
- 11 Go back to the PuTTY session to CLI-VM and verify the Application deployment readiness.

root@cli-vm:~# kubectl get all
- 12 Access vSphere Client.
- 13 Highlight vm-843f9e36-dcee-4b69-929b-a5355e3bec74 under RegionA01-COMP01/PKS-COMP cluster and select **Power** > **Power On** from the context menu.
- 14 Go back to the PuTTY session to CLI-VM and monitor the powered on worker node status by running the kubectl get all command. Wait until the worker node is shown as **Available** (it might take up to five minutes).

Test Output

- 1 Application deployment is running and ready (see Test script step 5).

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/frontend-7776b5c947-7clrm	1/1	Running	1	1h
pod/frontend-7776b5c947-96jlh	1/1	Running	0	1h
pod/frontend-7776b5c947-xsrv7	1/1	Running	0	1h
pod/redis-master-5c95c89d9-78ds4	1/1	Running	0	1h
pod/redis-slave-6548b4f8b9-5p92d	1/1	Running	0	1h
pod/redis-slave-6548b4f8b9-x9fd9	1/1	Running	0	1h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE				
service/frontend	ClusterIP	10.100.200.152	<none>	80/TCP
1h				
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP
2d				
service/redis-master	ClusterIP	10.100.200.55	<none>	6379/TCP
1h				
service/redis-slave	ClusterIP	10.100.200.194	<none>	6379/TCP
1h				

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE
AGE				
deployment.apps/frontend	3	3	3	3
1h				
deployment.apps/redis-master	1	1	1	1

Mandatory for Stateless Applications Only

1h				
deployment.apps/redis-slave	2	2	2	2
1h				
NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/frontend-7776b5c947	3	3	3	1h
replicaset.apps/redis-master-5c95c89d9	1	1	1	1h
replicaset.apps/redis-slave-6548b4f8b9	2	2	2	1h

- 2 Application deployment is running and ready except for the availability of the powered off node (see Test script step 11).

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/frontend-7776b5c947-7clrm	1/1	Running	1	1h
pod/frontend-7776b5c947-96jlh	1/1	Running	0	1h
pod/frontend-7776b5c947-xsrv7	1/1	Running	0	1h
pod/redis-master-5c95c89d9-78ds4	1/1	Running	0	1h
pod/redis-slave-6548b4f8b9-5p92d	1/1	Running	0	1h
pod/redis-slave-6548b4f8b9-x9fd9	1/1	Running	0	1h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE				
service/frontend	ClusterIP	10.100.200.152	<none>	80/TCP
1h				
service/kubernetes	ClusterIP	10.100.200.1	<none>	443/TCP
2d				
service/redis-master	ClusterIP	10.100.200.55	<none>	6379/TCP
1h				
service/redis-slave	ClusterIP	10.100.200.194	<none>	6379/TCP
1h				

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE
AGE				
deployment.apps/frontend	3	3	3	2
1h				
deployment.apps/redis-master	1	1	1	1
1h				
deployment.apps/redis-slave	2	2	2	2
1h				

NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/frontend-7776b5c947	3	3	2	1h
replicaset.apps/redis-master-5c95c89d9	1	1	1	1h
replicaset.apps/redis-slave-6548b4f8b9	2	2	2	1h

- 3 Application deployment is running and ready for all nodes (see Test script step 14).

```
root@cli-vm:~# kubectl get all
```

NAME	READY	STATUS	RESTARTS	AGE
pod/frontend-7776b5c947-7clrm	1/1	Running	2	1h
pod/frontend-7776b5c947-96jlh	1/1	Running	0	1h
pod/frontend-7776b5c947-xsrv7	1/1	Running	0	1h
pod/redis-master-5c95c89d9-78ds4	1/1	Running	0	1h
pod/redis-slave-6548b4f8b9-5p92d	1/1	Running	0	1h
pod/redis-slave-6548b4f8b9-x9fd9	1/1	Running	0	1h

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE				
service/frontend	ClusterIP	10.100.200.152	<none>	80/TCP
1h				

Mandatory for Stateless Applications Only

service/kubernetes 2d	ClusterIP	10.100.200.1	<none>	443/TCP
service/redis-master 1h	ClusterIP	10.100.200.55	<none>	6379/TCP
service/redis-slave 1h	ClusterIP	10.100.200.194	<none>	6379/TCP
NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE
AGE				
deployment.apps/frontend 1h	3	3	3	3
deployment.apps/redis-master 1h	1	1	1	1
deployment.apps/redis-slave 1h	2	2	2	2
NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/frontend-7776b5c947	3	3	3	1h
replicaset.apps/redis-master-5c95c89d9	1	1	1	1h
replicaset.apps/redis-slave-6548b4f8b9	2	2	2	1h

PKS.Master.PowerOffOn

This test verifies that running K8s cluster master node can rejoin the cluster after it is powered off/on, when the Stateless application is running.

Note This test is mandatory for Stateless applications only.

Mandatory for Stateless Applications Only

Test ID	PKS.Master.PowerOffOn
Configuration	PKS cluster is created. Application is running.

Mandatory for Stateless Applications Only

Test Script

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 2 Check the cluster existence.

```
root@cli-vm:~# pks clusters
```

Name	Plan Name	UUID	Status
k8s-cluster-101	small	53187740-e5cf-4ed7-9181-ab71bbdc1945	succeeded

- 3 Get credentials for the cluster.

```
root@cli-vm:~# pks get-credentials k8s-cluster-101
```

```
Fetching credentials for cluster k8s-cluster-101.
Context set for cluster k8s-cluster-101.
```

You can now switch between clusters by using:

```
$kubectl config use-context <cluster-name>
```

- 4 Set kubectl to use the context of the cluster.

```
root@cli-vm:~# kubectl config use-context k8s-cluster-101
```

```
Switched to context "k8s-cluster-101".
```

- 5 Verify the Application deployment readiness.

```
root@cli-vm:~# kubectl get all
```

- 6 Get the VM name for the master node under the power off/on test.

```
root@cli-vm:~# source A-BOSH.env
```

```
root@cli-vm:~# bosh vms
```

```
...
Deployment 'service-instance_a2a02b92-0ac5-4bfd-9706-fcfeae528d2b'

Instance                               Process State AZ
IPs      VM CID                        VM Type      Active
master/4a7a6e3c-fcf3-4fa7-950b-7031c68f67ba running      AZ-COMP
172.23.2.2 vm-04c898ea-6614-4b08-a22a-c50bed645293 medium.disk true
```

- 7 Access vSphere Client. See [Accessing vSphere Client](#) for details.
 - 8 Highlight vm-04c898ea-6614-4b08-a22a-c50bed645293 under RegionA01-COMP01/PKS-COMP cluster and select **Power > Power Off** from the context menu.
 - 9 Go back to the PuTTY session to CLI-VM and verify that the Application deployment cannot be accessed.
- ```
root@cli-vm:~# kubectl get all
```
- 10 Access vSphere Client.
  - 11 Highlight vm-04c898ea-6614-4b08-a22a-c50bed645293 under RegionA01-COMP01/PKS-COMP cluster and select **Power > Power On** from the context menu.
  - 12 Go back to the PuTTY session to CLI-VM and monitor the application deployment status by running the `kubectl get all` command. Wait until the master node is shown as **Available** (it might take up to five minutes).

## Test Output

- 1 Application deployment is running and ready (see Test script step 5).

**Mandatory for Stateless Applications Only**

```
root@cli-vm:~# kubectl get all
```

| NAME                             | READY | STATUS  | RESTARTS | AGE |
|----------------------------------|-------|---------|----------|-----|
| pod/frontend-7776b5c947-7clrm    | 1/1   | Running | 2        | 2h  |
| pod/frontend-7776b5c947-96jlh    | 1/1   | Running | 0        | 2h  |
| pod/frontend-7776b5c947-xsrv7    | 1/1   | Running | 0        | 2h  |
| pod/redis-master-5c95c89d9-78ds4 | 1/1   | Running | 0        | 2h  |
| pod/redis-slave-6548b4f8b9-5p92d | 1/1   | Running | 0        | 2h  |
| pod/redis-slave-6548b4f8b9-x9fd9 | 1/1   | Running | 0        | 2h  |

| NAME                 | TYPE      | CLUSTER-IP     | EXTERNAL-IP | PORT(S)  |
|----------------------|-----------|----------------|-------------|----------|
| service/frontend     | ClusterIP | 10.100.200.152 | <none>      | 80/TCP   |
| service/kubernetes   | ClusterIP | 10.100.200.1   | <none>      | 443/TCP  |
| service/redis-master | ClusterIP | 10.100.200.55  | <none>      | 6379/TCP |
| service/redis-slave  | ClusterIP | 10.100.200.194 | <none>      | 6379/TCP |

| NAME                         | DESIRED | CURRENT | UP-TO-DATE | AVAILABLE |
|------------------------------|---------|---------|------------|-----------|
| deployment.apps/frontend     | 3       | 3       | 3          | 3         |
| deployment.apps/redis-master | 1       | 1       | 1          | 1         |
| deployment.apps/redis-slave  | 2       | 2       | 2          | 2         |

| NAME                                   | DESIRED | CURRENT | READY | AGE |
|----------------------------------------|---------|---------|-------|-----|
| replicaset.apps/frontend-7776b5c947    | 3       | 3       | 3     | 2h  |
| replicaset.apps/redis-master-5c95c89d9 | 1       | 1       | 1     | 2h  |
| replicaset.apps/redis-slave-6548b4f8b9 | 2       | 2       | 2     | 2h  |

- 2 Application deployment cannot be accessed because the master node is down (see Test script step 9).

```
root@cli-vm:~# kubectl get all
```

```
No resources found.
Unable to connect to the server: dial tcp 24.24.24.10:8443: i/o timeout
Unable to connect to the server: dial tcp 24.24.24.10:8443: i/o timeout
Unable to connect to the server: dial tcp 24.24.24.10:8443: i/o timeout
Unable to connect to the server: dial tcp 24.24.24.10:8443: i/o timeout
Unable to connect to the server: dial tcp 24.24.24.10:8443: i/o timeout
Unable to connect to the server: dial tcp 24.24.24.10:8443: i/o timeout
Unable to connect to the server: dial tcp 24.24.24.10:8443: i/o timeout
Unable to connect to the server: dial tcp 24.24.24.10:8443: i/o timeout
Unable to connect to the server: dial tcp 24.24.24.10:8443: i/o timeout
Unable to connect to the server: dial tcp 24.24.24.10:8443: i/o timeout
```

- 3 Application deployment is running and ready for all nodes (see Test script step 12).

```
root@cli-vm:~# kubectl get all
```

| NAME                             | READY | STATUS  | RESTARTS | AGE |
|----------------------------------|-------|---------|----------|-----|
| pod/frontend-7776b5c947-7clrm    | 1/1   | Running | 2        | 2h  |
| pod/frontend-7776b5c947-96jlh    | 1/1   | Running | 0        | 2h  |
| pod/frontend-7776b5c947-xsrv7    | 1/1   | Running | 0        | 2h  |
| pod/redis-master-5c95c89d9-78ds4 | 1/1   | Running | 0        | 2h  |
| pod/redis-slave-6548b4f8b9-5p92d | 1/1   | Running | 0        | 2h  |

**Mandatory for Stateless Applications Only**

|                                  |           |                |             |           |    |
|----------------------------------|-----------|----------------|-------------|-----------|----|
| pod/redis-slave-6548b4f8b9-x9fd9 |           | 1/1            | Running     | 0         | 2h |
| NAME                             | TYPE      | CLUSTER-IP     | EXTERNAL-IP | PORT(S)   |    |
| AGE                              |           |                |             |           |    |
| service/frontend                 | ClusterIP | 10.100.200.152 | <none>      | 80/TCP    |    |
| 2h                               |           |                |             |           |    |
| service/kubernetes               | ClusterIP | 10.100.200.1   | <none>      | 443/TCP   |    |
| 2d                               |           |                |             |           |    |
| service/redis-master             | ClusterIP | 10.100.200.55  | <none>      | 6379/TCP  |    |
| 2h                               |           |                |             |           |    |
| service/redis-slave              | ClusterIP | 10.100.200.194 | <none>      | 6379/TCP  |    |
| 2h                               |           |                |             |           |    |
| NAME                             | DESIRED   | CURRENT        | UP-TO-DATE  | AVAILABLE |    |
| AGE                              |           |                |             |           |    |
| deployment.apps/frontend         | 3         | 3              | 3           | 3         |    |
| 2h                               |           |                |             |           |    |
| deployment.apps/redis-master     | 1         | 1              | 1           | 1         |    |
| 2h                               |           |                |             |           |    |
| deployment.apps/redis-slave      | 2         | 2              | 2           | 2         |    |
| 2h                               |           |                |             |           |    |

# Collecting Logs

## 8

After the validation is completed and all tests are passed, collect the logs from the PKS cluster nodes. Ensure that the Application solution is deployed and running in a good state after the validation tests are done. VMware might need to verify the Application solution deployment during the validation results review.

### Procedure

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#).

Run the following commands:

```
root@cli-vm:~# mkdir -p /home/vmware/app-cluster-logs/master
root@cli-vm:~# mkdir -p /home/vmware/app-cluster-logs/worker0
root@cli-vm:~# mkdir -p /home/vmware/app-cluster-logs/worker1
root@cli-vm:~# mkdir -p /home/vmware/app-cluster-logs/worker2
```

- 2 SSH to the PCF OpsManager VM.

```
root@cli-vm:~# ssh ubuntu@23.23.23.2
```

```
Enter passphrase for key '/root/.ssh/id_rsa': VMware1!
ubuntu@pks-ops-mgr:~#
```

- 3 To get the PKS cluster deployment ID and node information, run the bosh command:

```
ubuntu@pks-ops-mgr:~$ source A-BOSH.env
```

```
ubuntu@pks-ops-mgr:~$ bosh vms
```

```
Using environment '10.1.1.3' as client 'ops_manager'
```

```
Deployment 'harbor-container-registry-8d51221a98ebb37124f6'
```

| Instance                                        | Process     | State  | AZ      | IPs      | VM  |
|-------------------------------------------------|-------------|--------|---------|----------|-----|
| CID                                             | VM Type     | Active |         |          |     |
| harbor-app/5855fc38-e432-4585-818b-013bd2b839b3 | running     |        | AZ-MGMT | 10.1.1.5 | vm- |
| bccb123b-3d90-410b-97d3-fe6c011468a5            | xlarge.disk | true   |         |          |     |
| 1 vms                                           |             |        |         |          |     |

```
Deployment 'pivotal-container-service-0b648dc45ece40cd92e1'
```

| Instance | Process | State  | AZ | IPs |
|----------|---------|--------|----|-----|
| VM CID   | VM Type | Active |    |     |

```
pivotal-container-service/0d785096-429e-4267-b6b6-89b5d0295461 running AZ-MGMT 10.1.1.4
vm-4915010b-b6d7-4ce1-8fda-693efe81ba3e large true
1 vms

Deployment 'service-instance_1a385d84-c0cd-4dec-984a-f7735a9500af'
Instance Process State AZ IPs VM
CID VM Type Active
master/43ad8d9c-0469-42eb-b0d1-0ceeb8f99a16 running AZ-COMP 172.23.1.2
vm-5df9079f-9dbc-4a68-a2da-b0e39552a010 medium.disk true
worker/899f1f7b-28ff-4b2c-addb-b52fa56d07cb running AZ-COMP 172.23.1.3 vm-554d644c-
d8cc-4fd2-8caa-3a1c4937ca9f medium.disk true
worker/c7b24def-d3b9-465f-93af-93ea99d1faa1 running AZ-COMP 172.23.1.4 vm-38f0bbcf-
cef1-4afe-b99d-676b2677f62e medium.disk true
worker/ec488113-899d-409c-b1e0-6e8254868618 running AZ-COMP 172.23.1.5 vm-4e2a6d48-
bcb3-40aa-9048-1ac153a6ec9f medium.disk true
4 vms
Succeeded
```

- 4 From the PCF OpsManager VM, SSH to the PKS cluster master node.

```
ubuntu@pks-ops-mgr:~$ bosh -e 10.1.1.3 -d service-instance_1a385d84-
c0cd-4dec-984a-f7735a9500af ssh master/0
```

```
Using environment '10.1.1.3' as client 'ops_manager'
Using deployment 'service-instance_1a385d84-c0cd-4dec-984a-f7735a9500af'
Task 415084. Done
master/43ad8d9c-0469-42eb-b0d1-0ceeb8f99a16:~$
```

- 5 Collect the *kube-apiserver* and *kube-controller-manager* logs from the master node by sudo SCP transferring them into /home/vmware on the CLI-VM.

```
master/43ad8d9c-0469-42eb-b0d1-0ceeb8f99a16:~$ ls -l /var/vcap/sys/log/kube-
apiserver/
```

```
...
-rw----- 1 vcap vcap 12237564 Dec 8 02:03 kube-apiserver.stderr.log
-rw----- 1 vcap vcap 0 Dec 5 00:43 kube-apiserver.stdout.log
```

```
master/43ad8d9c-0469-42eb-b0d1-0ceeb8f99a16:~$ ls -l /var/vcap/sys/log/kube-
controller-manager/
```

```
...
-rw----- 1 vcap vcap 179295 Dec 8 01:59 kube-controller-manager.stderr.log
-rw----- 1 vcap vcap 0 Dec 5 00:43 kube-controller-manager.stdout.log
```

```
master/43ad8d9c-0469-42eb-b0d1-0ceeb8f99a16:~$ exit
```

```
logout
Connection to 172.23.1.2 closed.
Succeeded
```

sudo scp command example:



```
ubuntu@pks-ops-mgr:~$ sudo scp -r /var/vcap/sys/log/kube-apiserver/*
root@192.168.110.7:/home/vmware/app-cluster-logs/master/
```

- 6 From the PCF OpsManager VM, SSH to the PKS cluster worker node.

```
ubuntu@pks-ops-mgr:~$ bosh -d service-instance_1a385d84-c0cd-4dec-984a-
f7735a9500af ssh worker/0
```

```
Using environment '10.1.1.3' as client 'ops_manager'
Using deployment 'service-instance_1a385d84-c0cd-4dec-984a-f7735a9500af'
Task 415703. Done
worker/899f1f7b-28ff-4b2c-addb-b52fa56d07cb:~$
```

- 7 Collect the *kubelet* logs from the worker node by sudo SCP transferring them into /home/vmware on the CLI-VM.

```
worker/899f1f7b-28ff-4b2c-addb-b52fa56d07cb:~$ ls -l /var/vcap/sys/log/kubelet/
```

```
total 380
-rw-r--r-- 1 root root 1605 Dec 5 00:49 kubelet_ctl.stderr.log
-rw-r--r-- 1 root root 81 Dec 5 00:49 kubelet_ctl.stdout.log
-rw-r--r-- 1 root root 368080 Dec 8 02:02 kubelet.stderr.log
-rw-r--r-- 1 root root 31 Dec 5 00:49 kubelet.stdout.log
-rw-r----- 1 root root 0 Dec 5 00:49 post-start.stderr.log
-rw-r----- 1 root root 3180 Dec 5 00:52 post-start.stdout.log
```

sudo scp command example:

```
ubuntu@pks-ops-mgr:~$ sudo scp -r /var/vcap/sys/log/kubelet/*
root@192.168.110.7:/home/vmware/app-cluster-logs/worker0/
```

- 8 Repeat steps 6 and 7 for all worker nodes.
- 9 Create a *tgz* bundle with all logs from the master and worker nodes.

```
root@cli-vm:~# cd /home/vmware
```

```
root@cli-vm:/home/vmware# tar cvfz app-cluster-logs.tgz app-cluster-logs/
```

```
app-cluster-logs/
app-cluster-logs/worker0/
app-cluster-logs/worker0/kubelet.stderr.log
app-cluster-logs/worker0/kubelet_ctl.stderr.log
app-cluster-logs/master/
app-cluster-logs/master/kube-apiserver.stderr.log
app-cluster-logs/master/kube-controller-manager.stderr.log
app-cluster-logs/worker1/
app-cluster-logs/worker1/kubelet.stderr.log
app-cluster-logs/worker1/kubelet_ctl.stderr.log
app-cluster-logs/worker2/
app-cluster-logs/worker2/kubelet.stderr.log
app-cluster-logs/worker2/kubelet_ctl.stderr.log
```

- 10 Use WinSCP to transfer (see [Downloading and Transferring Files](#) for details) the *tgz* bundle from CLI-VM to the control center Windows VM.

#### **What to do next**

Submit the results.

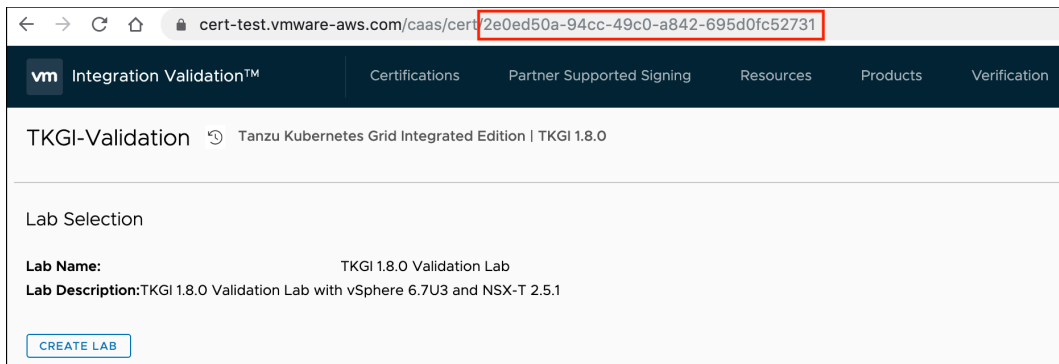
# Submitting Validation Results

## 9

After collecting the PKS cluster logs, you must submit the validation results to VMware.

### Procedure

- 1 Collect PKS Cluster logs. See [Chapter 8 Collecting Logs](#) for details.
- 2 Fill the VMware Tanzu Kubernetes Grid Integrated Edition Validation-in-Cloud - Validation Tests Check List.
- 3 Upload the PKS cluster logs and the VMware Tanzu Kubernetes Grid Integrated Edition Validation-in-Cloud - Validation Tests Check List to the DCPN case.
- 4 Enter the Viva Session ID (see [Lab Selection](#)) in the DCPN case - comment section to identify the validation session.



- 5 From the web browser of your Main Console VM, enter the URL: `http://10.148.172.7`. An internal IP address is displayed. Also, an internal IP address is displayed on the Main Console VM desktop. Report this IP address in the DCPN case.

# Requesting Partner-Ready VMware Tanzu Logo

10

After the VMware Validation team confirms in the DCPN case that your solution meets the validation requirements, perform the following steps:

## Procedure

- 1 Request for a Partner-Ready logo.

Write to [vsxalliance@vmware.com](mailto:vsxalliance@vmware.com) requesting for a Partner-Ready VMware Tanzu Logo.

- 2 Submit a support statement.

- a Go to <http://vmware-alliances.force.com/SubmitSupport>.
- b Enter the necessary information in the requested fields.
- c In the **Application Category** drop-down menu, select the appropriate application category.
- d Click **Add Support Statement**.
- e In the Support Statement page, read through the support statement, and click **Submit**.

# Working with the Lab

# 11

This chapter provides general information about VMware Validation Lab environment and ways to work with the lab components.

This chapter includes the following topics:

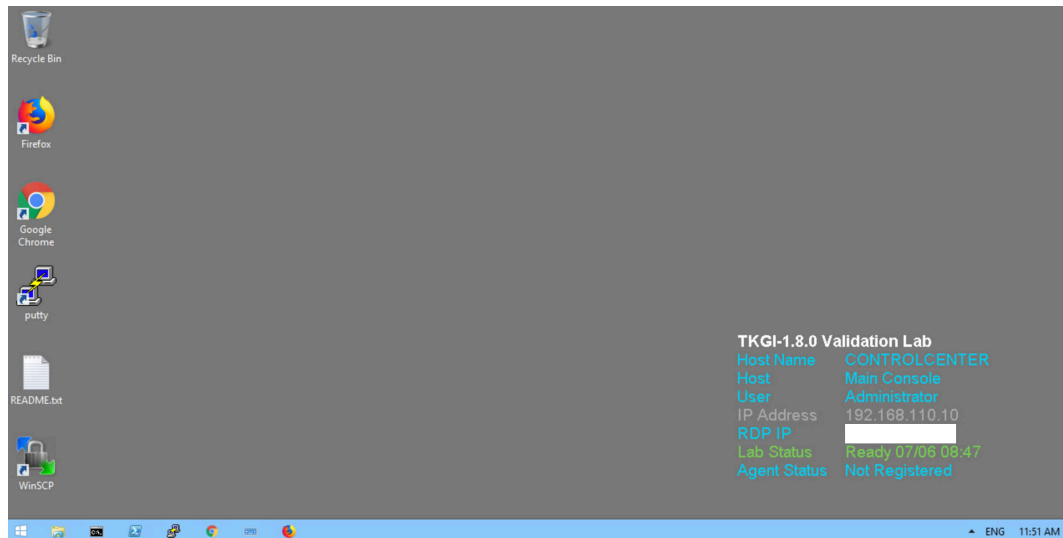
- [Utilities Installed on the Lab](#)
- [Downloading and Transferring Files](#)
- [Accessing CLI VM](#)
- [Accessing TKGI Environment](#)
- [Installing Helm Package Manager](#)
- [Accessing Harbor Repository](#)
- [Steps to Add Storage Class](#)
- [Accessing vSphere Client](#)
- [Load Balancers](#)
- [Adding Persistent Storage](#)

## Utilities Installed on the Lab

The following utilities and applications are pre-installed on the control center Windows VM;

- Firefox web browser
- Google Chrome web browser
- PuTTY
- WinSCP
- Notepad++ (go to **Start > All Programs**).

See the README.txt file on the Desktop for Software BOM and Lab component credentials.



## Downloading and Transferring Files

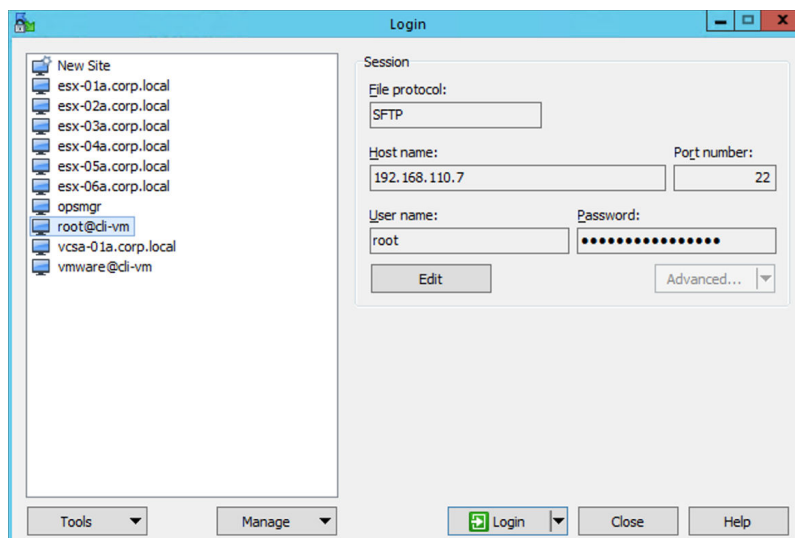
You download files from the Internet as follows:

- If your Application deployment requires downloading files from the Internet, run these commands on the CLI-VM:

```
git clone
wget
curl
```

- If your Application deployment requires downloading files from the Internet using a web interface, use the Firefox browser from the control center Windows VM.

To transfer files from Windows VM to the CLI-VM, use WinSCP from the control center Windows VM. Connect to the CLI-VM 192.168.110.7 (user name: *root* ; password: *VMware1!*) and copy the files into / root on CLI-VM.

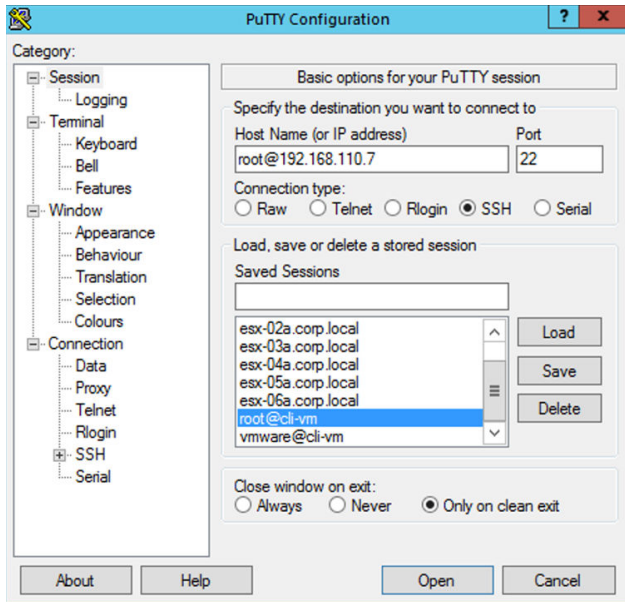


## Accessing CLI VM

To access the CLI-VM, perform the following steps:

### Procedure

- 1 Open PuTTY by clicking the PuTTY icon on the Desktop.
- 2 Select **root@cli-vm** or **Default Settings** in the Stored Sessions list, and click Load.



- 3 Click Open. A prompt opens requesting to enter the password for root@192.168.110.7.
- 4 Enter the password - *VMware1!*

### Results

The Command Prompt for CLI VM opens.

## Accessing TKGI Environment

This section provides information about how to create a PKS cluster by accessing the TKGI Lab through the VMware Learning Platform (VLP) control center Windows VM.

To log in to PKS, perform the following steps:

- 1 Open PuTTY session for the CLI-VM. See [Accessing CLI VM](#).
- 2 Log in to PKS.

```
root@cli-vm:~# pks login -a uaa.corp.local -u vmware -p VMware1! -k
```

```
API Endpoint: uaa.corp.local
User: vmware
```

## Creating a PKS Cluster

The cluster creation can take a significant amount of time to complete (about 45–60 minutes).

You can use the command - `pks create-cluster` and specify:

- one of the predefined cluster plans: `--plan [small|medium|large]`
- number of workers: `--num-nodes <number>`

---

**Note** Do not use the same name for internal and external cluster names.

---

For PKS cluster plan details, see [Chapter 15 Appendix C: TKGI Cluster Plans](#).

If the Application requires resources customization for the cluster, plan can be modified from PCF Ops Manager. See the following PKS documentation for details:

- [Installing Enterprise PKS on vSphere with NSX-T](#)
- [Plans](#)

Run this command:

```
root@cli-vm:~# pks create-cluster k8s-cluster-101 --external-hostname pks-
cluster-101.corp.local --plan small --num-nodes 3
```

```
Name: k8s-cluster-101
Plan Name: small
UUID: 53187740-e5cf-4ed7-9181-ab71bbdc1945
Last Action: CREATE
Last Action State: in progress
Last Action Description: Creating cluster
Kubernetes Master Host: pks-cluster-101.corp.local
Kubernetes Master Port: 8443
Worker Nodes: 3
Kubernetes Master IP(s): In Progress
Network Profile Name:
Use pks cluster k8s-cluster-101 to monitor the state of your cluster
```

To check the cluster creation task status (it might take up to 45–60 minutes), run the `pks cluster cluster-name` command.

## Checking PKS Cluster Status

After creating the PKS cluster, you must check the cluster name and other details.



## Procedure

- 1 To get the PKS cluster name and status, run this command:

```
root@cli-vm:~# pks clusters
```

| Name            | Plan Name | UUID                                 | Status    | Action |
|-----------------|-----------|--------------------------------------|-----------|--------|
| k8s-cluster-101 | small     | 53187740-e5cf-4ed7-9181-ab71bbdc1945 | succeeded | CREATE |

- 2 To get PKS cluster details, run this command:

```
root@cli-vm:~# pks cluster k8s-cluster-101
```

```
Name: k8s-cluster-101
Plan Name: small
UUID: 53187740-e5cf-4ed7-9181-ab71bbdc1945
Last Action: CREATE
Last Action State: succeeded
Last Action Description: Instance provisioning completed
Kubernetes Master Host: pks-cluster-101.corp.local
Kubernetes Master Port: 8443
Worker Nodes: 3
Kubernetes Master IP(s): 24.24.24.10
Network Profile Name:
```

- 3 The *pks-cluster-101.corp.local* entry must be added to the */etc/hosts* file, otherwise *kubect!* fails to communicate with *k8s-cluster-101*.

```
root@cli-vm:~# vi /etc/hosts
```

Add *24.24.24.10 pks-cluster-101.corp.local* line to the end of the */etc/hosts* file.

- 4 Fetch the credentials for cluster *k8s-cluster-101*:

```
root@cli-vm:~# pks get-credentials k8s-cluster-101
```

```
Fetching credentials for cluster k8s-cluster-101.
Context set for cluster k8s-cluster-101.
User can now switch between clusters by using:
$kubectl config use-context <cluster-name>
```

```
root@cli-vm:~# kubectl config use-context k8s-cluster-101
```

```
Switched to context "k8s-cluster-101".
```

## Validating Nodes for Cluster

To validate nodes for cluster *k8s-cluster-101*, run the following command:

```
root@cli-vm:~# kubectl get nodes -o wide
```

| NAME                                  | STATUS             | ROLES             | AGE                 | VERSION | INTERNAL-IP | EXTERNAL- |
|---------------------------------------|--------------------|-------------------|---------------------|---------|-------------|-----------|
| IP                                    | OS-IMAGE           | KERNEL-VERSION    | CONTAINER-RUNTIME   |         |             |           |
| afc7bfbdb-bea7-4044-8d81-2ed6121eec78 | Ready              | <none>            | 24m                 | v1.11.2 | 172.23.0.4  |           |
| 172.23.0.4                            | Ubuntu 16.04.5 LTS | 4.15.0-33-generic | docker://17.12.1-ce |         |             |           |
| cd48fdbbc-5f15-437c-9a1b-f70235c59a62 | Ready              | <none>            | 31m                 | v1.11.2 | 172.23.0.3  |           |
| 172.23.0.3                            | Ubuntu 16.04.5 LTS | 4.15.0-33-generic | docker://17.12.1-ce |         |             |           |
| dc8ca134-93bb-422a-89d0-89a2ba359685  | Ready              | <none>            | 17m                 | v1.11.2 | 172.23.0.5  |           |
| 172.23.0.5                            | Ubuntu 16.04.5 LTS | 4.15.0-33-generic | docker://17.12.1-ce |         |             |           |

## Installing Helm Package Manager

After the PKS cluster is created, you can install the Helm package manager. Run the following commands:

### Procedure

- 1 root@cli-vm:~# curl <https://raw.githubusercontent.com/kubernetes/helm/master/scripts/get> | bash

```
% Total % Received % Xferd Average Speed Time Time Time Current
 Dload Upload Total Spent Left Speed
100 7230 100 7230 0 0 4160 0 0:00:01 0:00:01 --:--:-- 4159
Downloading https://kubernetes-helm.storage.googleapis.com/helm-v2.12.3-linux-amd64.tar.gz
Preparing to install helm and tiller into /usr/local/bin
helm installed into /usr/local/bin/helm
tiller installed into /usr/local/bin/tiller
Run 'helm init' to configure helm.
```

- 2 root@cli-vm:~# kubectl -n kube-system create sa tiller && kubectl create clusterrolebinding tiller --clusterrole cluster-admin --serviceaccount=kube-system:tiller

```
serviceaccount/tiller created
clusterrolebinding.rbac.authorization.k8s.io/tiller created
```

- 3 root@cli-vm:~# helm init --skip-refresh --upgrade --service-account tiller

```
Creating /root/.helm
Creating /root/.helm/repository
...
Creating /root/.helm/cache/archive
Creating /root/.helm/repository/repositories.yaml
Adding stable repo with URL: https://kubernetes-charts.storage.googleapis.com
Adding local repo with URL: http://127.0.0.1:8879/charts
$HELM_HOME has been configured at /root/.helm.

Tiller (the Helm server-side component) has been installed into your Kubernetes Cluster.

Please note: by default, Tiller is deployed with an insecure 'allow unauthenticated users' policy.
```

To prevent this, run `helm init` with the `--tiller-tls-verify` flag.  
 For more information on securing your installation see:  
[https://docs.helm.sh/using\\_helm/#securing-your-helm-installation](https://docs.helm.sh/using_helm/#securing-your-helm-installation)  
 Happy Helming!

4 root@cli-vm:~# helm version

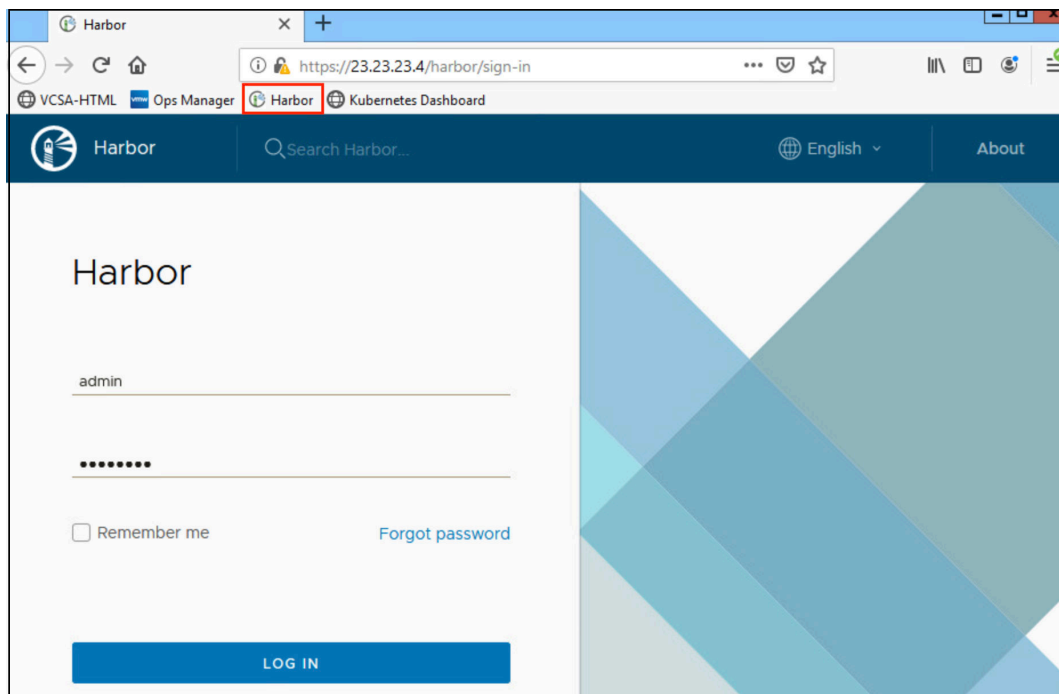
```
Client: &version.Version{SemVer:"v2.12.3", GitCommit:"2e55dbe1fdb5fdb96b75ff144a339489417b146b",
GitTreeState:"clean"}
Server: &version.Version{SemVer:"v2.12.3", GitCommit:"2e55dbe1fdb5fdb96b75ff144a339489417b146b",
GitTreeState:"clean"}
```

It might take some time (approximately 10 minutes) before the error message “Error: could not find a ready tiller pod” disappears and the server version is displayed.

## Accessing Harbor Repository

To access the harbor repository, open the Firefox browser, click the Harbor bookmark, and log in to Harbor.

- URL: *https://23.23.23.4*
- User name: *admin*
- Password: *VMware1!*



## Pushing Images into Harbor

To push unsigned images into Harbor from an external repository, perform the following steps:

## Procedure

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#).

- 2 Pull the images from the external repository (Google sample images are used as examples).

```
root@cli-vm:~# docker pull gcr.io/google-samples/gb-frontend:v4
```

```
root@cli-vm:~# docker pull gcr.io/google_containers/redis:e2e
```

```
root@cli-vm:~# docker pull gcr.io/google-samples/gb-redisslave:v1
```

To check the list of images, run this command:

```
root@cli-vm:~# docker images
```

| REPOSITORY<br>SIZE                           | TAG | IMAGE ID     | CREATED     |
|----------------------------------------------|-----|--------------|-------------|
| gcr.io/google-samples/gb-frontend<br>512MB   | v4  | e2b3e8542af7 | 2 years ago |
| gcr.io/google-samples/gb-redisslave<br>110MB | v1  | 5f026ddffa27 | 3 years ago |
| gcr.io/google_containers/redis<br>419MB      | e2e | e5e67996c442 | 3 years ago |

- 3 Tag the images.

```
root@cli-vm:~# docker tag gcr.io/google-samples/gb-frontend:v4 harbor.corp.local/
library/gb-frontend:v4
```

```
root@cli-vm:~# docker tag gcr.io/google-samples/gb-redisslave:v1
harbor.corp.local/library/gb-redisslave:v1
```

```
root@cli-vm:~# docker tag gcr.io/google_containers/redis:e2e harbor.corp.local/
library/redis:e2e
```

- 4 Push the images to Harbor.

```
root@cli-vm:~# docker push harbor.corp.local/library/gb-frontend:v4
```

```
root@cli-vm:~# docker push harbor.corp.local/library/gb-redisslave:v1
```

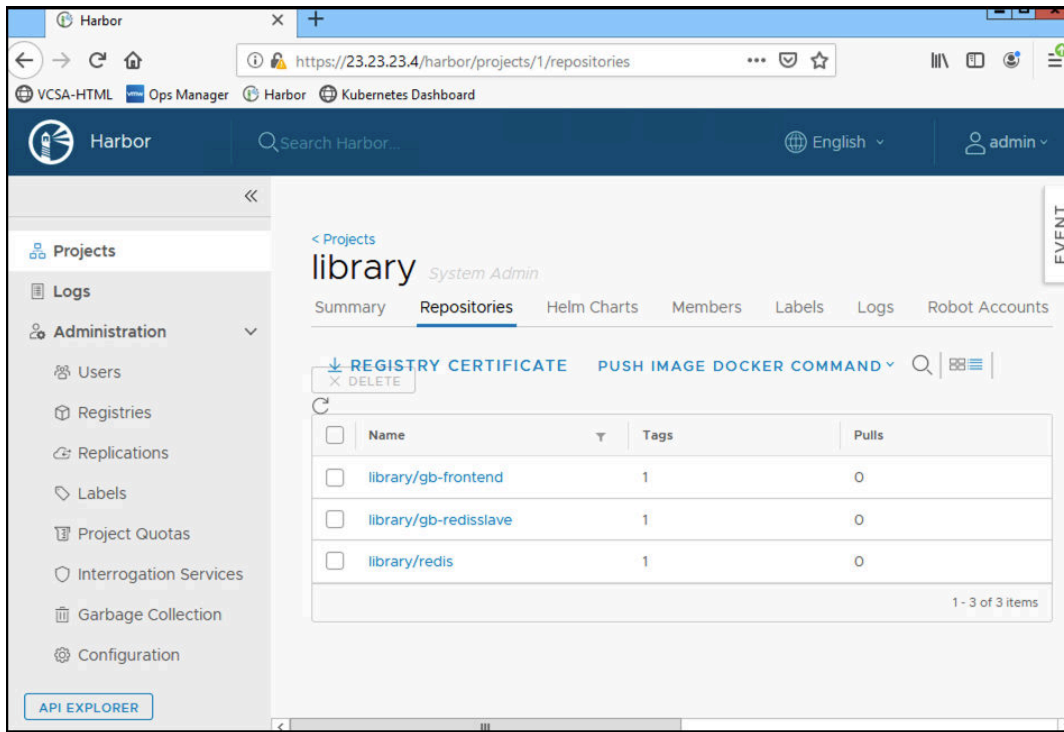
```
root@cli-vm:~# docker push harbor.corp.local/library/redis:e2e
```

- 5 On the Harbor web UI, verify that the three images are present.

Open the Firefox browser, click the Harbor bookmark (URL: <https://23.23.23.4>), and log in to Harbor.

- User name: *admin*
- Password: *VMware1!*

Go to **Projects > Repositories**.



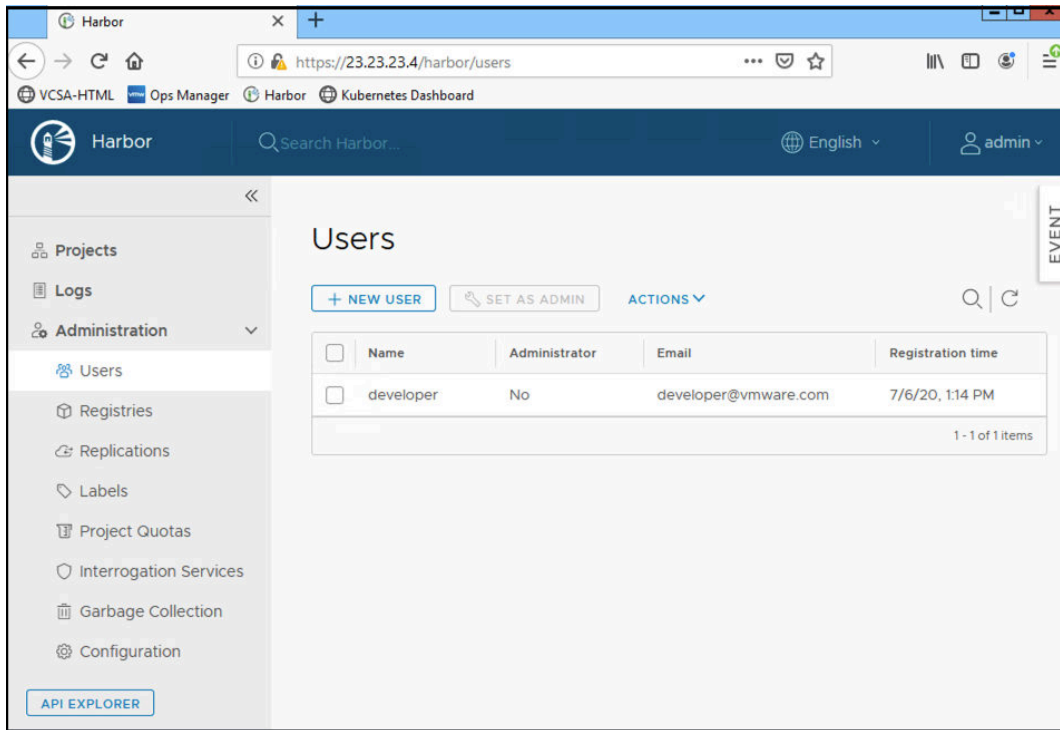
## Pushing Securely Signed Images to Harbor

To push securely signed images to Harbor, perform the following steps:

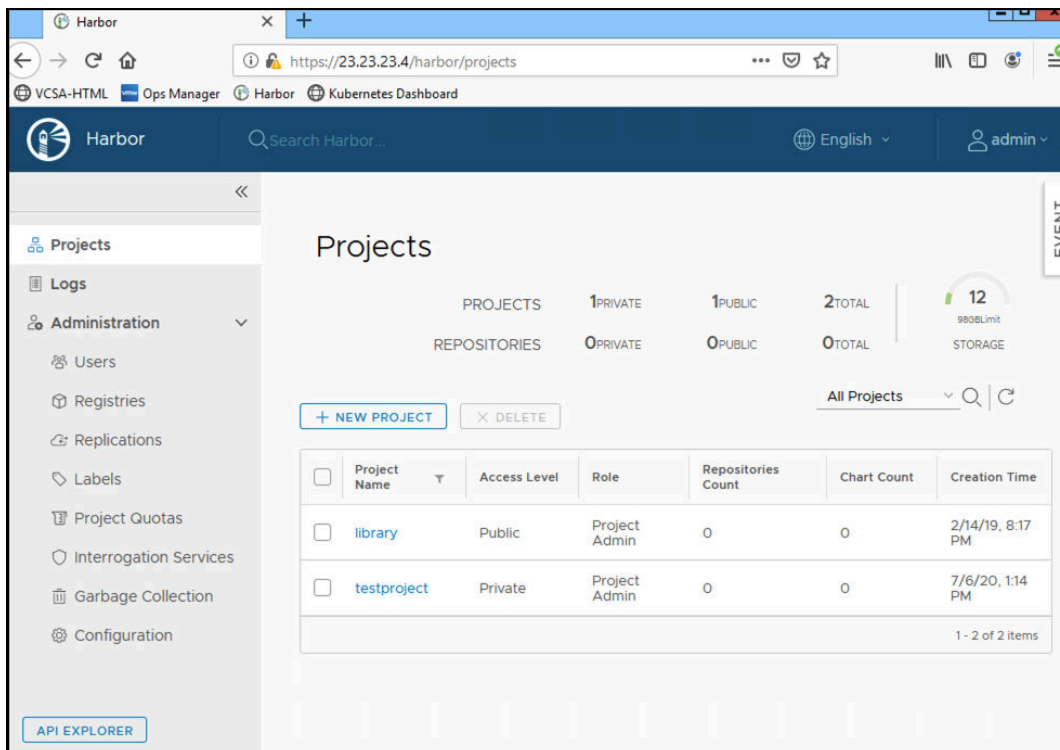
### Procedure

- 1 Log in to Harbor.

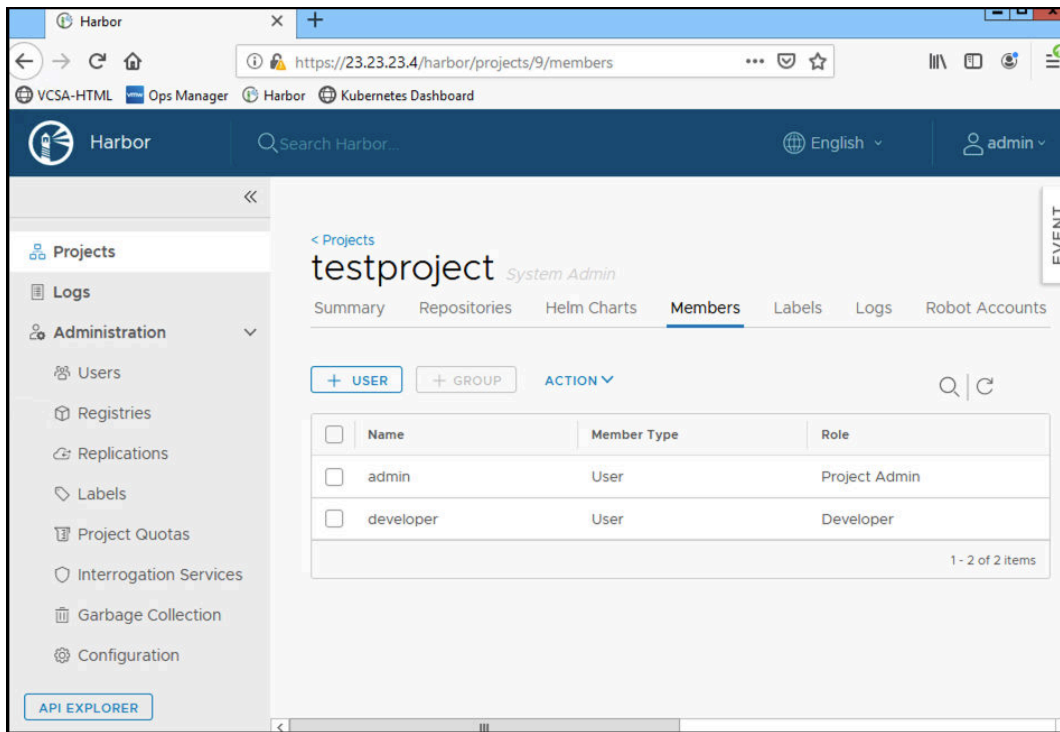
- 2 Go to **Administration > Users** and click **+NEW USER** (example: *developer*).



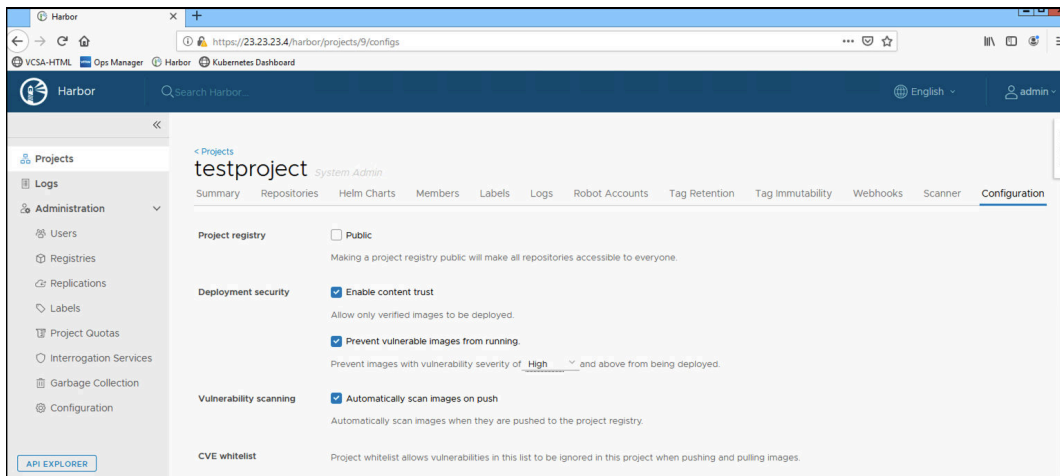
- 3 From the menu on the left, click **Projects** and create a new project by clicking **+NEW Project** (example: *testproject*).



- 4 Click the project that you created. Under the **Members** tab, click **+User** and add the user that you created in step 1 (*developer*). Ensure that the user is assigned the *Developer* role.



- 5 Go to the **Configuration** tab, set the values as follows, and click **SAVE**:
- Select **Enable Content Trust**.
  - Select **Prevent vulnerable images from running** and set the severity to “Low”
  - Select **Automatically scan images on push**.



- 6 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#) for details.
- 7 Pull the images from the external repository (example: Google sample images).
- ```
root@cli-vm:~# docker pull gcr.io/google-samples/gb-frontend:v4
```

```
root@cli-vm:~# docker pull gcr.io/google_containers/redis:e2e
root@cli-vm:~# docker pull gcr.io/google_samples/gb-redisslave:v1
```

- 8 Log in to Harbor registry as *developer*.

```
root@cli-vm:~# docker login harbor.corp.local -u developer
```

- 9 Set up two environment variables and tag the images.

```
root@cli-vm:~# export DOCKER_CONTENT_TRUST=1
root@cli-vm:~# export DOCKER_CONTENT_TRUST_SERVER=https://harbor.corp.local:4443
root@cli-vm:~# docker tag gcr.io/google_samples/gb-frontend:v4 harbor.corp.local/
testproject/gb-frontend:v4
root@cli-vm:~# docker tag gcr.io/google_samples/gb-redisslave:v1
harbor.corp.local/testproject/gb-redisslave:v1
root@cli-vm:~# docker tag gcr.io/google_containers/redis:e2e harbor.corp.local/
testproject/redis:e2e
```

- 10 Push the images into the Harbor repository. You are asked to enter a root key passphrase. You must enter the passphrase (*VMware1!*) every time you sign the image.

```
root@cli-vm:~# docker push harbor.corp.local/testproject/gb-frontend:v4
```

The push refers to repository [harbor.corp.local/testproject/gb-frontend]

```
3a31f3bf94a2: Layer already exists
cdc990c9b585: Layer already exists
...
816f1903c60f: Layer already exists
c12ecfd4861d: Layer already exists
v4: digest: sha256:aaa5b327ef3b4cb705513ab674fa40df66981616950c7de4912a621f9ee03dd4 size: 6968
Signing and pushing trust metadata
Enter passphrase for root key with ID ec1a991: VMware1!
Repeat passphrase for root key with ID ec1a991: VMware1!
Enter passphrase for new repository key with ID 7284e73: VMware1!
Repeat passphrase for new repository key with ID 7284e73: VMware1!
Finished initializing "harbor.corp.local/testproject/gb-frontend"
Successfully signed harbor.corp.local/testproject/gb-frontend:v4
```

```
root@cli-vm:~# docker push harbor.corp.local/testproject/gb-redisslave:v1
```

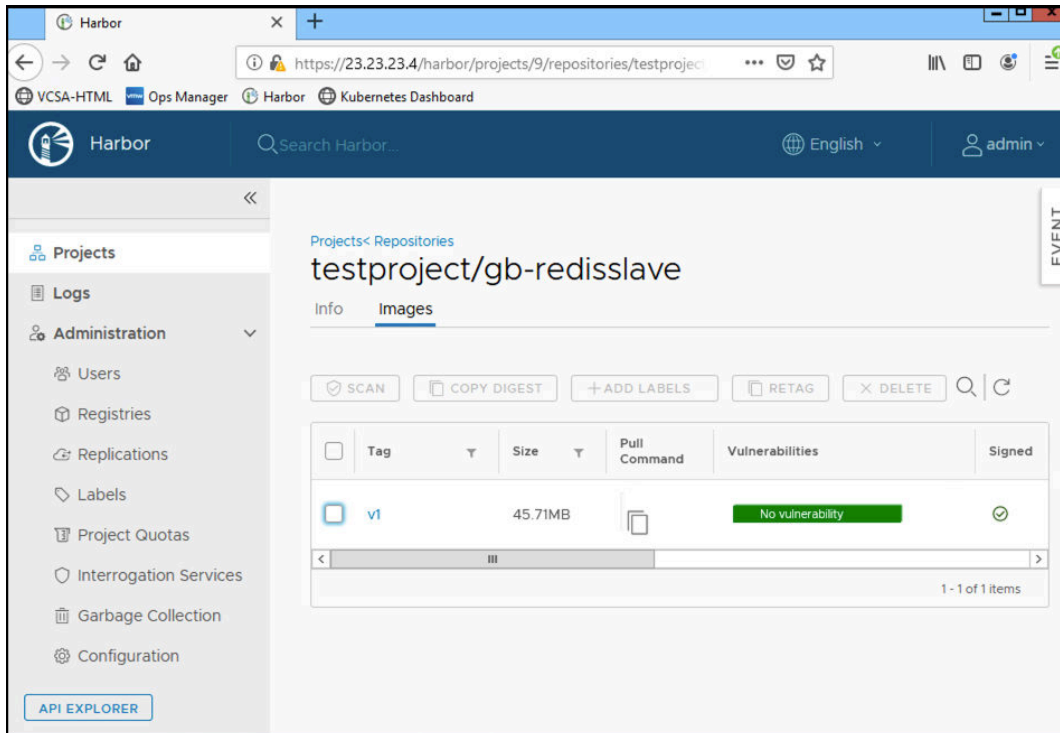
```
...
Signing and pushing trust metadata
Enter passphrase for root key with ID ec1a991: VMware1!
Enter passphrase for new repository key with ID 63e9a04: VMware1!
Repeat passphrase for new repository key with ID 63e9a04: VMware1!
Finished initializing "harbor.corp.local/testproject/gb-redisslave"
Successfully signed harbor.corp.local/testproject/gb-redisslave:v1
```



```
root@cli-vm:~# docker push harbor.corp.local/testproject/redis:e2e
```

```
...
Signing and pushing trust metadata
Enter passphrase for root key with ID ec1a991: VMware1!
Enter passphrase for new repository key with ID e7477bb: VMware1!
Repeat passphrase for new repository key with ID e7477bb: VMware1!
Finished initializing "harbor.corp.local/testproject/redis"
Successfully signed harbor.corp.local/testproject/redis:e2e
```

- 11 The tags (v4, v1, e2e in example) are marked as Signed in the Harbor UI. Go to **Projects > Repositories > testproject/gb-redisslave** (testproject/gb-frontend, or testproject/redis).



Steps to Add Storage Class

Procedure

- 1 Start a PuTTY session to CLI-VM. See [Accessing CLI VM](#).
- 2 Create a storage-class.yaml file.

```
root@cli-vm:~# vi storage-class.yaml
```

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: standard-sc
  annotations:
```

```
storageclass.kubernetes.io/is-default-class: "true"
provisioner: kubernetes.io/vsphere-volume
parameters:
  diskformat: thin
```

3 Save and exit the `storage-class.yaml` file.

4 In the Command Prompt, run the following commands:

```
root@cli-vm:~# kubectl apply -f storage-class.yaml
```

```
storageclass.storage.k8s.io/standard-sc created
```

```
root@cli-vm:~# kubectl get sc
```

NAME	PROVISIONER	AGE
standard-sc (default)	kubernetes.io/vsphere-volume	8s

For reference, see:

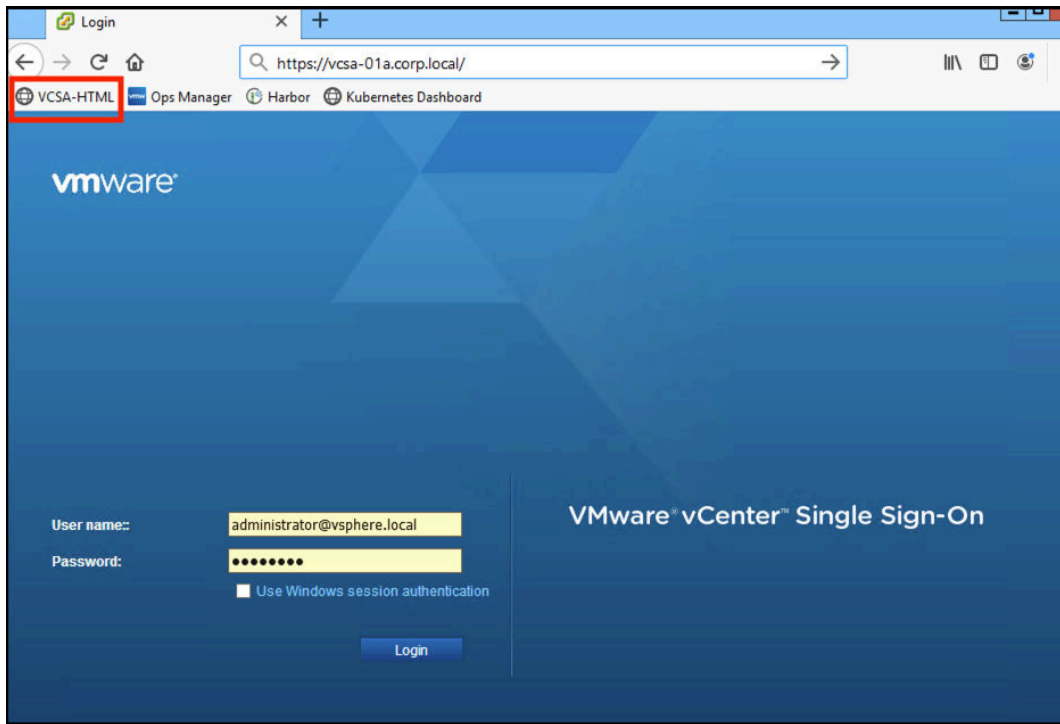
- Deploy a Kubernetes Stateful Application on PKS - WordPress: [https://code.vmware.com/samples/4835/deploy-a-kubernetes-stateful-application-on-pks---wordpress?h=Pivotal%20Container%20Service%20\(PKS\)](https://code.vmware.com/samples/4835/deploy-a-kubernetes-stateful-application-on-pks---wordpress?h=Pivotal%20Container%20Service%20(PKS)).
- Storage Classes: <https://kubernetes.io/docs/concepts/storage/storage-classes/>.

Accessing vSphere Client

To access the vSphere Client, perform the following steps:

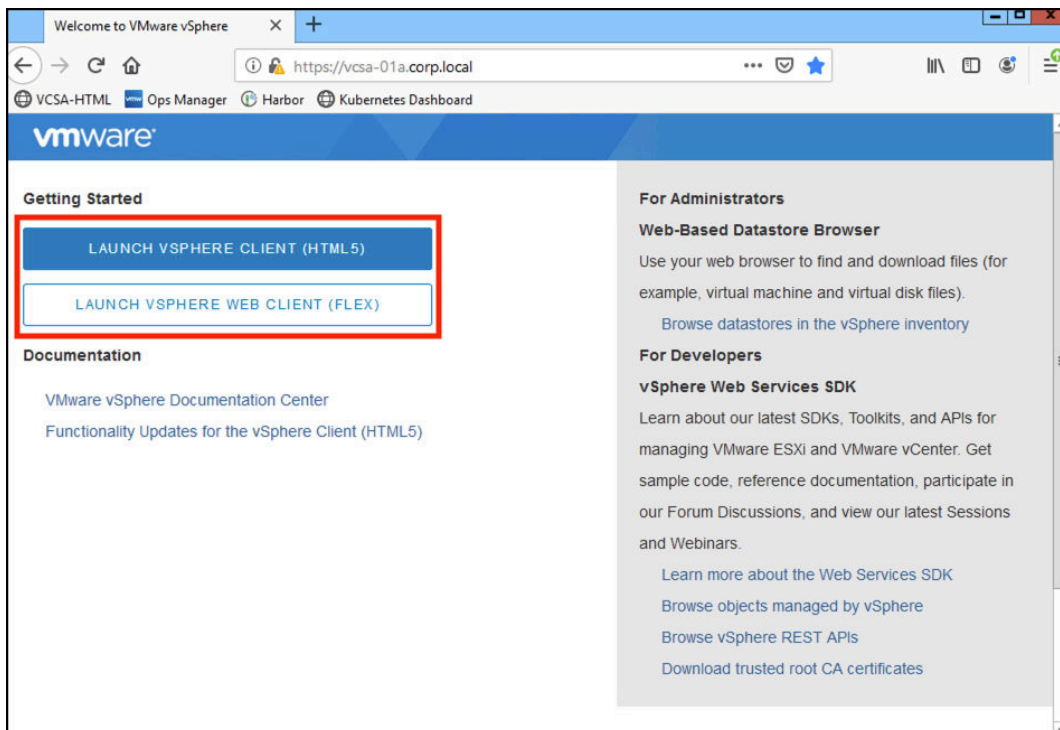
Procedure

- 1 From the control center Windows VM, open the Firefox web browser.
 - a Click the VCSA-HTML bookmark. The vCenter Server Appliance page opens.
 - b Log in to vCenter.
 - URL: *https://vcsa-01a.corp.local*
 - User name: *administrator@vsphere.local*
 - Password: *VMware1!*



2 Click any one of the following links:

- a *vSphere Web Client (Flex)*
- b *vSphere Client (HTML5)*



Load Balancers

There are no load balancers that are provided in the VMware Validation Lab for TKGI. You may, however, bring your own into the environment.

NSX-T IP pools are used for the associated external IP addresses that are issued to the load balancer service. There might be two IP addresses assigned to the load balancer. If more than one IP address is issued from the IP pools, only one IP address is used to contact the deployed application. The other IP address is used only for communication between the load balancer and the NSX-T router. Ensure that the correct IP address is used.

Note Do not specify any load balancer IP addresses during deployment (in YAML files, Helm charts, or others) as these IP addresses can conflict with the IPs that are issued from the IP pools.

Two IP addresses might be issued from NSX-T:

- 100.64.x.x/31 (This IP address is only accessible from the NSX-T router - DO NOT use this IP address).
- 24.24.24.x/24 (This IP address is the application IP and is routable from the Main Console VM).

Example output from `kubectl` (note down the EXTERNAL-IP field for gateway-external):

```
root@cli-vm:~# kubectl --namespace=test get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP
Service1	ClusterIP	10.100.200.159	<none>
TCP			9093/
gateway	ClusterIP	10.100.200.130	<none>
TCP			8080/
gateway-external	LoadBalancer	10.100.200.183	100.64.176.7, 24.24.24.22 8080:31924/TCP
Service2	ClusterIP	10.100.200.71	<none>
TCP			4222/
Service3	ClusterIP	10.100.200.197	<none>
TCP			9090/

Adding Persistent Storage

The VMware Enterprise PKS environment has limited storage available, which can be used as part of the deployment. The entire amount of available storage can be determined by using vSphere Client, but it will be at least a few 100 GBs.

Adding storage devices using thin provisioning allows the PKS environment to see larger devices and reduces errors in provisioning.

Adding Unmounted Block Devices to the Worker and Master Nodes

To add unmounted block devices to the PKS nodes in your cluster, simply add vSphere disks to the appropriate cluster VMs.

To determine the VM names of the PKS cluster nodes, run the following commands:

```
root@cli-vm:~# source A-BOSH.env
```

```
root@cli-vm:~# bosh vms
```

```
Using environment '10.1.1.3' as client 'ops_manager'
```

```
Task 36350
Task 36352
Task 36351
Task 36350 done
Task 36351 done
Task 36352 done
...
...
Deployment 'service-instance_ecd48568-cd51-4cca-a31e-558dc4442853'
```

Instance CID	Process State	VM Type	AZ	IPs	VM
master/4ca975d0-e083-4784-975f-19cd625132e1	running		AZ-COMP	172.23.1.2	
vm-6959007d-8fca-4433-9c86-d796a4cdb579	medium.disk true				
worker/029af88e-9436-4c81-b13e-00d32f138a48	running		AZ-COMP	172.23.1.4	vm-95309c86-
f989-4704-8d23-31fcbfd0bdee	medium.disk true				
worker/04bd7e45-885a-485f-877f-2fa3dc3fb103	running		AZ-COMP	172.23.1.3	vm-fdebbec0-
d556-4b69-9641-fa9121f20445	medium.disk true				

```
3 vms
```

```
root@cli-vm:~#
```

Using vSphere Client, add virtual disks to the appropriate VMs. It is recommended to use thin provisioning with these disks. The VMs can see the block devices, but they are unmounted on the cluster VMs.

To add the additional block device, perform the following steps:

Procedure

- 1 Right-click the VM and click **Edit Settings**.
- 2 Click **ADD NEW DEVICE** and select **Hard Disk**.
- 3 Choose the appropriate size for the virtual device.
- 4 Expand the options for the device and select **Thin Provision** for **Disk Provisioning**.

Repeat steps 1–4 for each disk that is to be added to the VM.

- 5 After adding all the disks, click **OK**.

Repeat steps 1–5 for each VM that requires additional unmounted block devices.

The screenshot shows the 'Edit Settings' window for a VM with ID 'vm-95309c86-f989-4704-8d23-31fcbfd0bdee'. The 'Virtual Hardware' tab is active. A table lists the hardware components:

Component	Value	Unit
CPU	2	
Memory	4	GB
Hard disk 1	3	GB
Hard disk 2	32	GB
Hard disk 3	50	GB
New Hard disk *	50	GB

Below the table, the 'Maximum Size' is 199.07 GB. The 'VM storage policy' is 'Datastore Default'. The 'Location' is 'Store with the virtual machine'. The 'Disk Provisioning' dropdown is set to 'Thin Provision'. The 'ADD NEW DEVICE' button is in the top right, and 'CANCEL' and 'OK' buttons are at the bottom right.

Adding Persistent Storage to a Pod Using Persistent Volume Claim

There are many different ways and different types of storage that can be configured as part of deployments. This section provides information about one statically provisioned storage example and one dynamically provisioned storage example.

The following subsections provide guidance on how to address additional storage requirements in the validation environment.

To configure persistent volumes with your CNA deployment, each YAML file must be applied using `kubectl`:

```
root@cli-vm:~# kubectl apply -f <filename>.yaml
```

Dynamic Persistent Volumes

Perform the following steps for a deployment:

Procedure

- 1 Create a storage class.
- 2 Define the Persistent Volume Claims.
- 3 Define the deployment.

Creating a Storage Class for Dynamically Provisioned Volumes

To dynamically provision storage with your pod, you must use a storage class. The following example shows a YAML manifest file to define a storage class. This file uses the shared datastore in the environment and defines the provisioned disks to be thin-provisioned. The value for the name field (*example-thin-disk*) will be used to reference this storage class when defining the persistent value claims.

```
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: example-thin-disk
provisioner: kubernetes.io/vsphere-volume
parameters:
  datastore: RegionA01-ISCSI01-COMP01
  diskformat: thin
  fstype: ext4
```

Defining Persistent Volume Claim for Deployments

For a dynamically created volume, the previously created storage class must be referenced in the Persistent Volume Claim (PVC) definition.

The following is an example PVC YAML manifest file for a dynamically provisioned volume:

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: example-disk-claim
  annotations:
    volume.beta.kubernetes.io/storage-class: example-thin-disk
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 50Gi
```

Defining a Deployment

The following is an example of a deployment YAML manifest file that references the PVC (this file focuses on the storage aspect of the deployment). This file can work with either the steps to create a static volume, or the steps to create a dynamic volume.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: example-app
spec:
  selector:
    ...
  template:
    ...
    spec:
```

```
...
volumes:
- name: example-data-disk
  persistentVolumeClaim:
    claimName: example-disk-claim
```

Statically Allocated Persistent Volumes

Statically allocating the virtual disks for persistent volumes can be a straight-forward way to make the virtual disks to be attached to the pods. To use statically allocated volumes, perform the following steps:

Procedure

- 1 Create the appropriate virtual disks.
- 2 Define the Persistent Volumes.
- 3 Define the Persistent Volume Claims.
- 4 Define the deployment.

Creating Virtual Disks for Statically Provisioned Volumes

Instead of using statically provisioned persistent volumes, you must first create the appropriate virtual disks. The following example assumes that the virtual disks are all created in their own directory on the datastore.

Log in to one of the ESXi hosts that has access to the shared storage for the environment (this example creates one virtual disk of 50 GB, thin provisioned):

```
root@esx-01a:~] cd /vmfs/volumes/RegionA01-ISCSI01-COMP01
root@esx-01a:~] mkdir kubestor
root@esx-01a:~] cd kubestor
root@esx-01a:~] vmkfstools -c 50G -d thin deployment_disk.vmdk
Create: 100% done.
```

Defining a Persistent Volume for Statically Provisioned Volumes

For static volumes, the next step is to define a persistent volume. It is important to match the capacity of the volume (spec field in the following example) to the size of the static disk created.

The following is an example of a Persistent Volume YAML file:

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: example-disk
spec:
  capacity:
    storage: 50Gi
  accessModes:
    - ReadWriteOnce
```



```

persistentVolumeReclaimPolicy: Retain
vsphereVolume:
  volumePath: "[RegionA01-ISCSI01-COMP01] kubestor/deployment_disk"
  fsType: ext4

```

Defining Persistent Volume Claim for Deployments

The following is an example Persistent Volume Claim (PVC) YAML manifest file for a statically provisioned volume:

```

apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: example-disk-claim
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 50Gi

```

Defining a Deployment

The following is an example of a deployment YAML manifest file that references the PVC (this file focuses on the storage aspect of the deployment). This file can work with either the steps to create a static volume, or the steps to create a dynamic volume.

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: example-app
spec:
  selector:
    ...
  template:
    ...
    spec:
      ...
      volumes:
        - name: example-data-disk
          persistentVolumeClaim:
            claimName: example-disk-claim

```

Tips and Troubleshooting

12

This chapter provides tips and troubleshooting information for using Validation-in-Cloud through the VMware Integration Validation (VIVA) service and the VMware Learning Platform (VLP).

- The VLP hosts the lab and also hosts the VMware Hands-on Labs (HOL). For HOL Online FAQs, see <https://communities.vmware.com/docs/DOC-24916>.
- After you receive access to the lab, ensure that you validate the infrastructure against your requirement.
- From the left pane, you can access any of the other VMs in the lab through the **Consoles** tab.
- For more screen space, click the **Toggle Full Screen Mode** button on the top right of the console.
- Adjust the screen resolution of the Main Console VM to a higher resolution. The web console matches the resolution output by the VM and fits it within the console window.
- Copy and Paste operations between the console window and your desktop are disabled for security reasons. However, you can use the **Send Text** button on the top left of the console to send text from your desktop to the VM.
- If you are a first-time user of the system, use the Guide icon for a guided tour of the console interface.
- To leave your lab and continue later, close the browser window.
- To delete the current lab instance and begin a new lab instance, click the **Reset** button.
- If you encounter issues with your browser, clear cookies from the browser and try again, or switch to a different browser. You can also use the Incognito/Private mode in the Chrome/Firefox browser.
- You might encounter an issue with the VLP when attempting to resume your validation lab, such as a response of *"Maintenance Mode, please try again later"*. In this case, close your browser, wait for a few minutes, and reopen the VLP site in a new browser window.
- For VMware support purposes only: from the Web browser on your Main Console VM, browse to <http://10.148.172.7>. This will return an internal IP address. Also, an internal IP Address is shown on the Main Console VM desktop. Report this IP address in the DCPN ticket as part of your submission case description.

This chapter includes the following topics:

- [Troubleshooting PKS Cluster Creation](#)
- [Accessing PCF Ops Manager](#)

- [Accessing NSX-T Manager](#)
- [Accessing Kubernetes Dashboard](#)

Troubleshooting PKS Cluster Creation

This section provides information about how to troubleshoot issues when creating a PKS cluster.

Restart PKS Instance

To make sure that all services started on pivotal-container-service work correctly, restart the PKS instance from Bosh.

Procedure

- 1 Open PuTTY session to CLI-VM. See [Accessing CLI VM](#).
- 2 On the Command Prompt, run the following commands:

```
root@cli-vm:~# source A-BOSH.env
root@cli-vm:~# bosh vms
```

```
Using environment '10.1.1.3' as client 'ops_manager'
...
Deployment 'harbor-container-registry-8d51221a98ebb37124f6'
harbor-app/5855fc38-e432-4585-818b-013bd2b839b3 running AZ-MGMT 10.1.1.5 vm-
bccb123b-3d90-410b-97d3-fe6c011468a5 xlarge.disk true
1 vms

Deployment 'pivotal-container-service-0b648dc45ece40cd92e1'
pivotal-container-service/0d785096-429e-4267-b6b6-89b5d0295461 running AZ-MGMT 10.1.1.4
vm-4915010b-b6d7-4ce1-8fda-693efe81ba3e large true
1 vms
```

- 3 Restart the PKS instance.

For example, if the service and the service ID are as follows:

pivotal-container-service/0d785096-429e-4267-b6b6-89b5d0295461 instance from deployment '*pivotal-container-service-0b648dc45ece40cd92e1*', run this command:

```
root@cli-vm:~# bosh restart pivotal-container-service/0d785096-429e-4267-
b6b6-89b5d0295461 -d pivotal-container-service-0b648dc45ece40cd92e1
```

```
Using environment '10.1.1.3' as client 'ops_manager'
Using deployment 'pivotal-container-service-0b648dc45ece40cd92e1'

Continue? [yN]: y
Task 79281 | 21:07:07 | Updating instance pivotal-container-service:
pivotal-container-service/0d785096-429e-4267-b6b6-89b5d0295461 (0) (canary) (00:02:41)
```

Get UAAC Token and Log In to PKS

The original UAAC token expires after 24 hours and it might prevent you from logging in to PKS. To get a UAAC token, perform the following steps:

Procedure

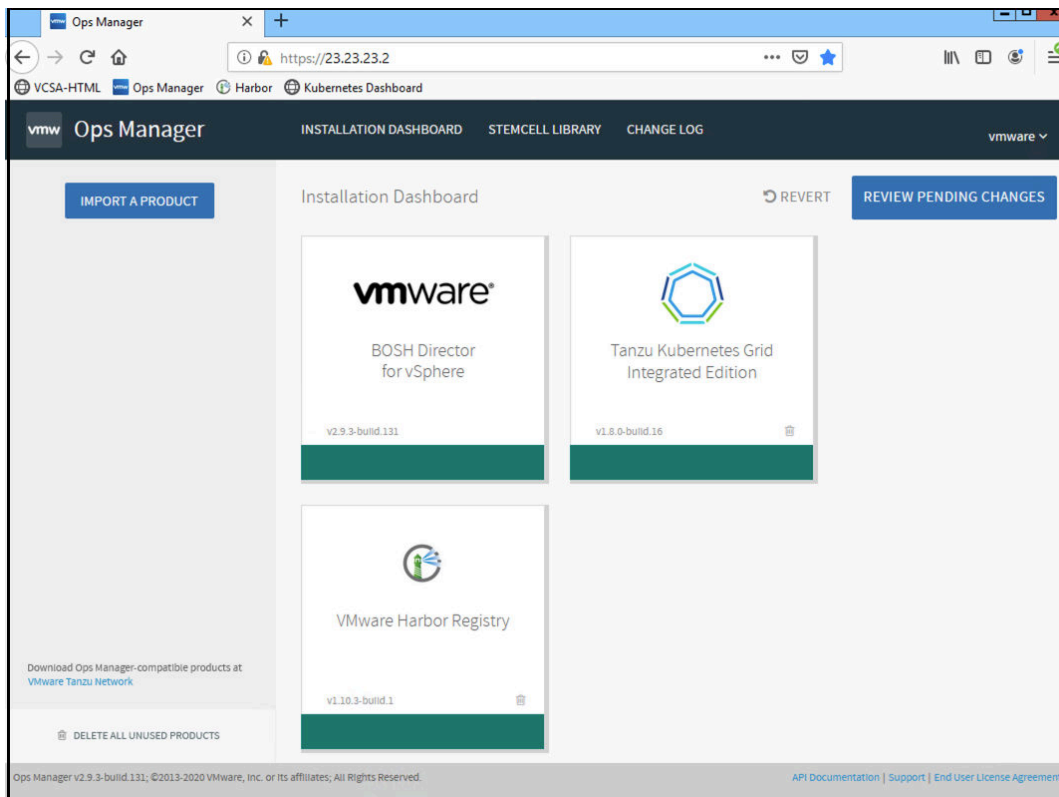
- 1 Open PuTTY session to CLI-VM. See [Accessing CLI VM](#).

- 2 Set the UAAC target.

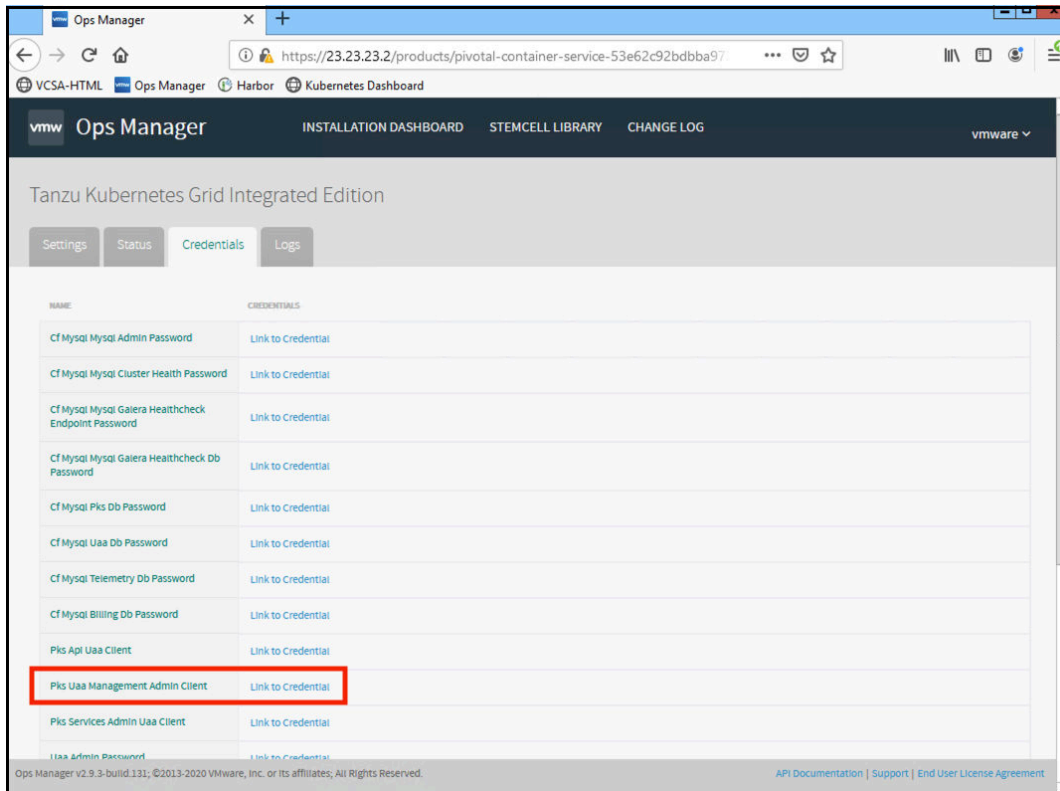
```
root@cli-vm:~# uaac target https://uaa.corp.local:8443 --skip-ssl-validation
```

```
Target: https://uaa.corp.local:8443  
Context: admin, from client admin
```

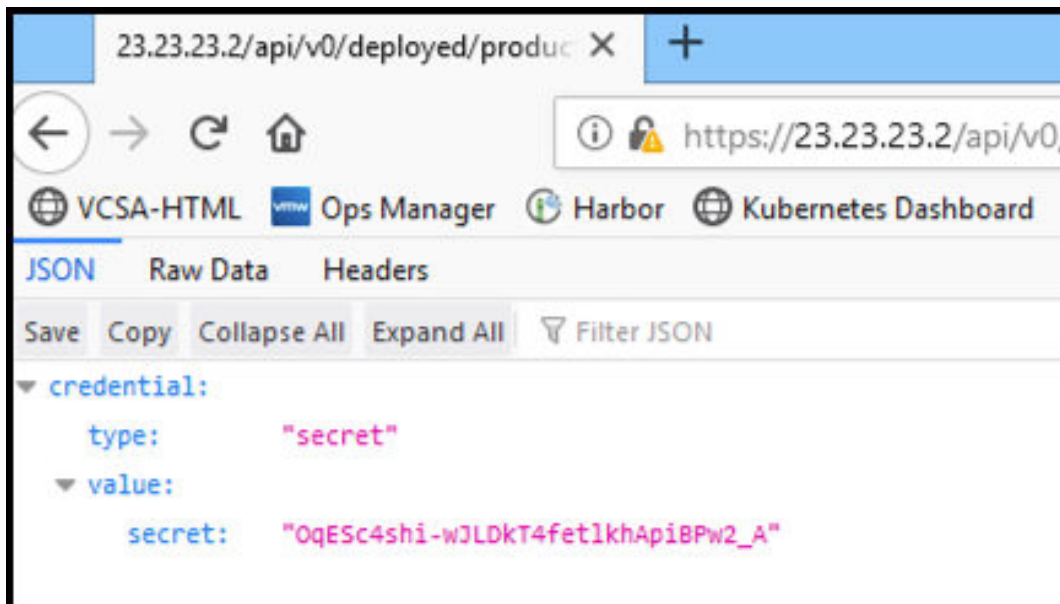
- 3 Log in to PCF Ops Manager. See [Accessing PCF Ops Manager](#).



- Go to **Tanzu Kubernetes Grid Integrated Edition > Credentials > Pks Uaa Management Admin Client** and click **Link to Credential**.



- Copy the secret value (*OqESc4shi-wJLDkT4fetlkhApiBPw2_A*) and paste it on a clipboard.



Note The secret value depends on the TKGI lab version.

- 6 Go to the CLI-VM and get the UAAC token by running this command:

```
root@cli-vm:~# uaac token client get admin -s <secret_value>
```

For example:

```
root@cli-vm:~# uaac token client get admin -s 0qESc4shi-wJLDkT4fetlkhApiBPw2_A
```

```
Successfully fetched token via client credentials grant.
Target: https://uaa.corp.local:8443
Context: admin, from client admin
```

- 7 Log in to PKS

```
root@cli-vm:~# pks login -a uaa.corp.local -u vmware -p VMware1! -k
```

```
API Endpoint: uaa.corp.local
User: vmware
```

Accessing PCF Ops Manager


The PCF Ops Manager can be accessed using the following details provided. The IP address provided is constant and will not change.

- 1 From the control center Windows VM, open Firefox, and click the **Ops Manager** bookmark (<https://23.23.23.2/uaa/login>).

Login credentials:

- User name: *vmware*
- Password: *VMware1!*

- 2 If you are asked for passphrase before providing the username or password, enter the passphrase as *VMware1!*

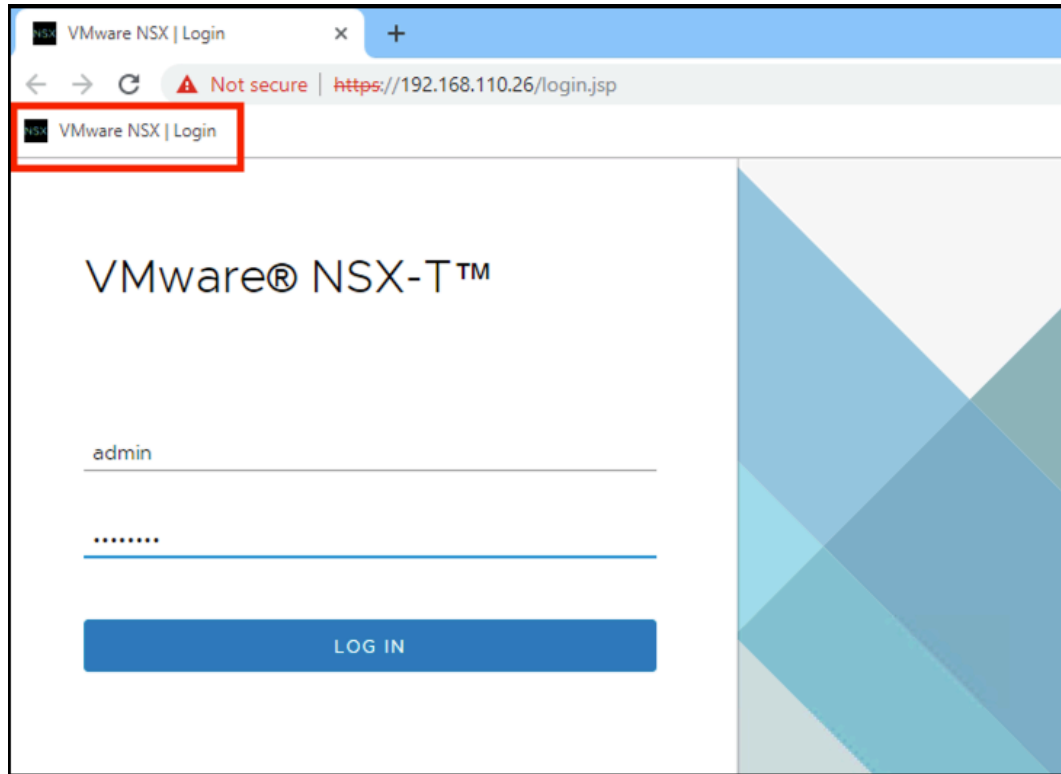


Accessing NSX-T Manager

From the control center Windows VM, open Google Chrome, and click the **NSX-Mgr** bookmark (<https://192.168.110.26>).

Login credentials:

- User name: *admin*
- Password: *VMware1!*



Accessing Kubernetes Dashboard

To access the Kubernetes dashboard, perform the following steps:

Procedure

- 1 Start a PuTTY session to root@cli-vm.
 - a Go to **Expand Connection > SSH > Tunnels**.
 - b Add a forwarded port with source port being 8001 and destination being localhost:8001. Click **Add**.
 - c Click **Open**. A prompt asking you to enter the password for root@192.168.110.7 appears. Enter the password (*VMware1!*).

- 2 Create a cluster (*test-cluster*). For details, see [Accessing TKGI Environment](#) and [Creating a PKS Cluster](#) sections.

```
root@cli-vm:~# pks create-cluster test-cluster --external-hostname ext-test-cluster.corp.local --plan small --num-nodes 3
```

Note The number of nodes and plan configuration might differ according to your application requirements.

- 3 Configure the cluster (*test-cluster*) to be used with your application.

```
root@cli-vm:~# pks cluster test-cluster
```


Check output for Kubernetes Master IP(s).

Edit the `/etc/hosts` file and add the line

'Kubernetes_Master_IP(s) ext-test-cluster.corp.local'.

```
root@cli-vm:~# vi /etc/hosts
```

- 4 Deploy Kubernetes Dashboard.

```
root@cli-vm:~# kubectl apply -f https://raw.githubusercontent.com/kubernetes/dashboard/v2.0.0-beta8/aio/deploy/recommended.yaml
```

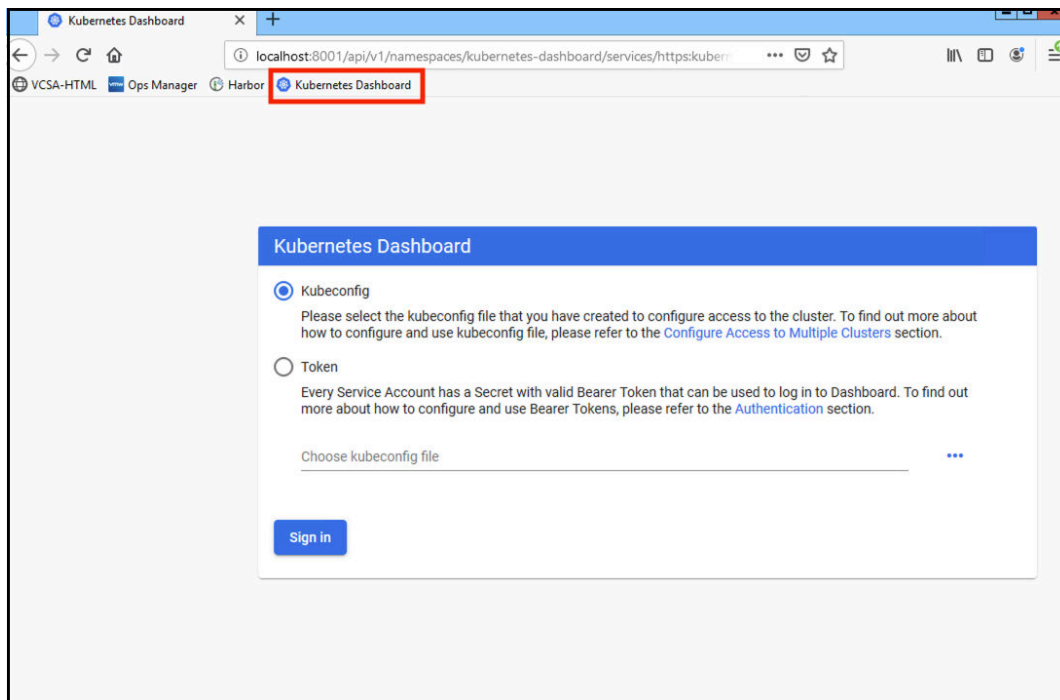
- 5 Start kubectl proxy.

```
root@cli-vm:~# kubectl proxy
```

Starting to serve on 127.0.0.1:8001

- 6 From the control center Windows VM, open Firefox, and click the **Kubernetes Dashboard** bookmark.

<http://localhost:8001/api/v1/namespaces/kubernetes-dashboard/services/https:kubernetes-dashboard:/proxy/#!/login>



- 7 From the **Kubernetes Dashboard** selection, select **Kubeconfig**, and provide the config file path. The config file is located at `/root/.kube/config` on the CLI-VM. Copy the file from the CLI-VM to the control center Windows VM using WinSCP.
- 8 Click **SIGN IN**.

- 9 Select your application-related Namespace from the drop-down menu in the Dashboard and check whether your application-related information is shown in the Kubernetes Dashboard correctly in green.

The screenshot shows the Kubernetes Dashboard interface. The left sidebar contains a navigation menu with sections: Cluster, Namespace (set to 'default'), Overview, Workloads, Discovery and Load Balancing, and Config and Storage. The 'Workloads' section is active, showing 'Deployments' and 'Pods'.

Workload Status

Three large green circles represent the status of Deployments, Pods, and Replica Sets, all indicating a healthy state.

Deployments

Name	Namespace	Labels	Pods	Age	Images
frontend	default	-	3 / 3	5 minutes	harbor.corp.local/testproject/gb-frontend:v4
redis-slave	default	-	2 / 2	5 minutes	harbor.corp.local/testproject/gb-redis-slave:v1
redis-master	default	-	1 / 1	5 minutes	harbor.corp.local/testproject/redis-e2e

Pods

Name	Namespace	Labels	Node	Status	Restart	CPU Usage (cores)	Memory Usage (bytes)	Age
frontend-c7f699c8d-k95zc	default	app: guestbook, pod-template-hash: c7f699	d54f34fec5-49a0bc-72	Running	0	1.00m	9.70Mi	5 minutes

Appendix A: Sample Docker Bench Security Script Output

13

```
# -----
# Docker Bench for Security v1.3.4
#
# Docker, Inc. (c) 2015-
#
# Checks for dozens of common best-practices around deploying Docker containers in production.
# Inspired by the CIS Docker Community Edition Benchmark v1.1.0.
# -----

Initializing Fri Nov 30 16:25:57 PST 2018

[INFO] 1 - Host Configuration
[WARN] 1.1 - Ensure a separate partition for containers has been created
[NOTE] 1.2 - Ensure the container host has been Hardened
[INFO] 1.3 - Ensure Docker is up to date
[INFO]      * Using 18.06.1, verify is it up to date as deemed necessary
[INFO]      * Your operating system vendor may provide support and security maintenance for Docker
[INFO] 1.4 - Ensure only trusted users are allowed to control Docker daemon
[INFO]      * docker:x:999:
[WARN] 1.5 - Ensure auditing is configured for the Docker daemon
[WARN] 1.6 - Ensure auditing is configured for Docker files and directories - /var/lib/docker
[WARN] 1.7 - Ensure auditing is configured for Docker files and directories - /etc/docker
[WARN] 1.8 - Ensure auditing is configured for Docker files and directories - docker.service
[WARN] 1.9 - Ensure auditing is configured for Docker files and directories - docker.socket
[WARN] 1.10 - Ensure auditing is configured for Docker files and directories - /etc/default/docker
[INFO] 1.11 - Ensure auditing is configured for Docker files and directories - /etc/docker/
daemon.json
[INFO]      * File not found
[WARN] 1.12 - Ensure auditing is configured for Docker files and directories - /usr/bin/docker-
containerd
[WARN] 1.13 - Ensure auditing is configured for Docker files and directories - /usr/bin/docker-runc

[INFO] 2 - Docker daemon configuration
[WARN] 2.1 - Ensure network traffic is restricted between containers on the default bridge
[PASS] 2.2 - Ensure the logging level is set to 'info'
[PASS] 2.3 - Ensure Docker is allowed to make changes to iptables
[PASS] 2.4 - Ensure insecure registries are not used
[PASS] 2.5 - Ensure aufs storage driver is not used
[INFO] 2.6 - Ensure TLS authentication for Docker daemon is configured
```

```

[INFO]      * Docker daemon not listening on TCP
[INFO] 2.7  - Ensure the default ulimit is configured appropriately
[INFO]      * Default ulimit doesn't appear to be set
[WARN] 2.8  - Enable user namespace support
[PASS] 2.9  - Ensure the default cgroup usage has been confirmed
[PASS] 2.10 - Ensure base device size is not changed until needed
[WARN] 2.11 - Ensure that authorization for Docker client commands is enabled
[WARN] 2.12 - Ensure centralized and remote logging is configured
[INFO] 2.13 - Ensure operations on legacy registry (v1) are Disabled (Deprecated)
[WARN] 2.14 - Ensure live restore is Enabled
[WARN] 2.15 - Ensure Userland Proxy is Disabled
[PASS] 2.16 - Ensure daemon-wide custom seccomp profile is applied, if needed
[PASS] 2.17 - Ensure experimental features are avoided in production
[WARN] 2.18 - Ensure containers are restricted from acquiring new privileges

[INFO] 3 - Docker daemon configuration files
[PASS] 3.1  - Ensure that docker.service file ownership is set to root:root
[PASS] 3.2  - Ensure that docker.service file permissions are set to 644 or more restrictive
[PASS] 3.3  - Ensure that docker.socket file ownership is set to root:root
[PASS] 3.4  - Ensure that docker.socket file permissions are set to 644 or more restrictive
[PASS] 3.5  - Ensure that /etc/docker directory ownership is set to root:root
[PASS] 3.6  - Ensure that /etc/docker directory permissions are set to 755 or more restrictive
[PASS] 3.7  - Ensure that registry certificate file ownership is set to root:root
[WARN] 3.8  - Ensure that registry certificate file permissions are set to 444 or more restrictive
[WARN]      * Wrong permissions for /etc/docker/certs.d/
[INFO] 3.9  - Ensure that TLS CA certificate file ownership is set to root:root
[INFO]      * No TLS CA certificate found
[INFO] 3.10 - Ensure that TLS CA certificate file permissions are set to 444 or more restrictive
[INFO]      * No TLS CA certificate found
[INFO] 3.11 - Ensure that Docker server certificate file ownership is set to root:root
[INFO]      * No TLS Server certificate found
[INFO] 3.12 - Ensure that Docker server certificate file permissions are set to 444 or more
restrictive
[INFO]      * No TLS Server certificate found
[INFO] 3.13 - Ensure that Docker server certificate key file ownership is set to root:root
[INFO]      * No TLS Key found
[INFO] 3.14 - Ensure that Docker server certificate key file permissions are set to 400
[INFO]      * No TLS Key found
[PASS] 3.15 - Ensure that Docker socket file ownership is set to root:docker
[PASS] 3.16 - Ensure that Docker socket file permissions are set to 660 or more restrictive
[INFO] 3.17 - Ensure that daemon.json file ownership is set to root:root
[INFO]      * File not found
[INFO] 3.18 - Ensure that daemon.json file permissions are set to 644 or more restrictive
[INFO]      * File not found
[PASS] 3.19 - Ensure that /etc/default/docker file ownership is set to root:root
[PASS] 3.20 - Ensure that /etc/default/docker file permissions are set to 644 or more restrictive

[INFO] 4 - Container Images and Build File
[WARN] 4.1  - Ensure a user for the container has been created
[WARN]      * Running as root: competent_feynman
[NOTE] 4.2  - Ensure that containers use trusted base images
[NOTE] 4.3  - Ensure unnecessary packages are not installed in the container
[NOTE] 4.4  - Ensure images are scanned and rebuilt to include security patches

```

```

[PASS] 4.5 - Ensure Content trust for Docker is Enabled
[WARN] 4.6 - Ensure HEALTHCHECK instructions have been added to the container image
[WARN]      * No Healthcheck found: [gcr.io/google-samples/gb-frontend:v4 harbor.corp.local/
testproject/gb-frontend:v4]
[WARN]      * No Healthcheck found: [gcr.io/google-samples/gb-frontend:v4 harbor.corp.local/
testproject/gb-frontend:v4]
[INFO] 4.7 - Ensure update instructions are not use alone in the Dockerfile
[INFO]      * Update instruction found: [gcr.io/google-samples/gb-frontend:v4 harbor.corp.local/
testproject/gb-frontend:v4]
[INFO]      * Update instruction found: [gcr.io/google-samples/gb-frontend:v4 harbor.corp.local/
testproject/gb-frontend:v4]
[NOTE] 4.8 - Ensure setuid and setgid permissions are removed in the images
[INFO] 4.9 - Ensure COPY is used instead of ADD in Dockerfile
[INFO]      * ADD in image history: [gcr.io/google-samples/gb-frontend:v4 harbor.corp.local/
testproject/gb-frontend:v4]
[INFO]      * ADD in image history: [gcr.io/google-samples/gb-frontend:v4 harbor.corp.local/
testproject/gb-frontend:v4]
[NOTE] 4.10 - Ensure secrets are not stored in Dockerfiles
[NOTE] 4.11 - Ensure verified packages are only Installed

[INFO] 5 - Container Runtime
[PASS] 5.1 - Ensure AppArmor Profile is Enabled
[WARN] 5.2 - Ensure SELinux security options are set, if applicable
[WARN]      * No SecurityOptions Found: competent_feynman
[PASS] 5.3 - Ensure Linux Kernel Capabilities are restricted within containers
[PASS] 5.4 - Ensure privileged containers are not used
[PASS] 5.5 - Ensure sensitive host system directories are not mounted on containers
[PASS] 5.6 - Ensure ssh is not run within containers
[PASS] 5.7 - Ensure privileged ports are not mapped within containers
[NOTE] 5.8 - Ensure only needed ports are open on the container
[PASS] 5.9 - Ensure the host's network namespace is not shared
[WARN] 5.10 - Ensure memory usage for container is limited
[WARN]      * Container running without memory restrictions: competent_feynman
[WARN] 5.11 - Ensure CPU priority is set appropriately on the container
[WARN]      * Container running without CPU restrictions: competent_feynman
[WARN] 5.12 - Ensure the container's root filesystem is mounted as read only
[WARN]      * Container running with root FS mounted R/W: competent_feynman
[PASS] 5.13 - Ensure incoming container traffic is binded to a specific host interface
[WARN] 5.14 - Ensure 'on-failure' container restart policy is set to '5'
[WARN]      * MaximumRetryCount is not set to 5: competent_feynman
[PASS] 5.15 - Ensure the host's process namespace is not shared
[PASS] 5.16 - Ensure the host's IPC namespace is not shared
[PASS] 5.17 - Ensure host devices are not directly exposed to containers
[INFO] 5.18 - Ensure the default ulimit is overwritten at runtime, only if needed
[INFO]      * Container no default ulimit override: competent_feynman
[PASS] 5.19 - Ensure mount propagation mode is not set to shared
[PASS] 5.20 - Ensure the host's UTS namespace is not shared
[PASS] 5.21 - Ensure the default seccomp profile is not Disabled
[NOTE] 5.22 - Ensure docker exec commands are not used with privileged option
[NOTE] 5.23 - Ensure docker exec commands are not used with user option
[PASS] 5.24 - Ensure cgroup usage is confirmed
[WARN] 5.25 - Ensure the container is restricted from acquiring additional privileges
[WARN]      * Privileges not restricted: competent_feynman
[WARN] 5.26 - Ensure container health is checked at runtime

```

```

[WARN]      * Health check not set: competent_feynman
[INFO] 5.27 - Ensure docker commands always get the latest version of the image
[WARN] 5.28 - Ensure PIDs cgroup limit is used
[WARN]      * PIDs limit not set: competent_feynman
[INFO] 5.29 - Ensure Docker's default bridge docker0 is not used
[INFO]      * Container in docker0 network: competent_feynman
[PASS] 5.30 - Ensure the host's user namespaces is not shared
[PASS] 5.31 - Ensure the Docker socket is not mounted inside any containers

[INFO] 6 - Docker Security Operations
[INFO] 6.1 - Avoid image sprawl
[INFO]      * There are currently: 1 images
[INFO] 6.2 - Avoid container sprawl
[INFO]      * There are currently a total of 1 containers, with 1 of them currently running

[INFO] 7 - Docker Swarm Configuration
[PASS] 7.1 - Ensure swarm mode is not Enabled, if not needed
[PASS] 7.2 - Ensure the minimum number of manager nodes have been created in a swarm (Swarm mode not
enabled)
[PASS] 7.3 - Ensure swarm services are binded to a specific host interface (Swarm mode not enabled)
[PASS] 7.4 - Ensure data exchanged between containers are encrypted on different nodes on the
overlay network
[PASS] 7.5 - Ensure Docker's secret management commands are used for managing secrets in a Swarm
cluster (Swarm mode not enabled)
[PASS] 7.6 - Ensure swarm manager is run in auto-lock mode (Swarm mode not enabled)
[PASS] 7.7 - Ensure swarm manager auto-lock key is rotated periodically (Swarm mode not enabled)
[PASS] 7.8 - Ensure node certificates are rotated as appropriate (Swarm mode not enabled)
[PASS] 7.9 - Ensure CA certificates are rotated as appropriate (Swarm mode not enabled)
[PASS] 7.10 - Ensure management plane traffic has been separated from data plane traffic (Swarm mode
not enabled)

[INFO] Checks: 105
[INFO] Score: 19

```

Appendix B: TKGI Software BOM

14

This chapter includes the following topics:

- [TKGI Software BOM – 1.7 Lab](#)
- [TKGI Software BOM – 1.8 Lab](#)

TKGI Software BOM – 1.7 Lab

The following table provides the version and build information of different TKGI setup components. This lab is based on the TKGI 1.7 GA version.

Product	Version	Build
vSphere (ESXi)	6.7 U2	ESXi - 13006603
vSphere (VC)	6.7 GA	vCenter - 8217866
NSX-T	2.5.1	2.5.1.0.0.15314288
TKGI	1.7.0	1.7.0-build.26
PCF Ops Manager	2.9.0	v2.9.0-build.106
Harbor	1.10.1	v1.10.1-build.7
Bosh	2.9.0	v2.9.0-build.106
Kubernetes	1.16.7	

TKGI Software BOM – 1.8 Lab

The following table provides the version and build information of different TKGI setup components. This lab is based on the TKGI 1.8 GA version.

Product	Version	Build
vSphere (ESXi)	6.7 U3	ESXi -14320388
vSphere (VC)	6.7 GA	vCenter - 8217866
NSX-T	2.5.1	2.5.1.0.0.15314288
TKGI	1.8.0	1.8.0-build.16
PCF Ops Manager	2.9.3	v2.9.3-build.131

Product	Version	Build
Harbor	1.10.3	v1.10.3-build.1
Bosh	2.9.3	v2.9.3-build.131
Kubernetes	1.17.5	

Appendix C: TKGI Cluster Plans

15

Each master and worker nodes have the following plans.

TKGI Cluster Plans - 1.7 and 1.8 Labs

The following table provides information about the cluster plans for the TKGI 1.7 and 1.8 labs.

Plan Settings Name	Plan	CPU	Memory (GB)	Disk Size (GB)	Number of Nodes	
					Master	Worker
Plan 1	small	2	4	32	1	3
Plan 3	large	8	8	32	1	3
Plan 4	xlarge	4 - master	16 - master	32 - master	1	2
		4 - worker	16 - worker	128 - worker		
Plan 5	small6workers	2	4	32	1	6
Plan 7	large6workers	8	8	32	1	6
Plan 9	small3masters	2	4	32	3	4
Plan 10	large3masters	8	8	32	3	4